

Introduction to Machine Learning

Jiří Materna



About me

- Ph.D. in Natural Language Processing and Artificial Intelligence at Masaryk University
- 10 years at Seznam.cz (last 8 years as Head Of Research)
- Founder and lecturer at ML College
- Founder and co-organizer of ML Prague
- ML Freelance and consultant

Outline

Day 1

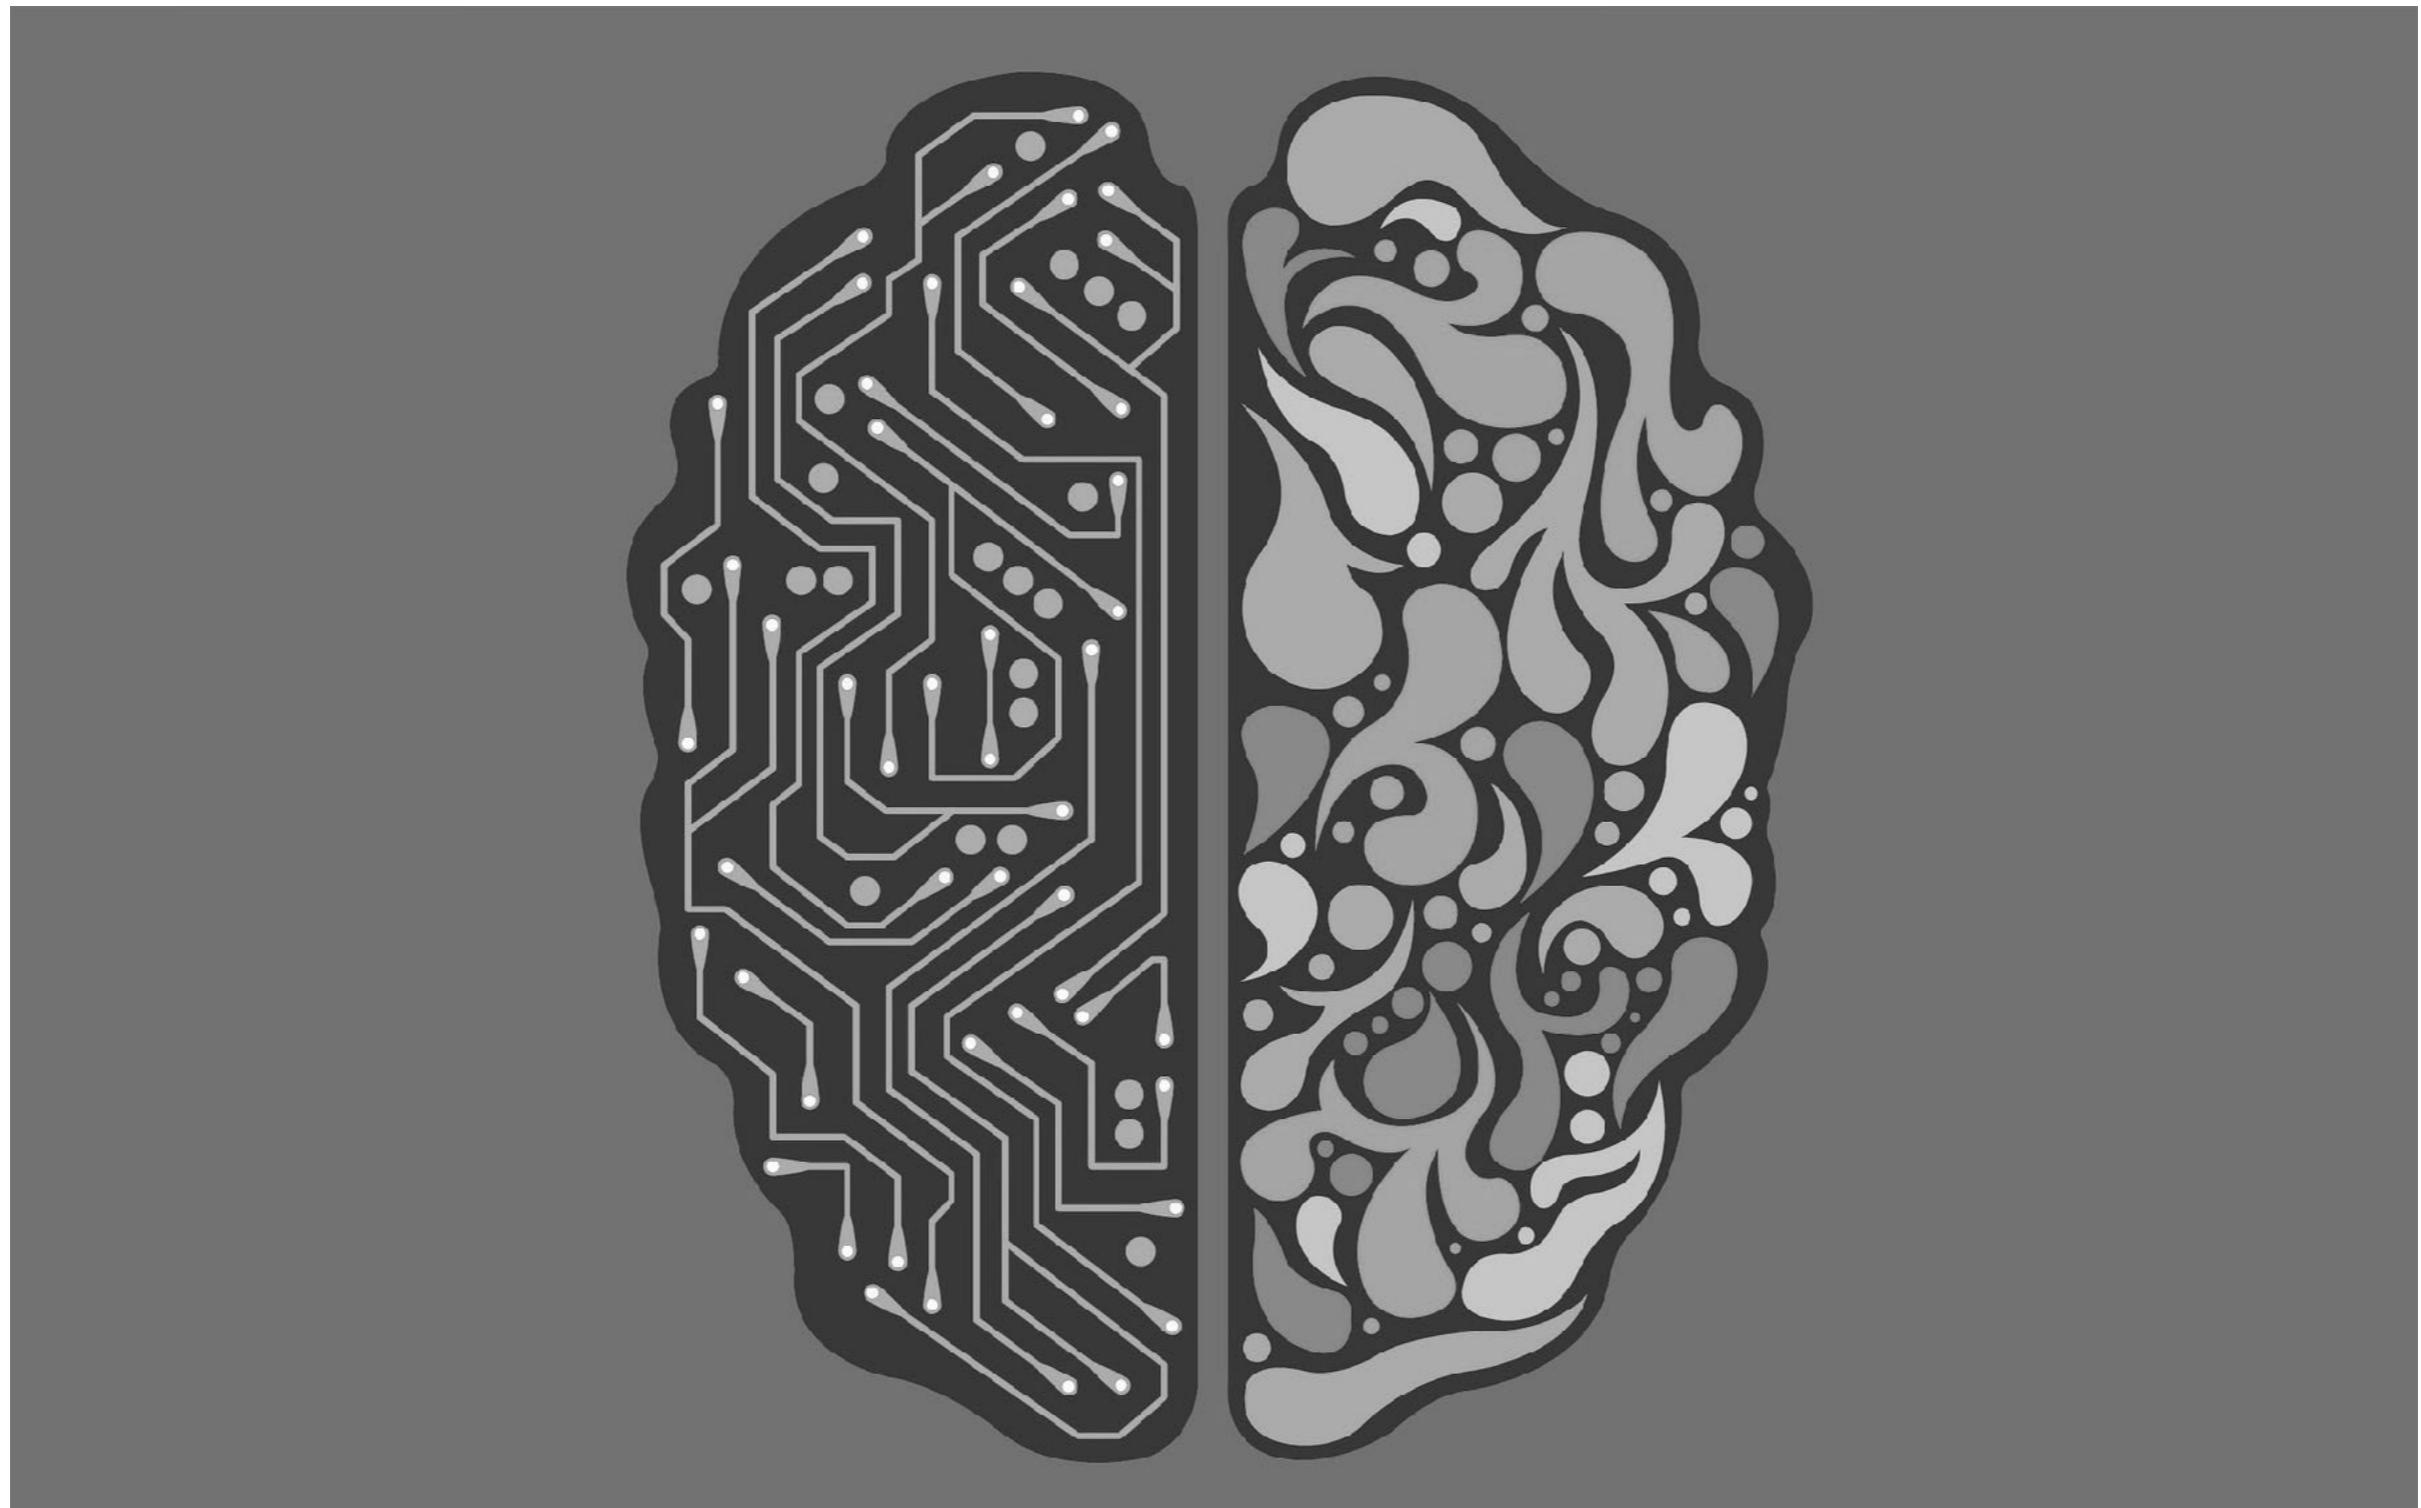
- Introduction to machine learning
- Types of ML tasks
- Data preparation
- Model evaluation
- Basic classification algorithms
- Scikit-learn tutorial
- Practical classification task
- Basic regression algorithms
- Regression model evaluation
- Practical regression task

Outline

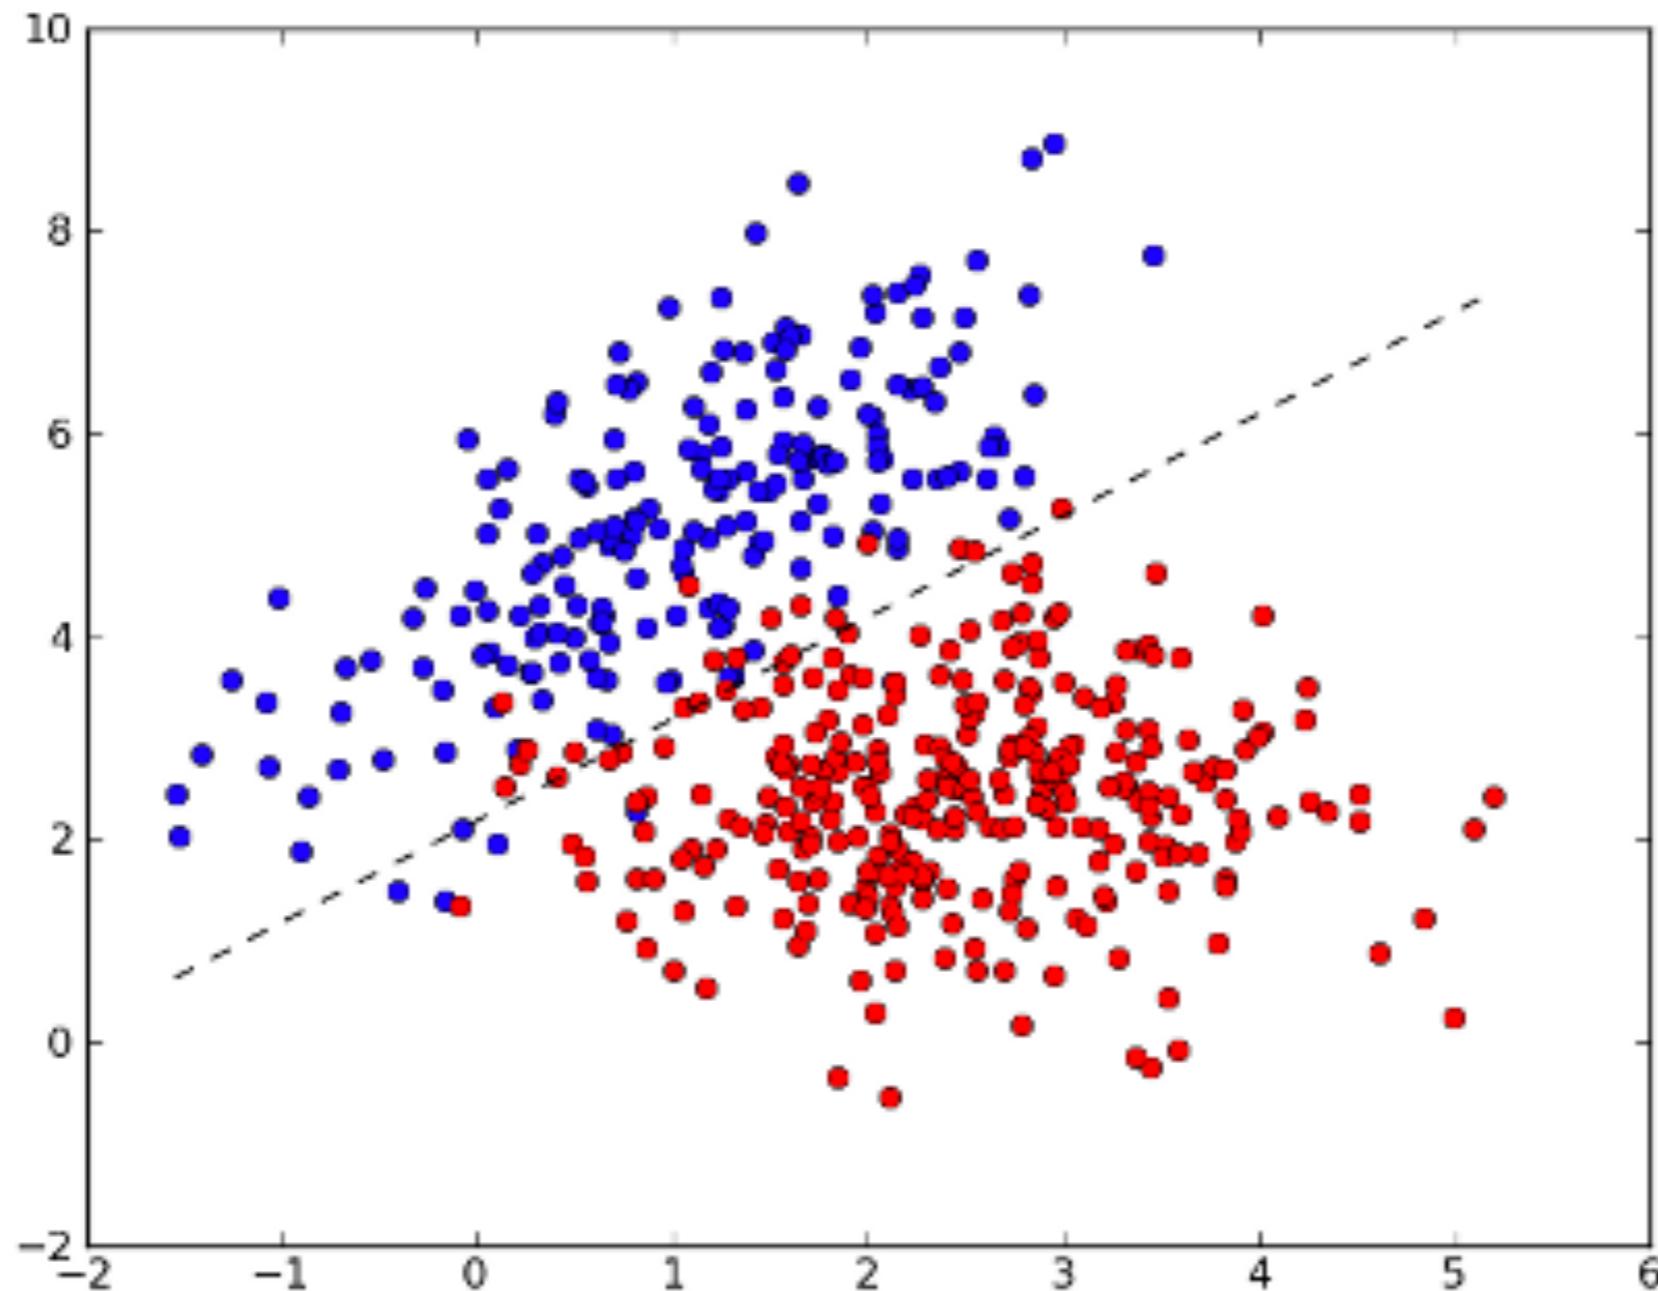
Day 2

- Dimensionality reduction
- Basic clustering algorithms
- Practical clustering task
- Introduction to neural networks
- Activation functions for neural networks
- Multilayered neural networks
- Methods for training neural networks
- Keras tutorial
- Practical classification and regression tasks solved using neural networks

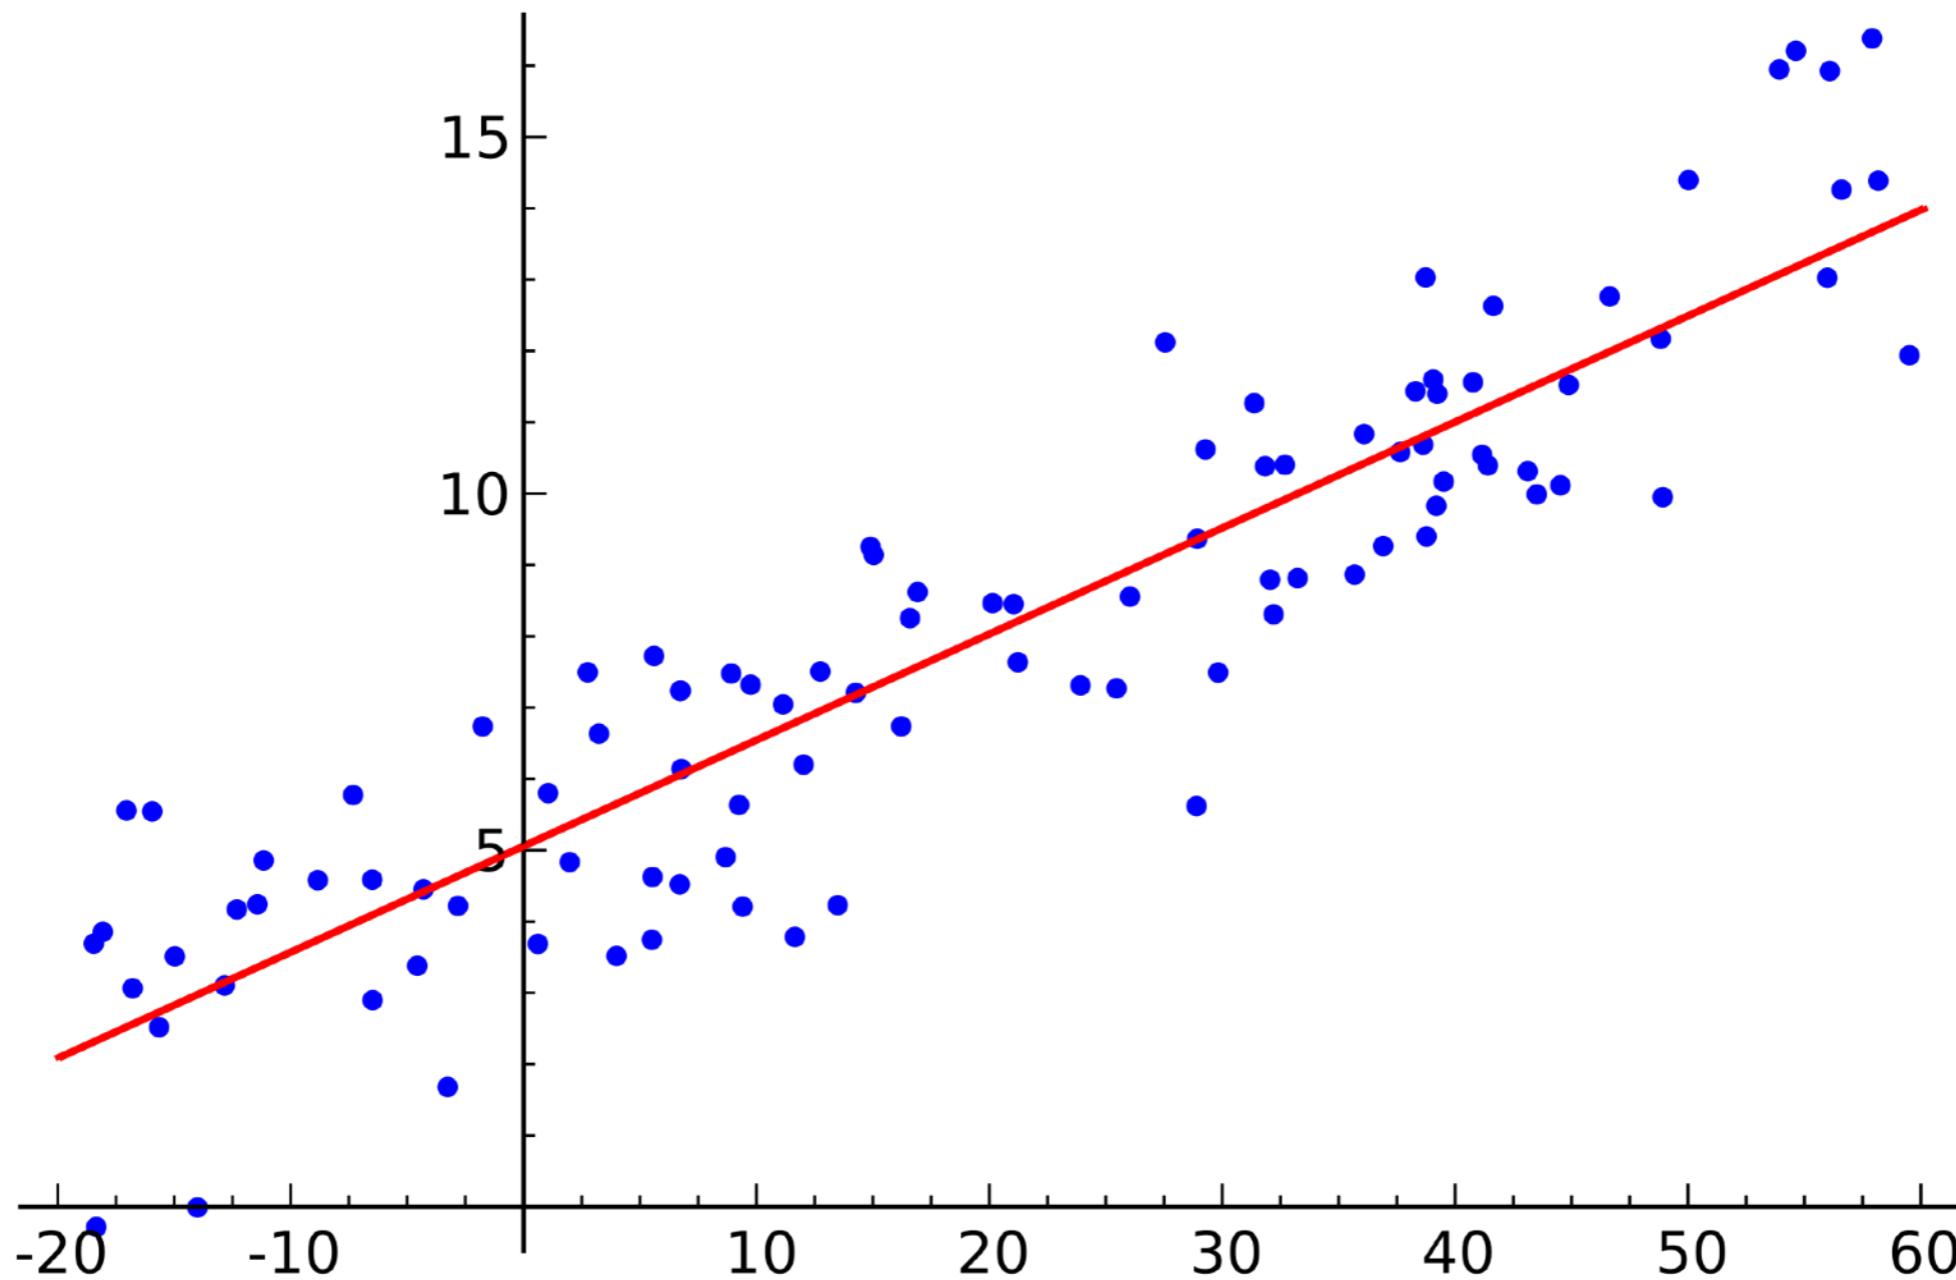
What is (not) machine learning?



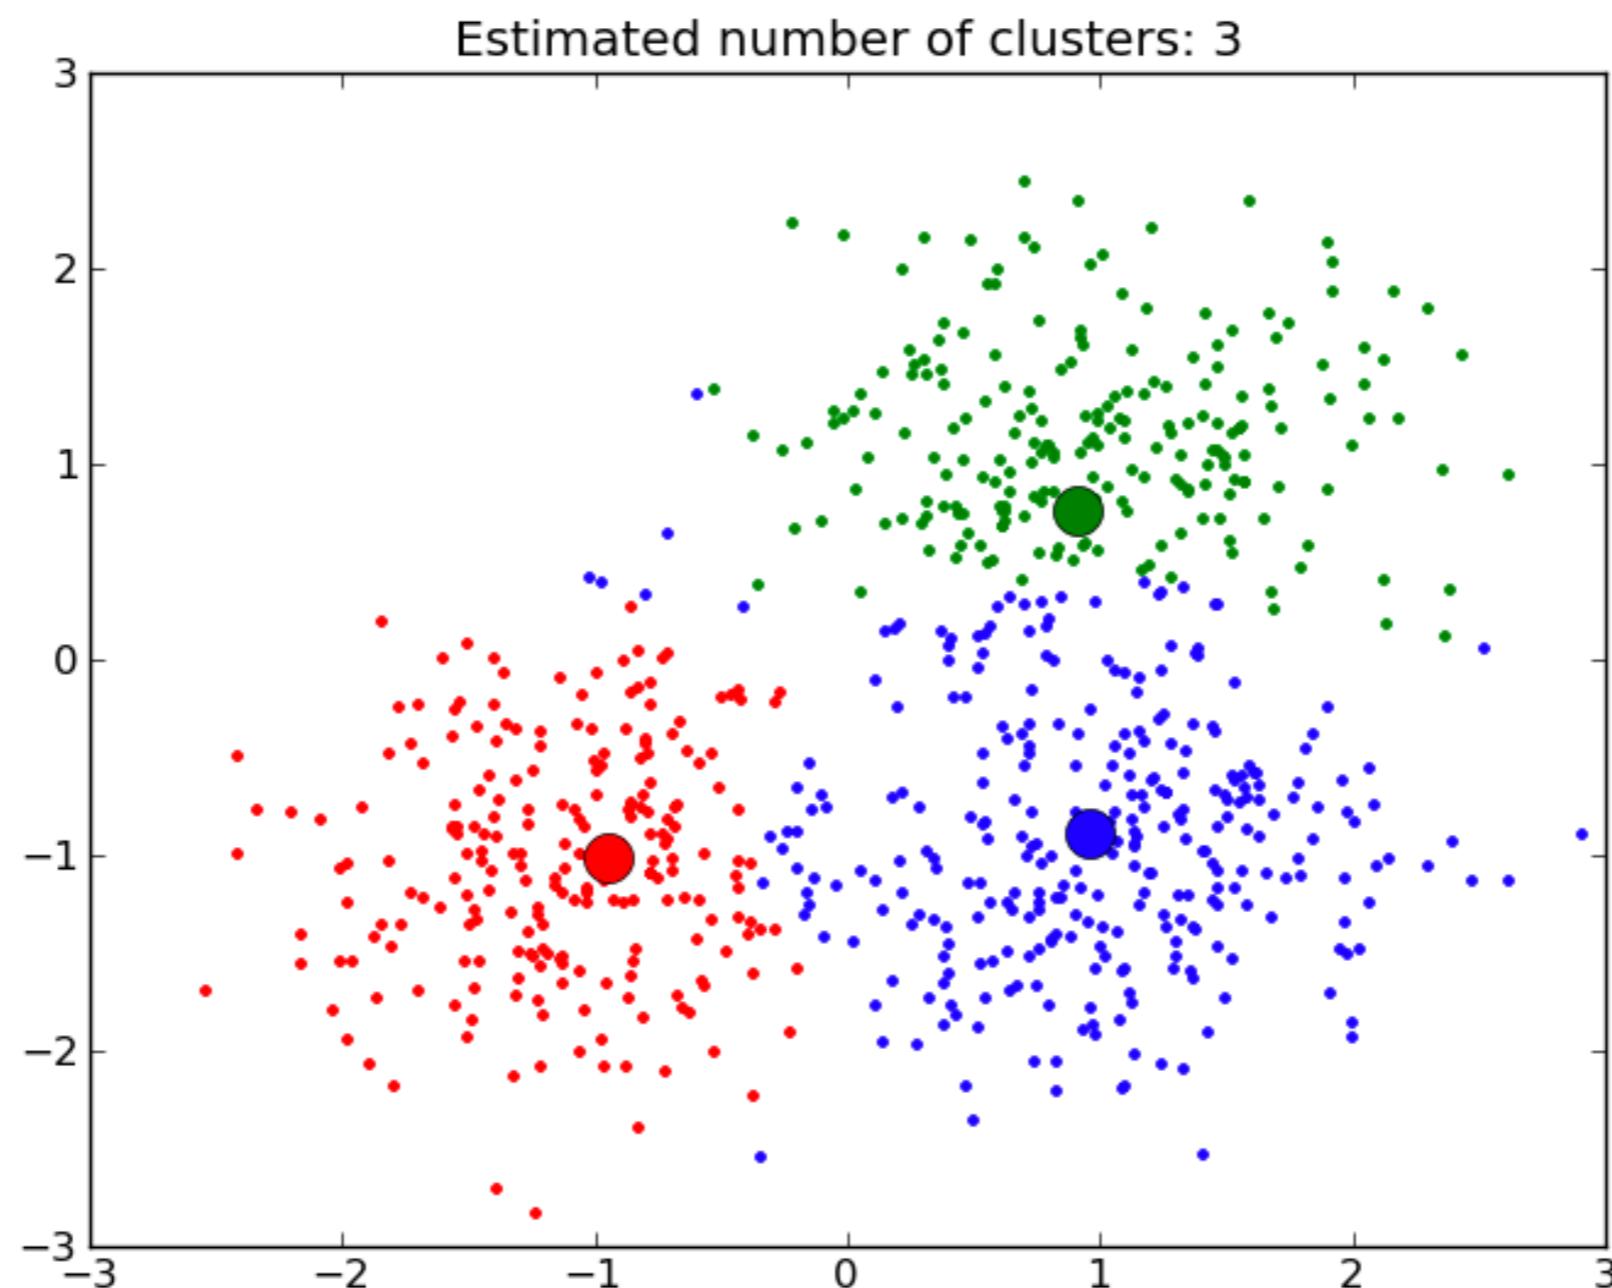
Classification



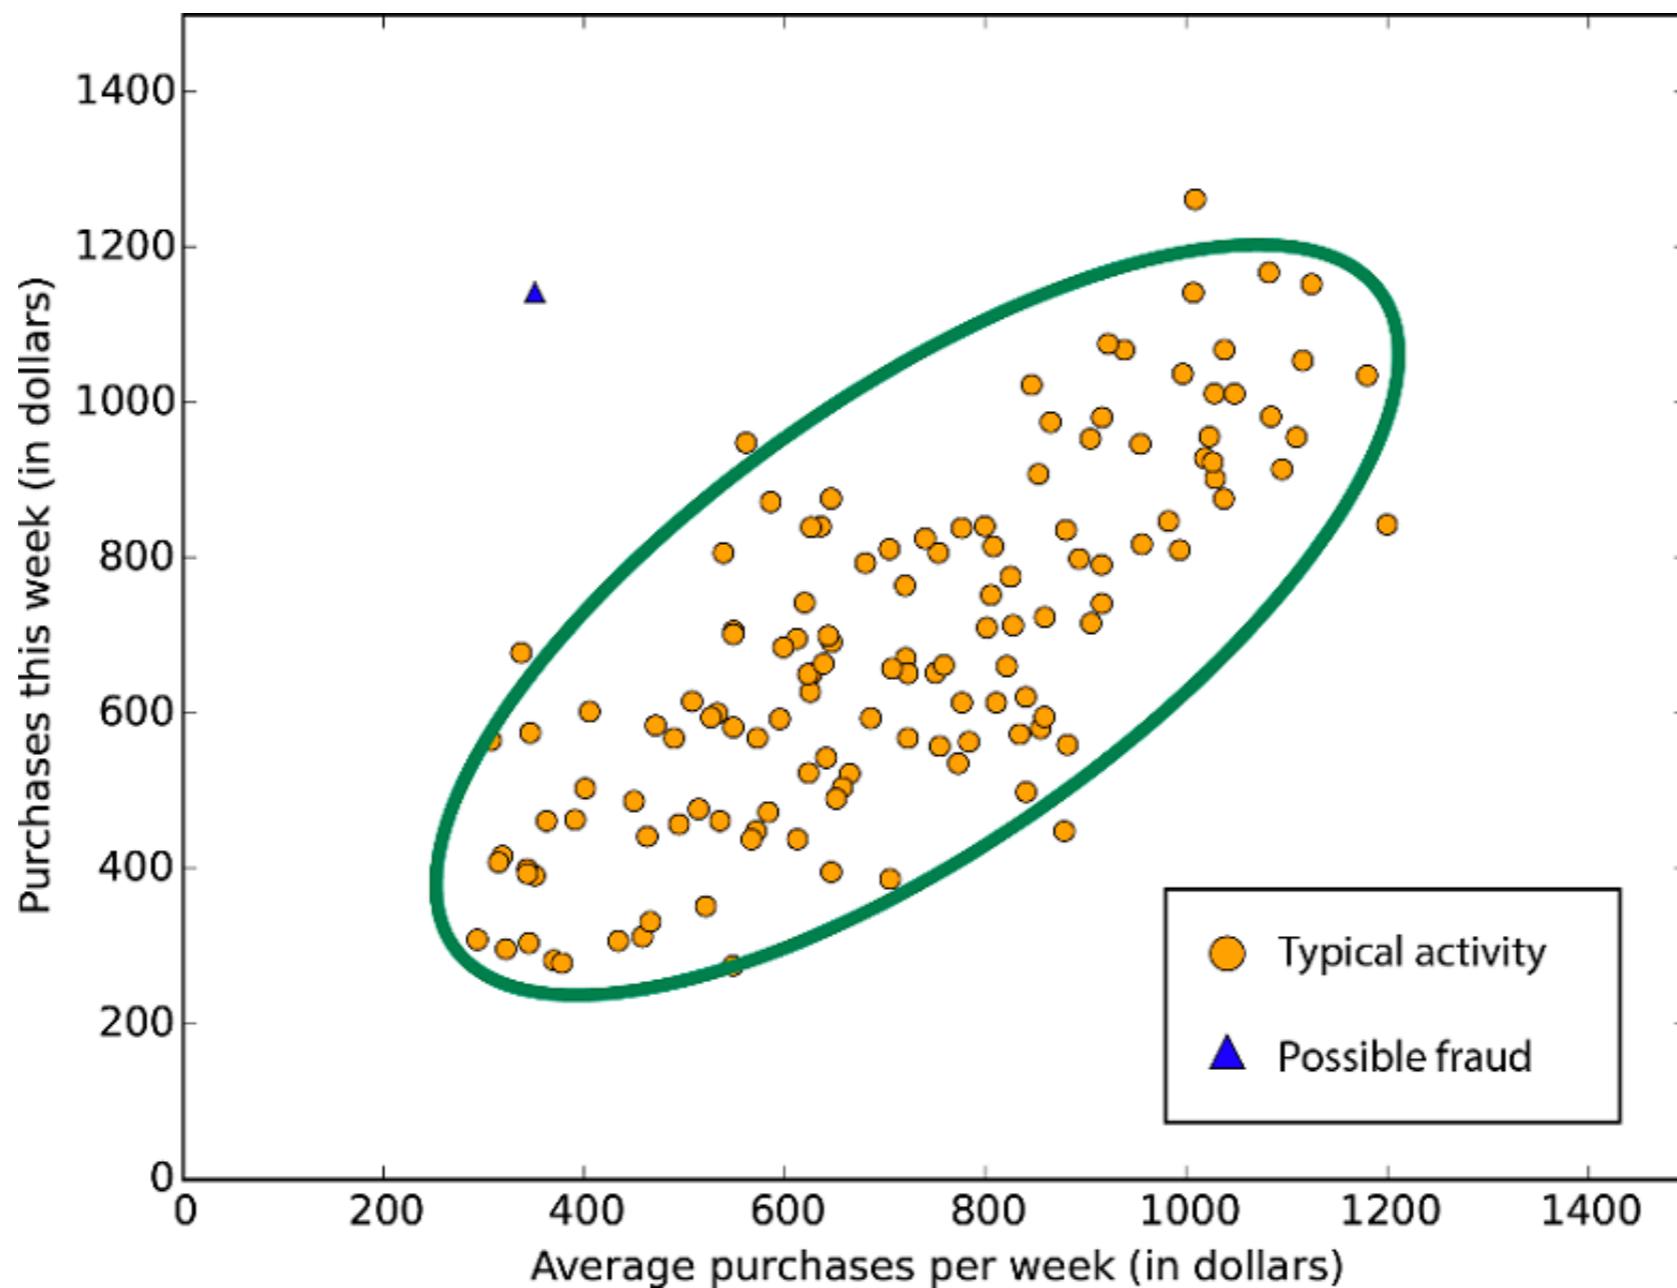
Regression



Clustering

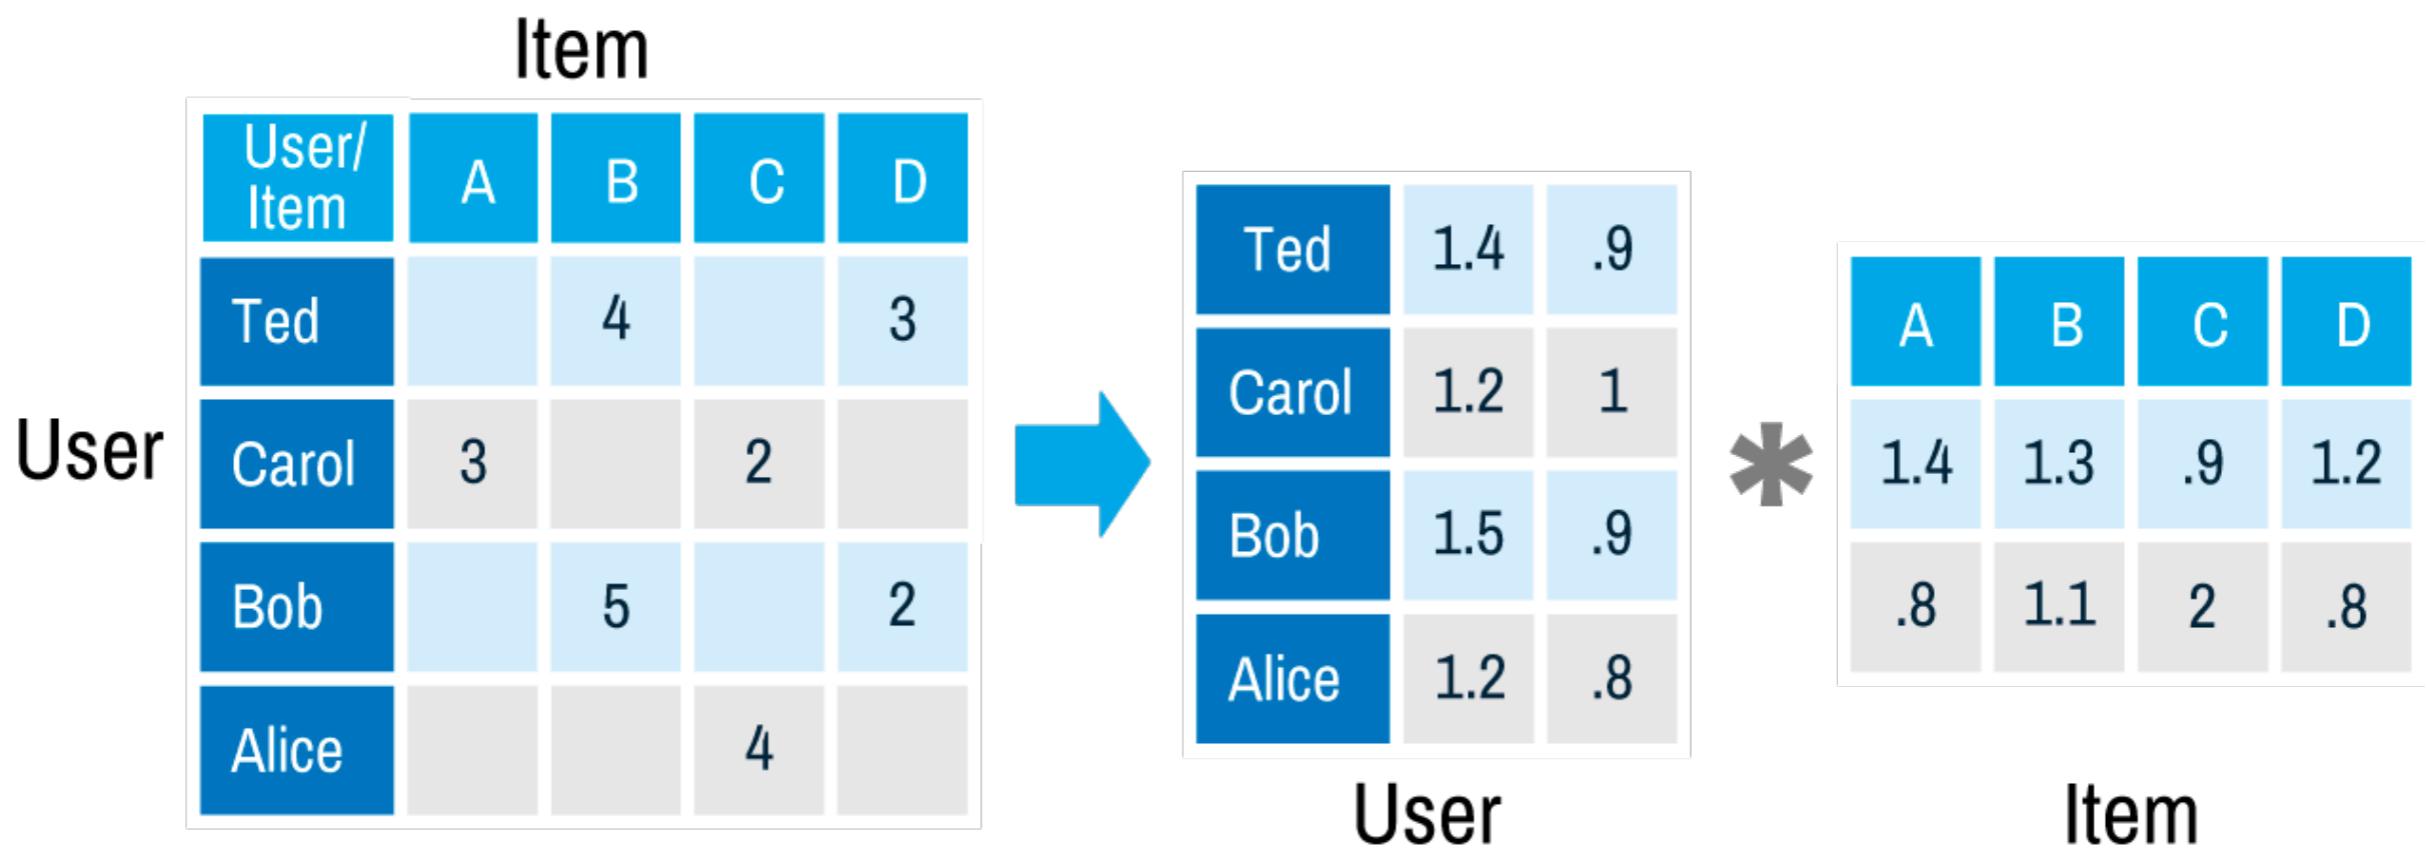


Anomaly detection

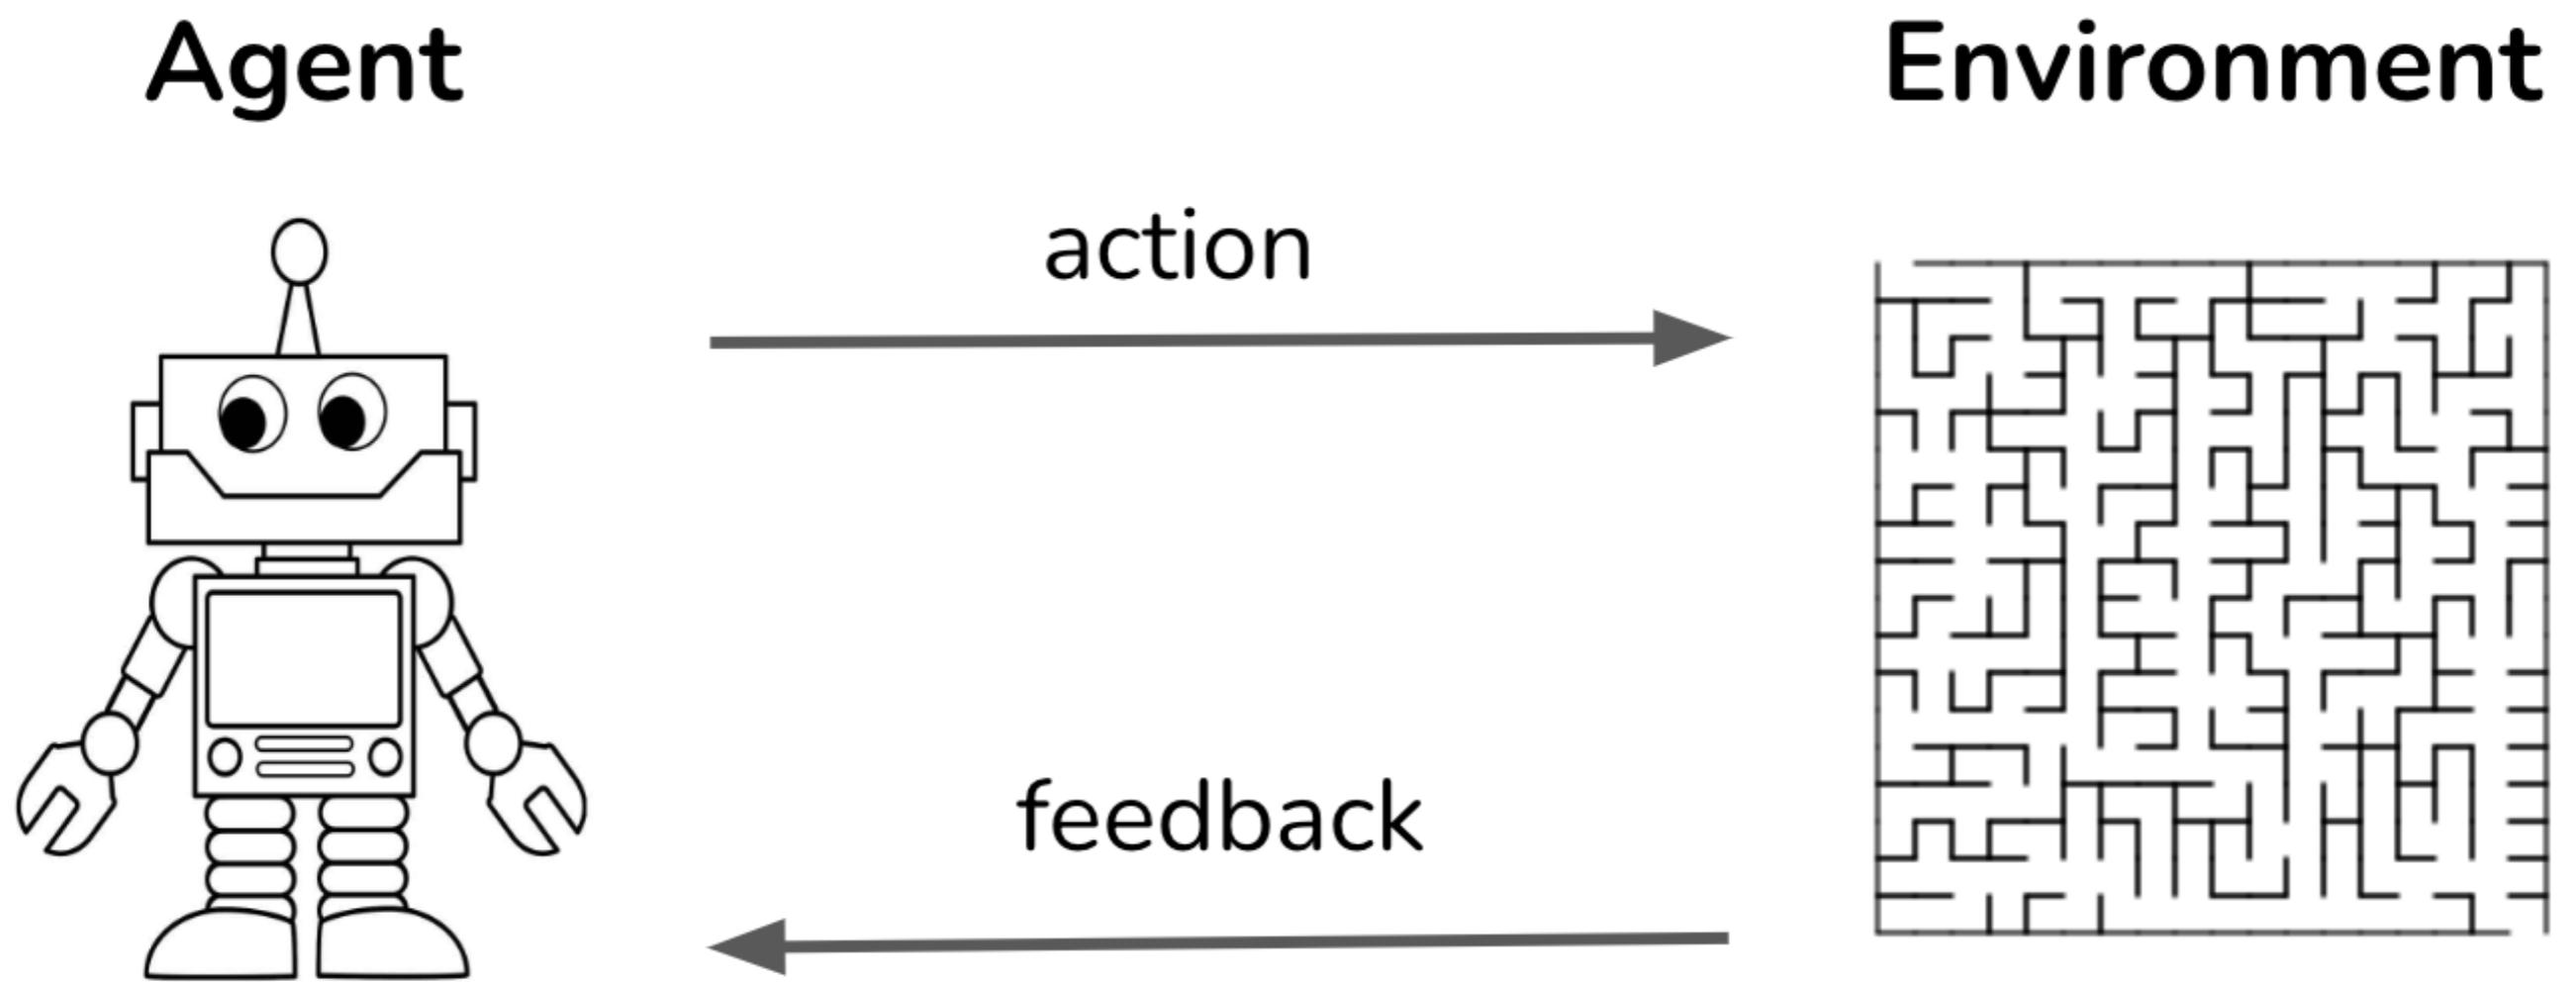


Source: <https://docs.microsoft.com/en-us/azure>

Recommendation

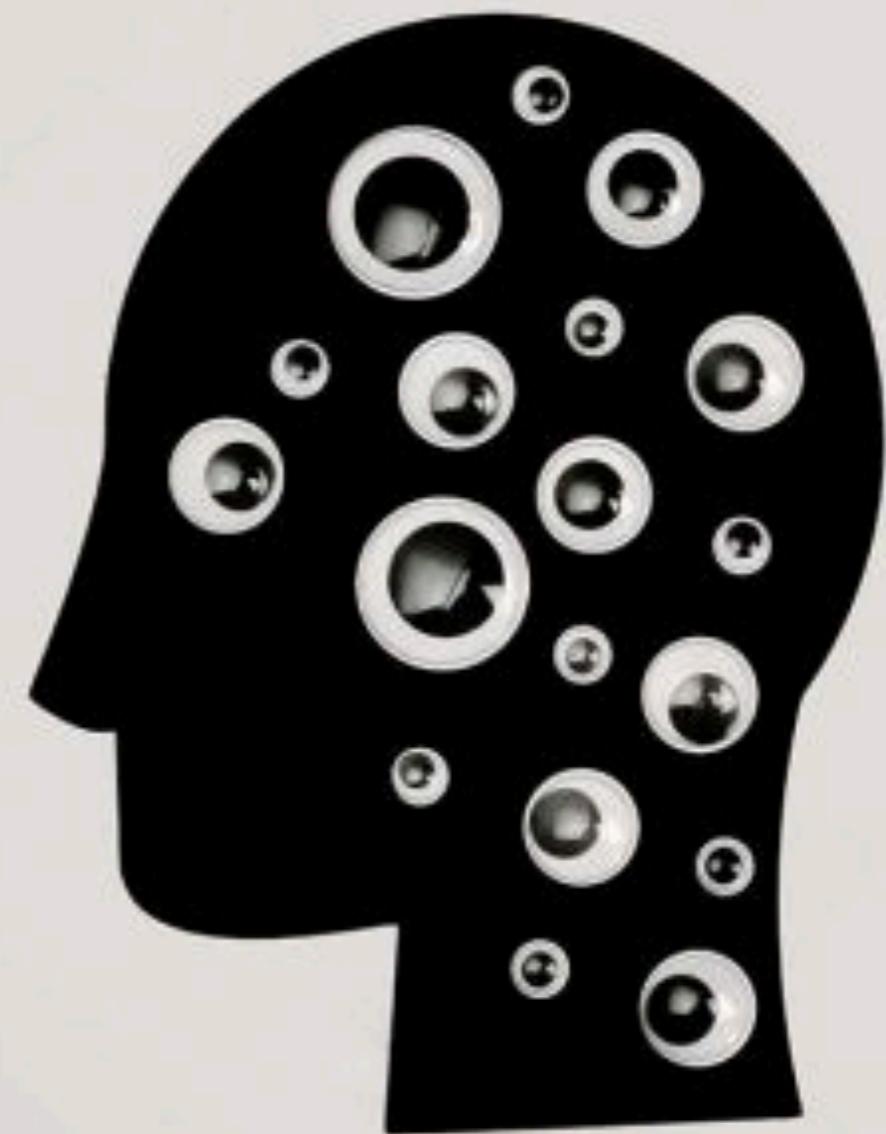


Reinforcement learning



Generative

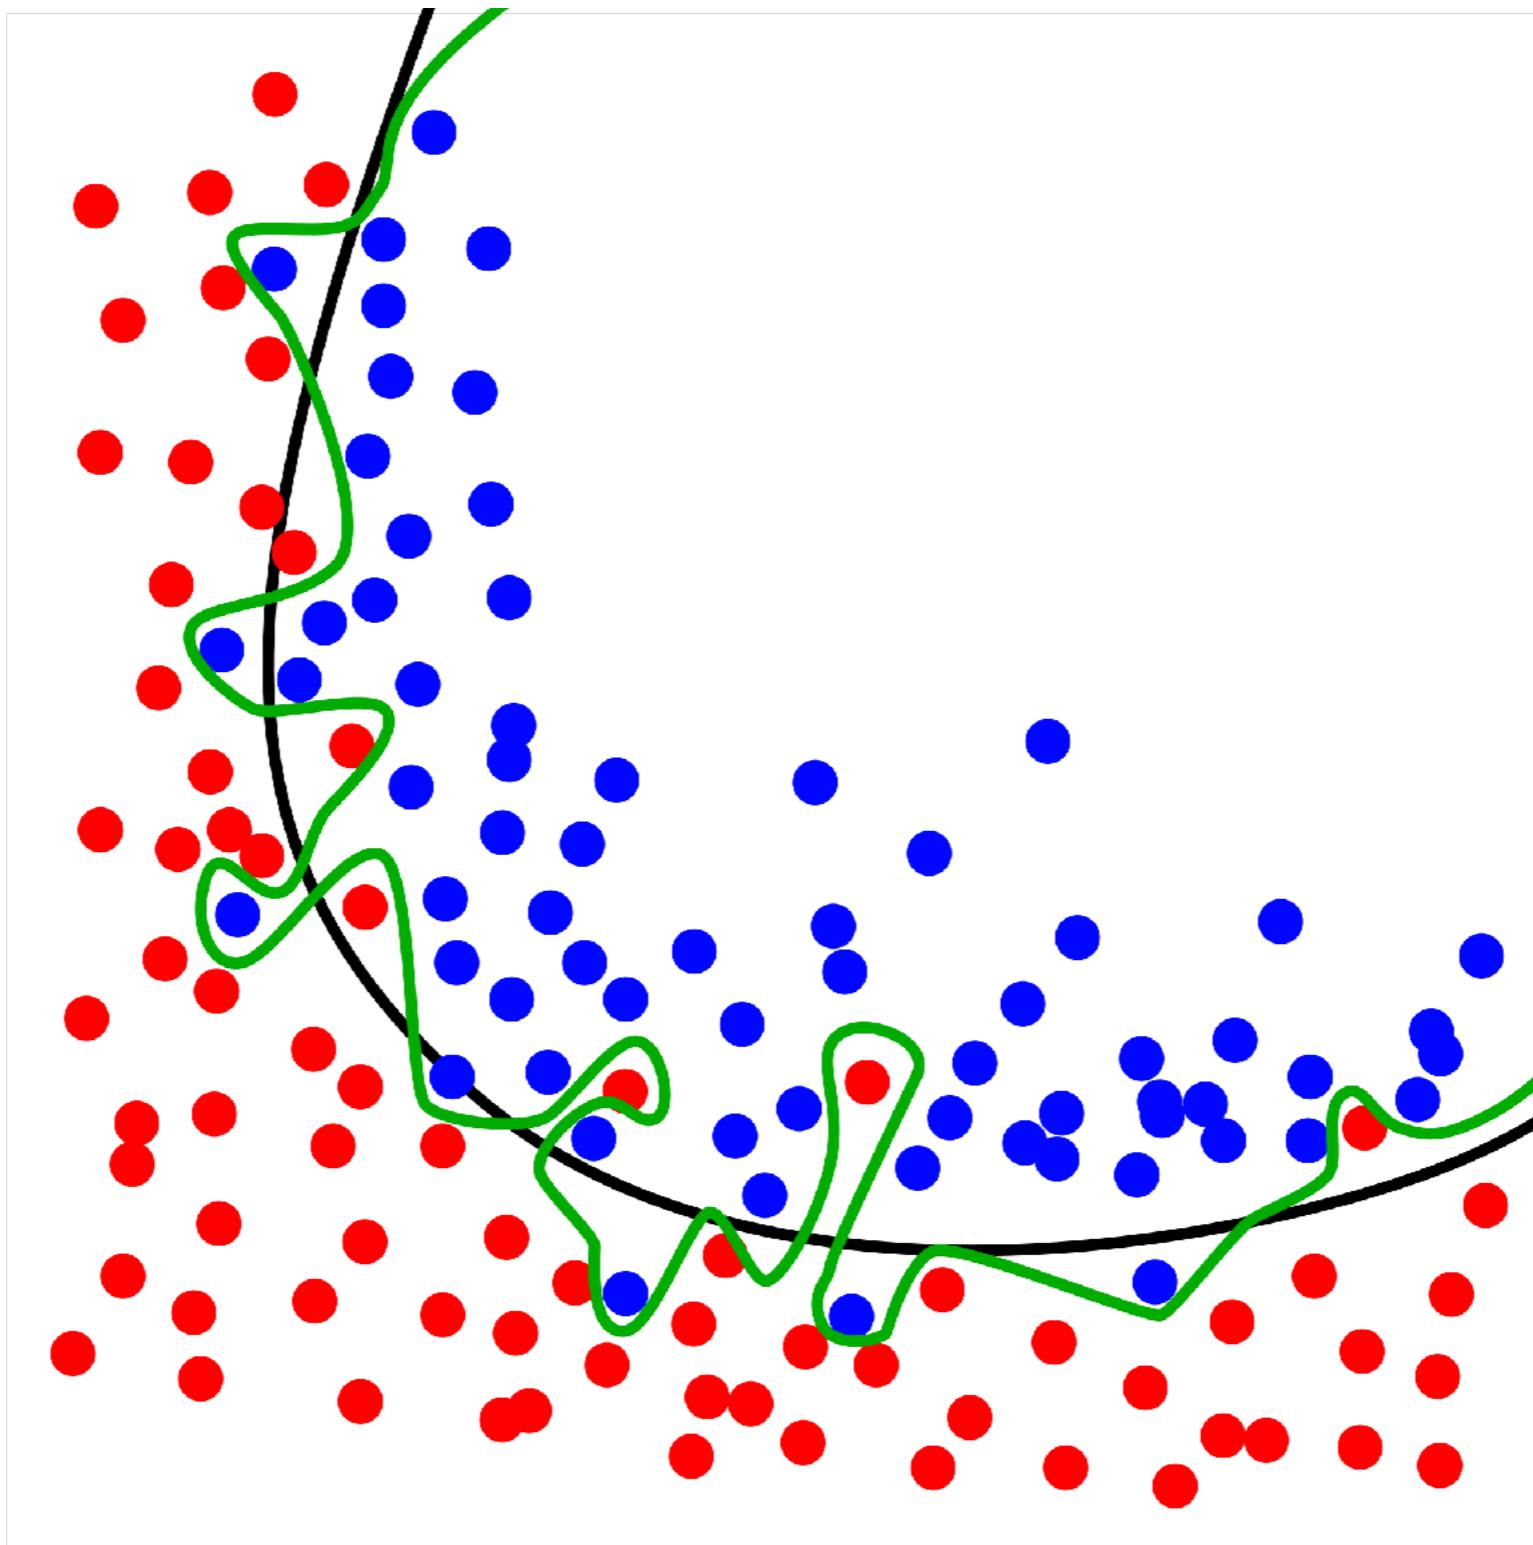
AI



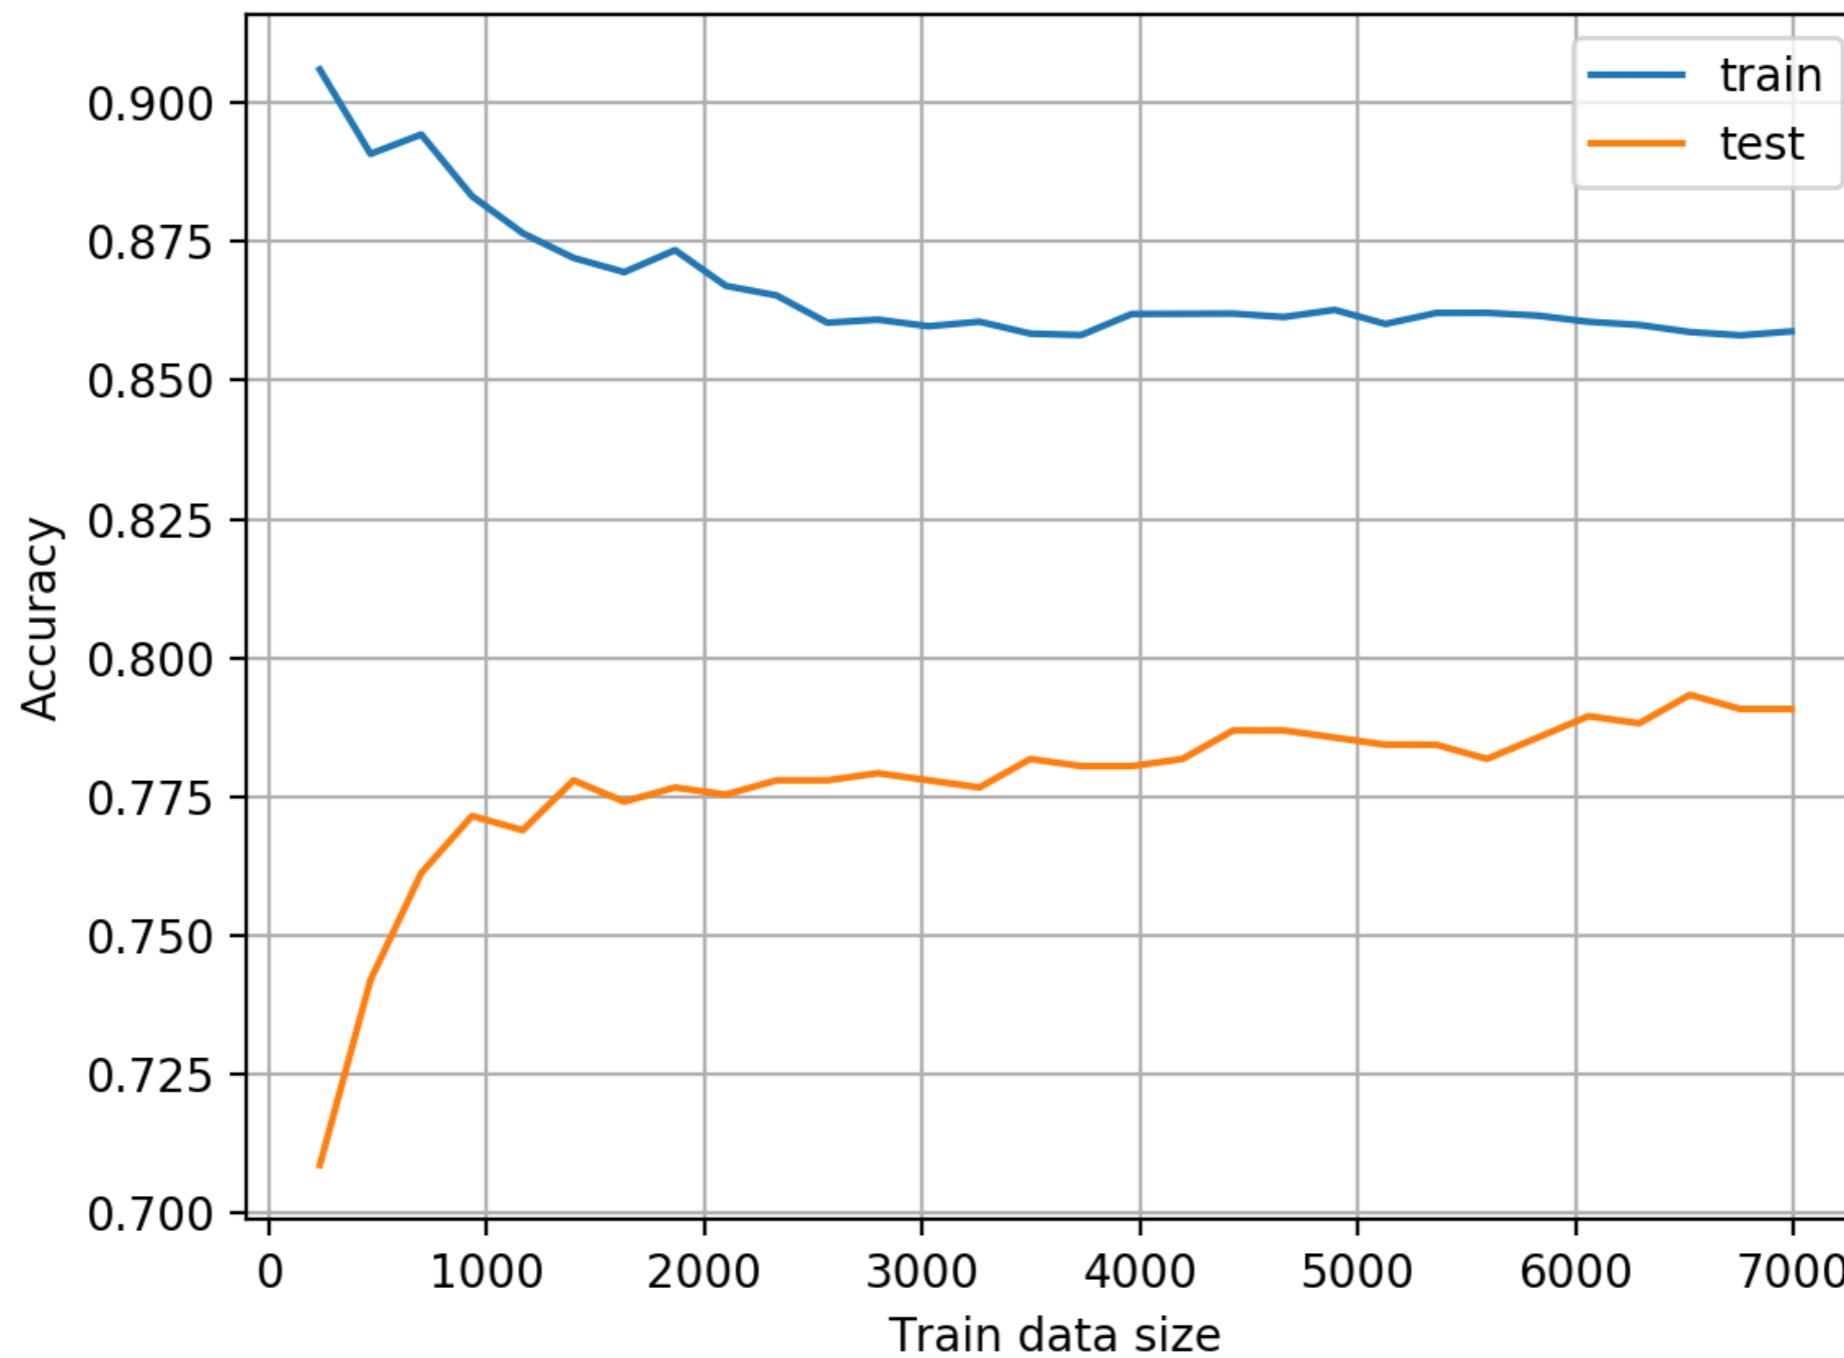
Data

- Train, validation, test data sets
- Cross-validation
- Imbalanced data sets
- Baseline models

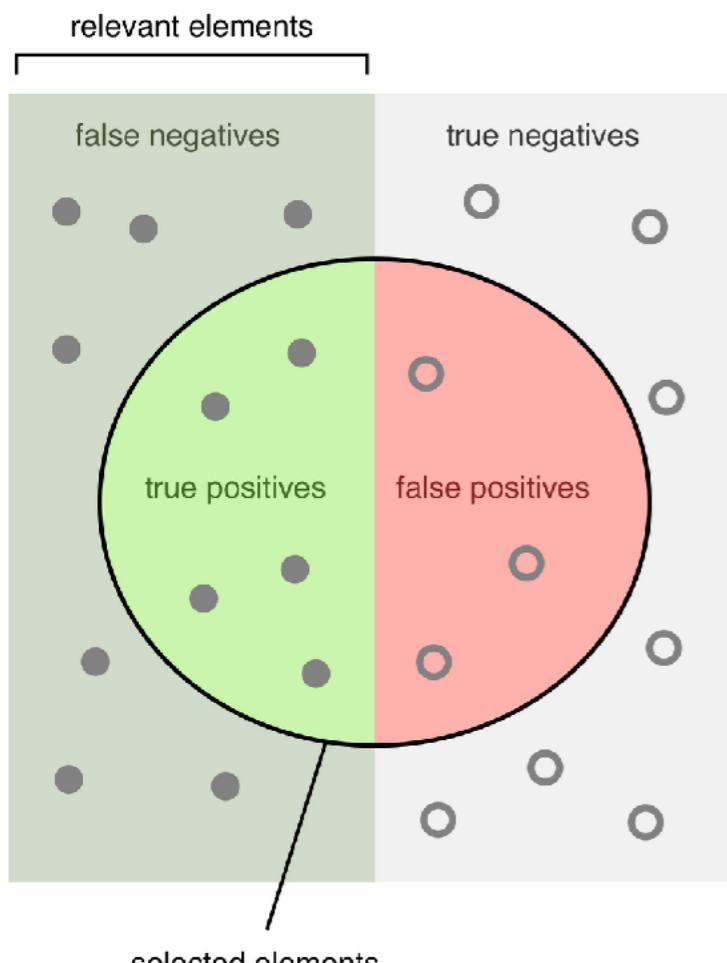
Overfitting



Overfitting detection



Classification Model evaluation



$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fp + fn}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

How many selected items are relevant?

$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{grey}}$$

Conditional probability

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(A)P(B|A)}{P(B)}$$

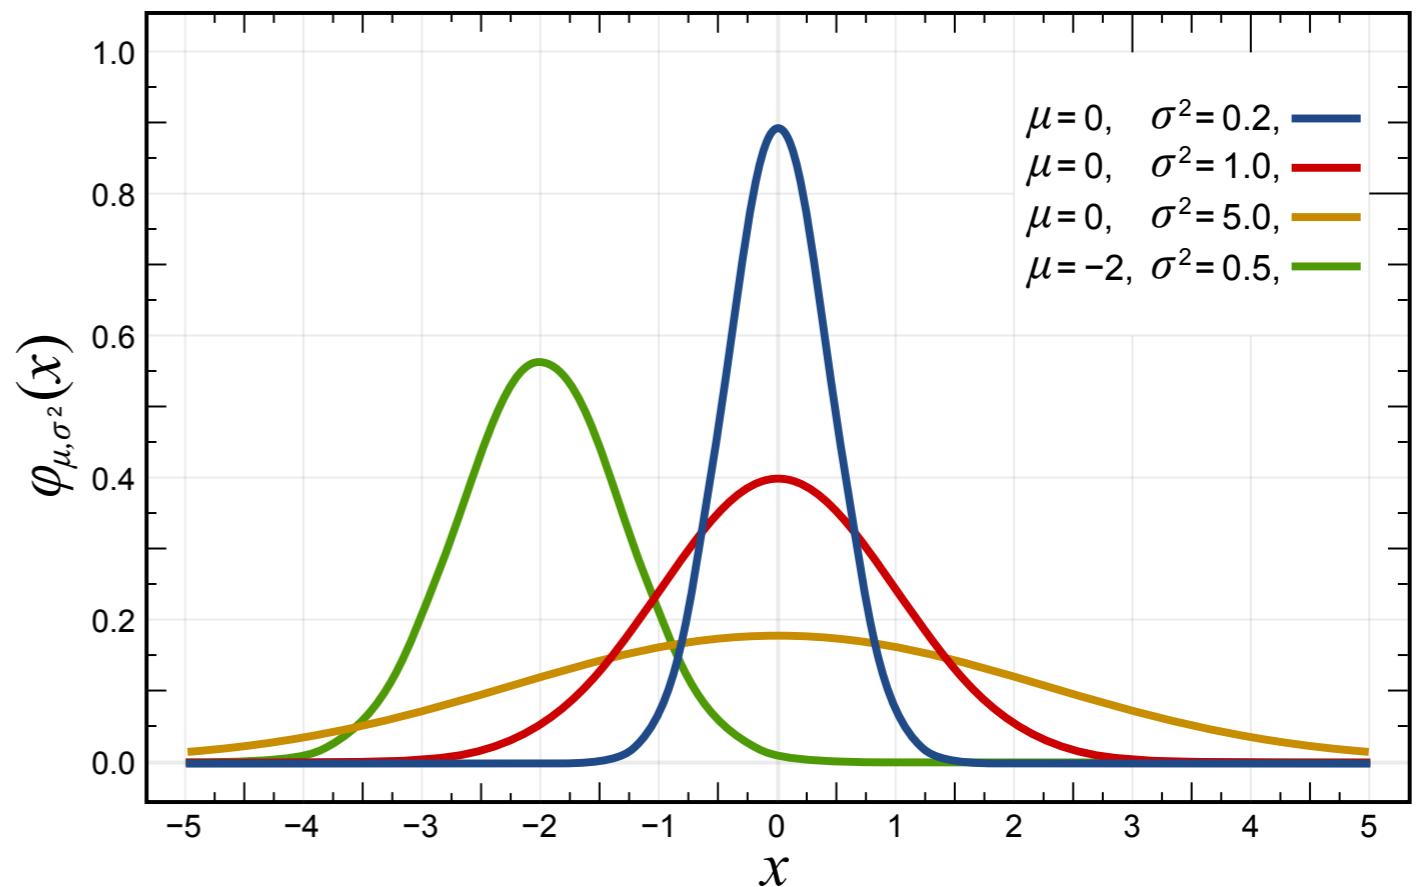
$$A \perp B \iff P(A \cap B) = P(A)P(B)$$

Naive Bayes Classifier

client	balance	income	sex	unemployed	loan
1	H	H	F	F	T
2	L	H	M	F	T
3	L	L	M	T	F
4	H	L	F	T	T
5	L	L	F	T	F
6	H	L	M	F	?

Gaussian Naive Bayes Classifier

$$p(x = v|C_k) = \frac{1}{\sqrt{2\pi\sigma_k^2}} e^{-\frac{(v - \mu_k)^2}{2\sigma_k^2}}$$



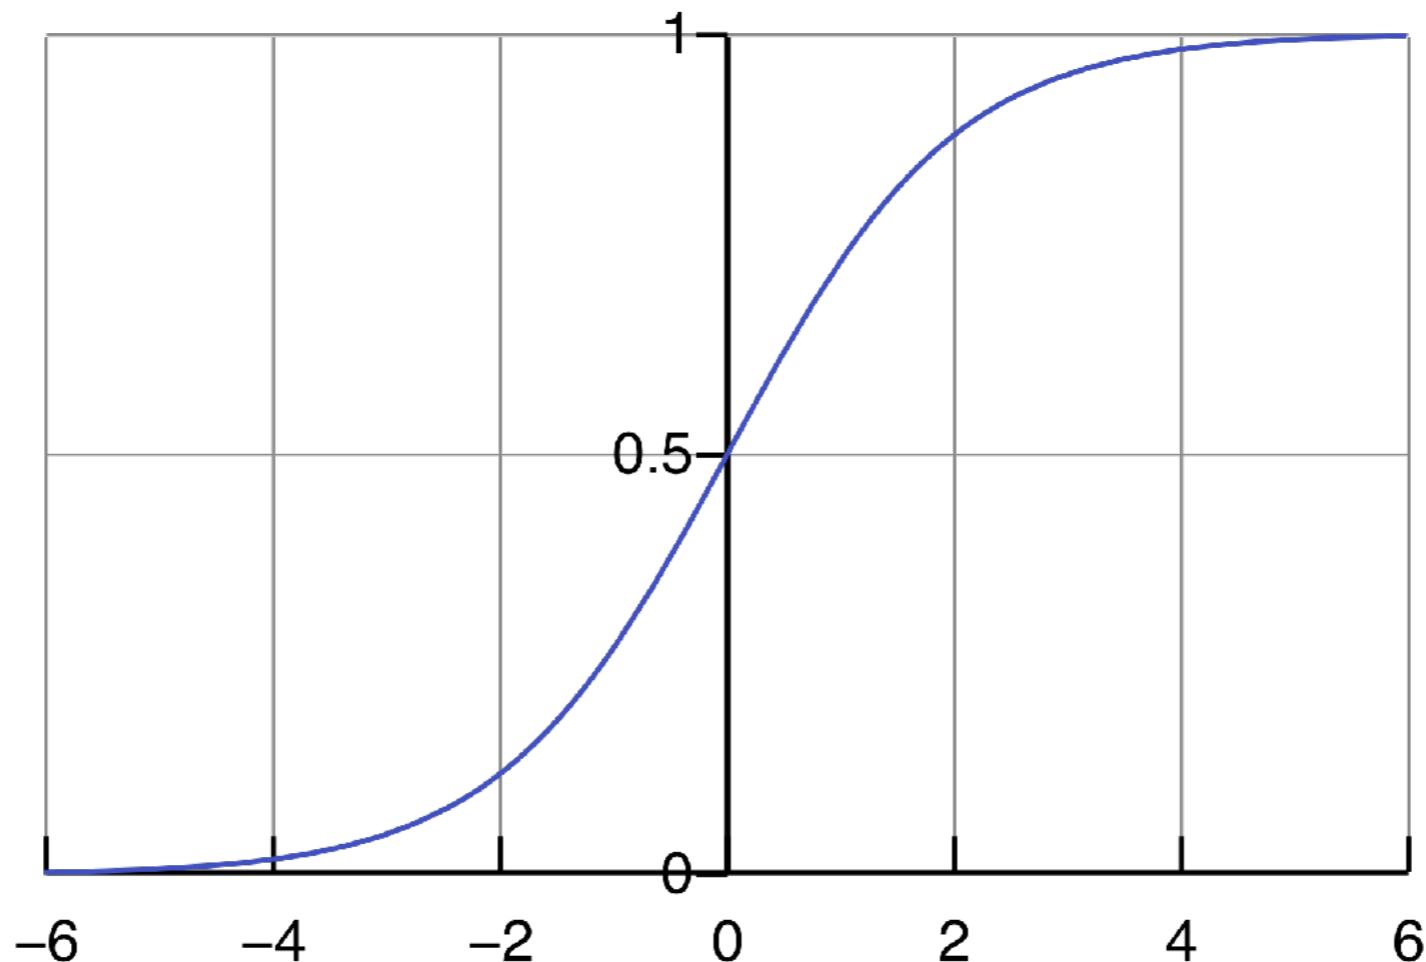
Scikit-learn tutorial

<http://scikit-learn.org/stable/>

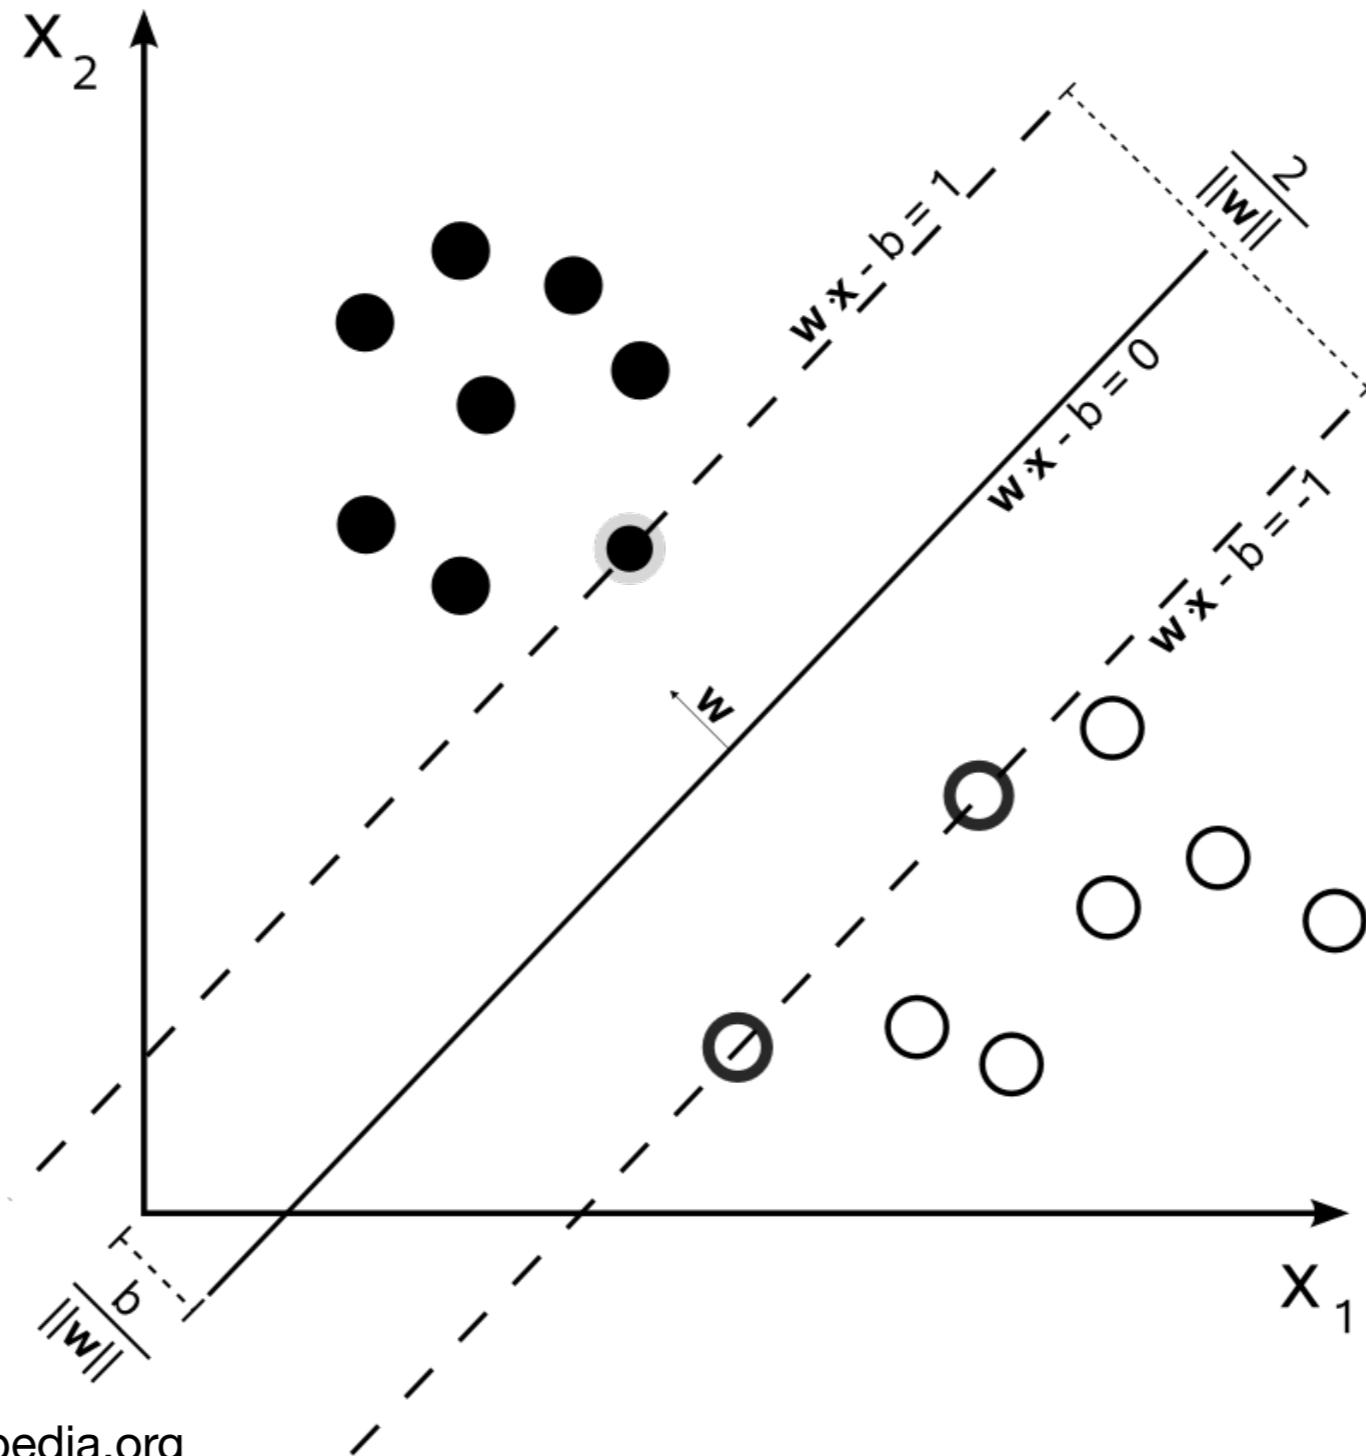
[01-Scikit-introduction.ipynb](#)

Logistic regression

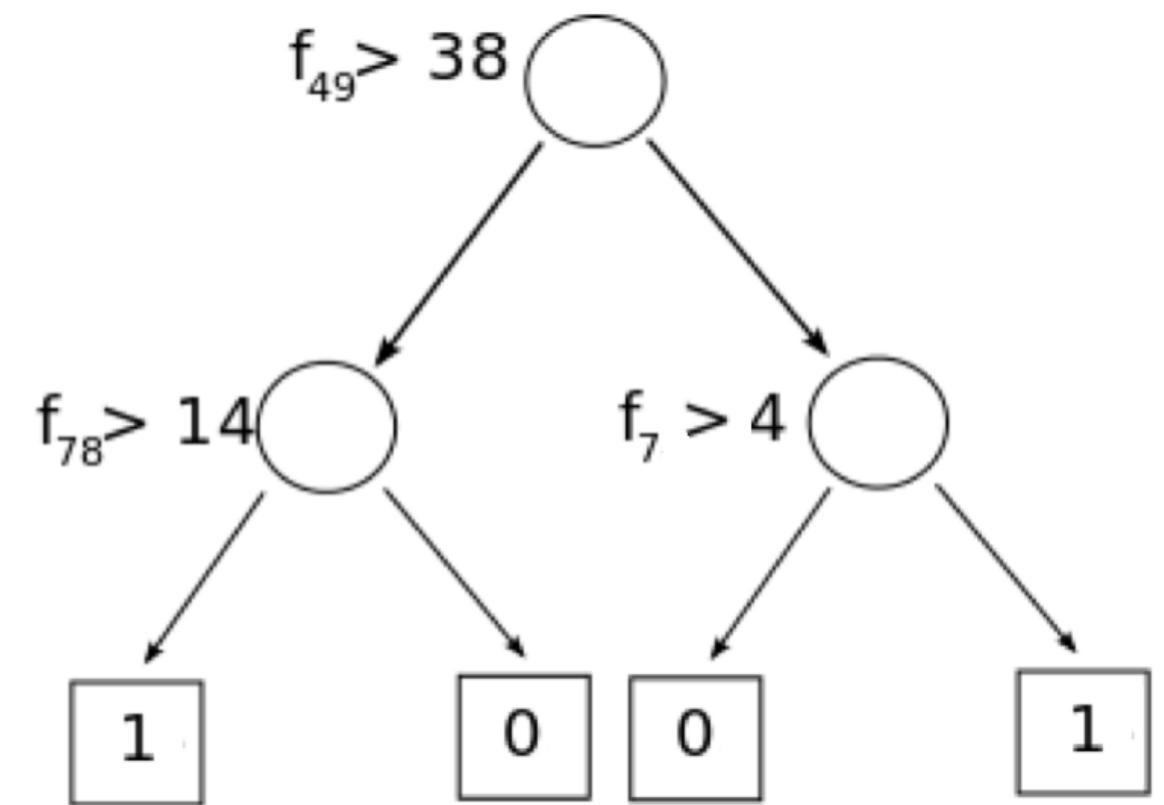
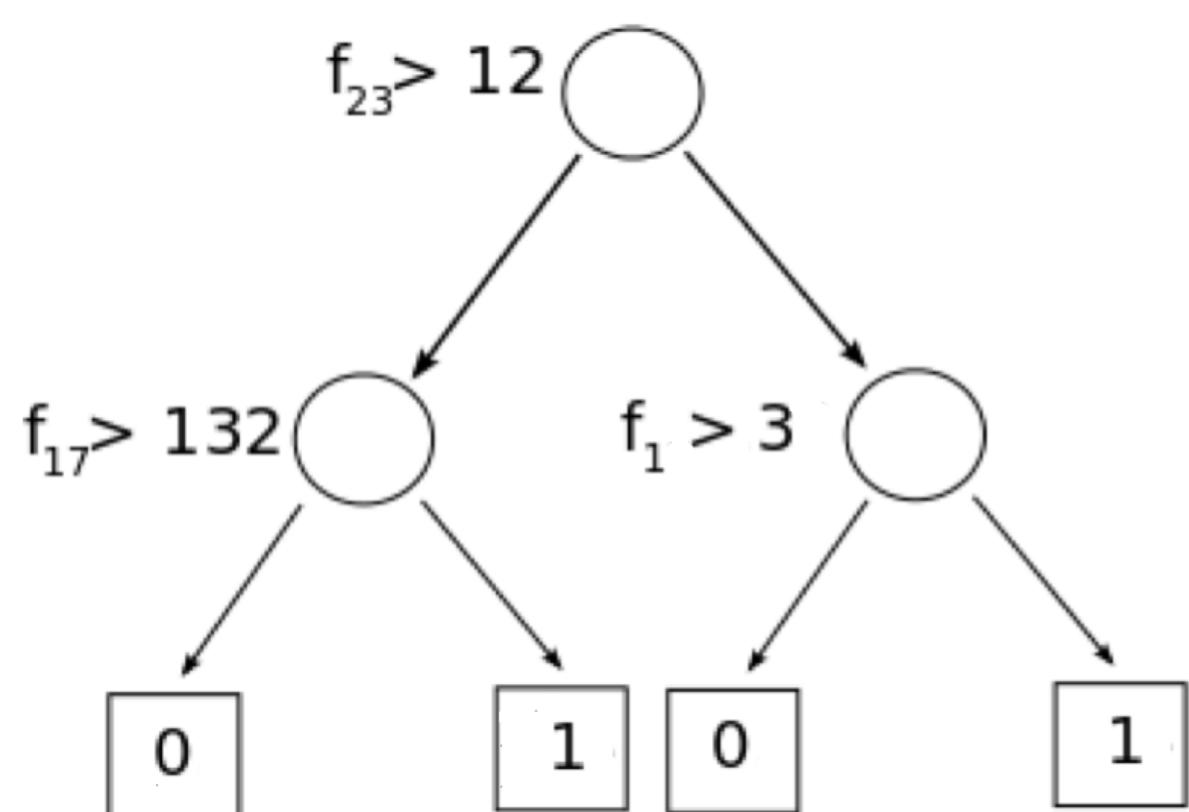
$$P(y|\vec{x}) = \frac{1}{1 + e^{-(\vec{x}\vec{w} + w_0)}}$$



Support Vector Machines



Boosted Decision Trees

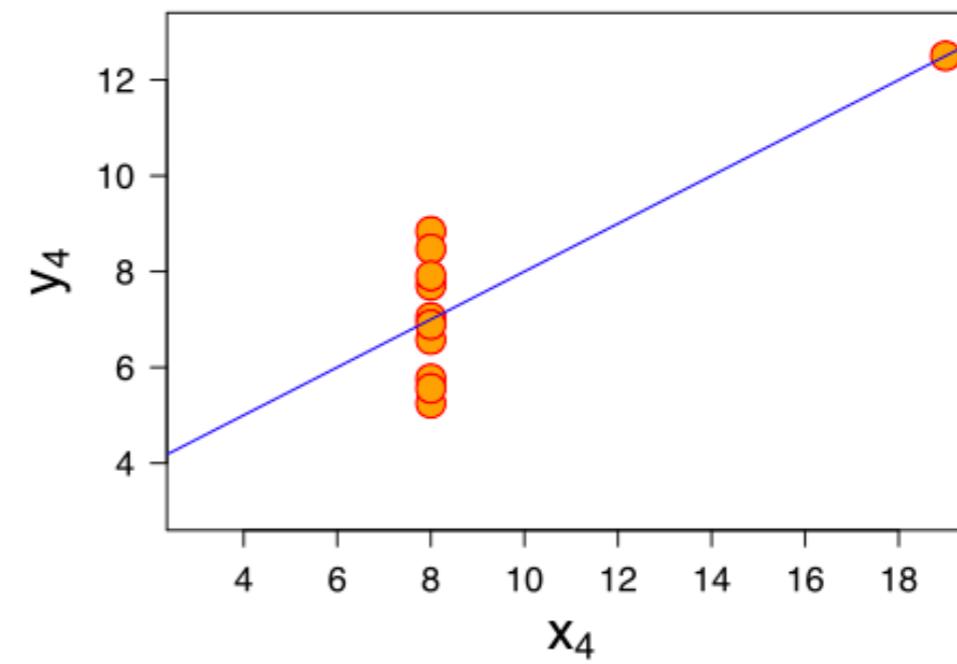
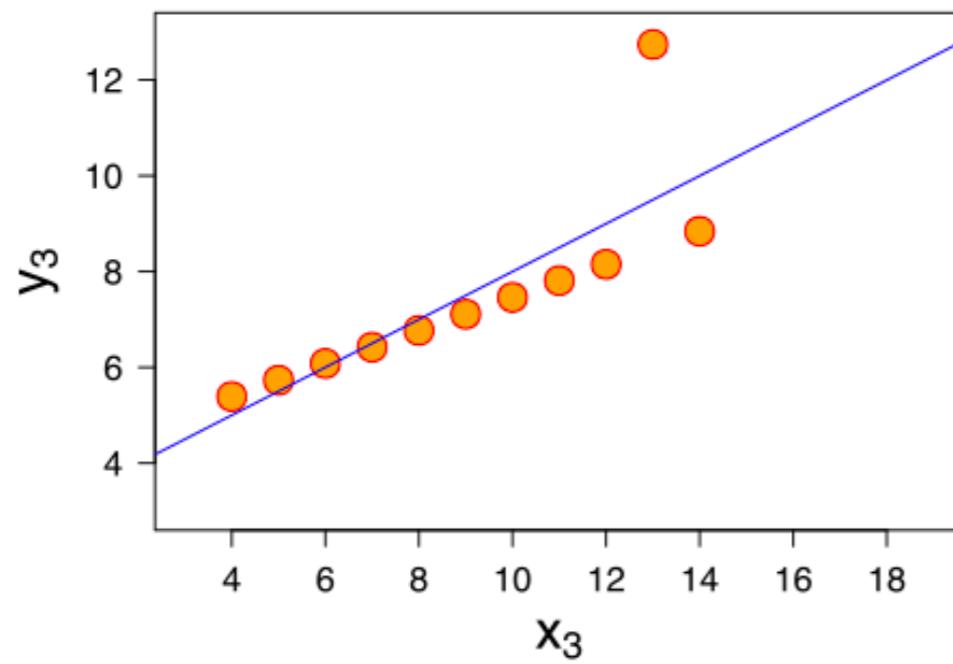
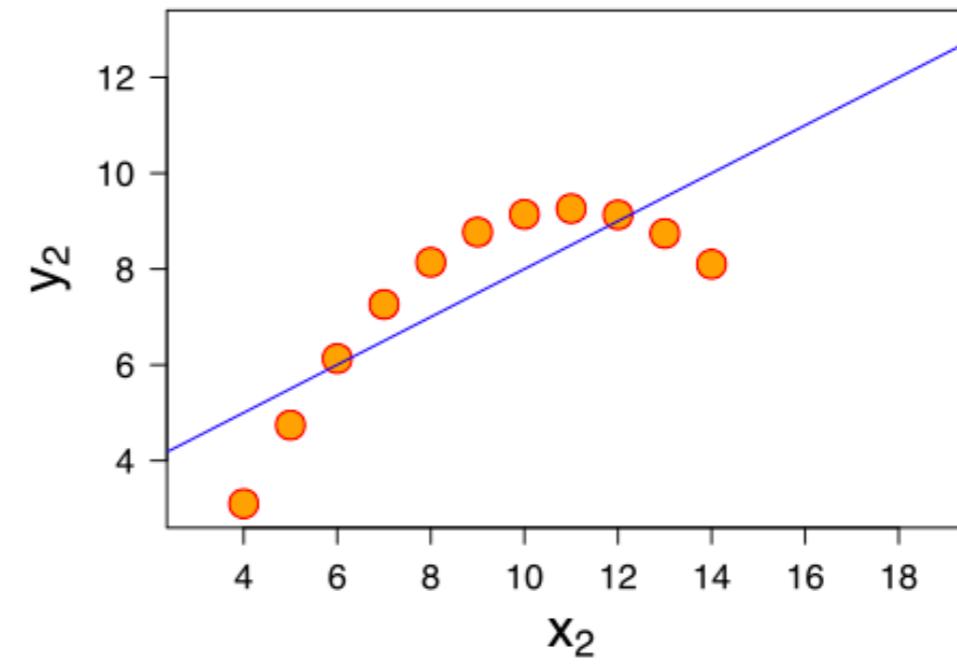
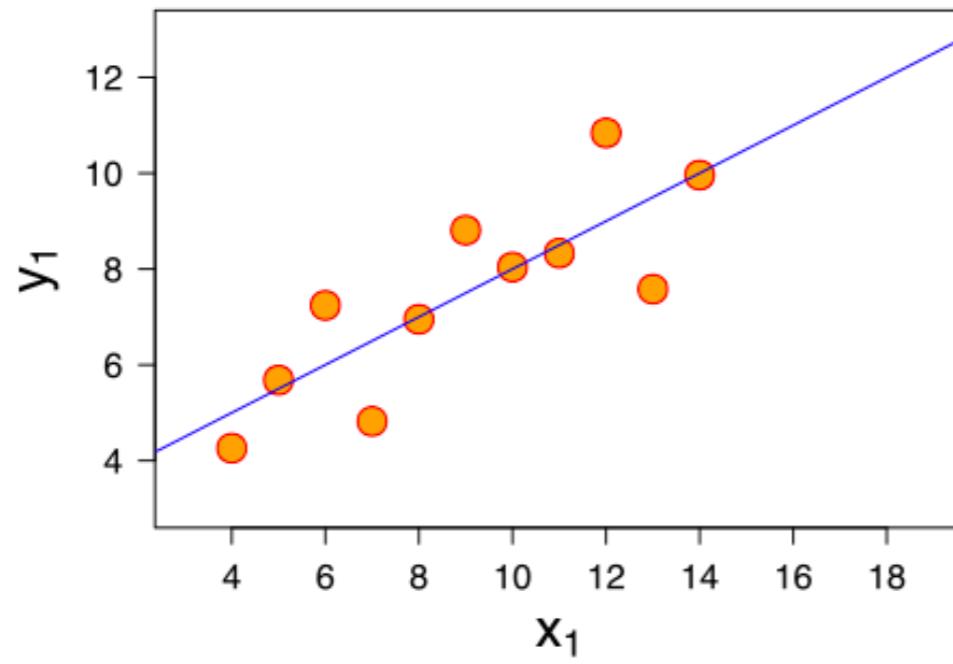


Classification task

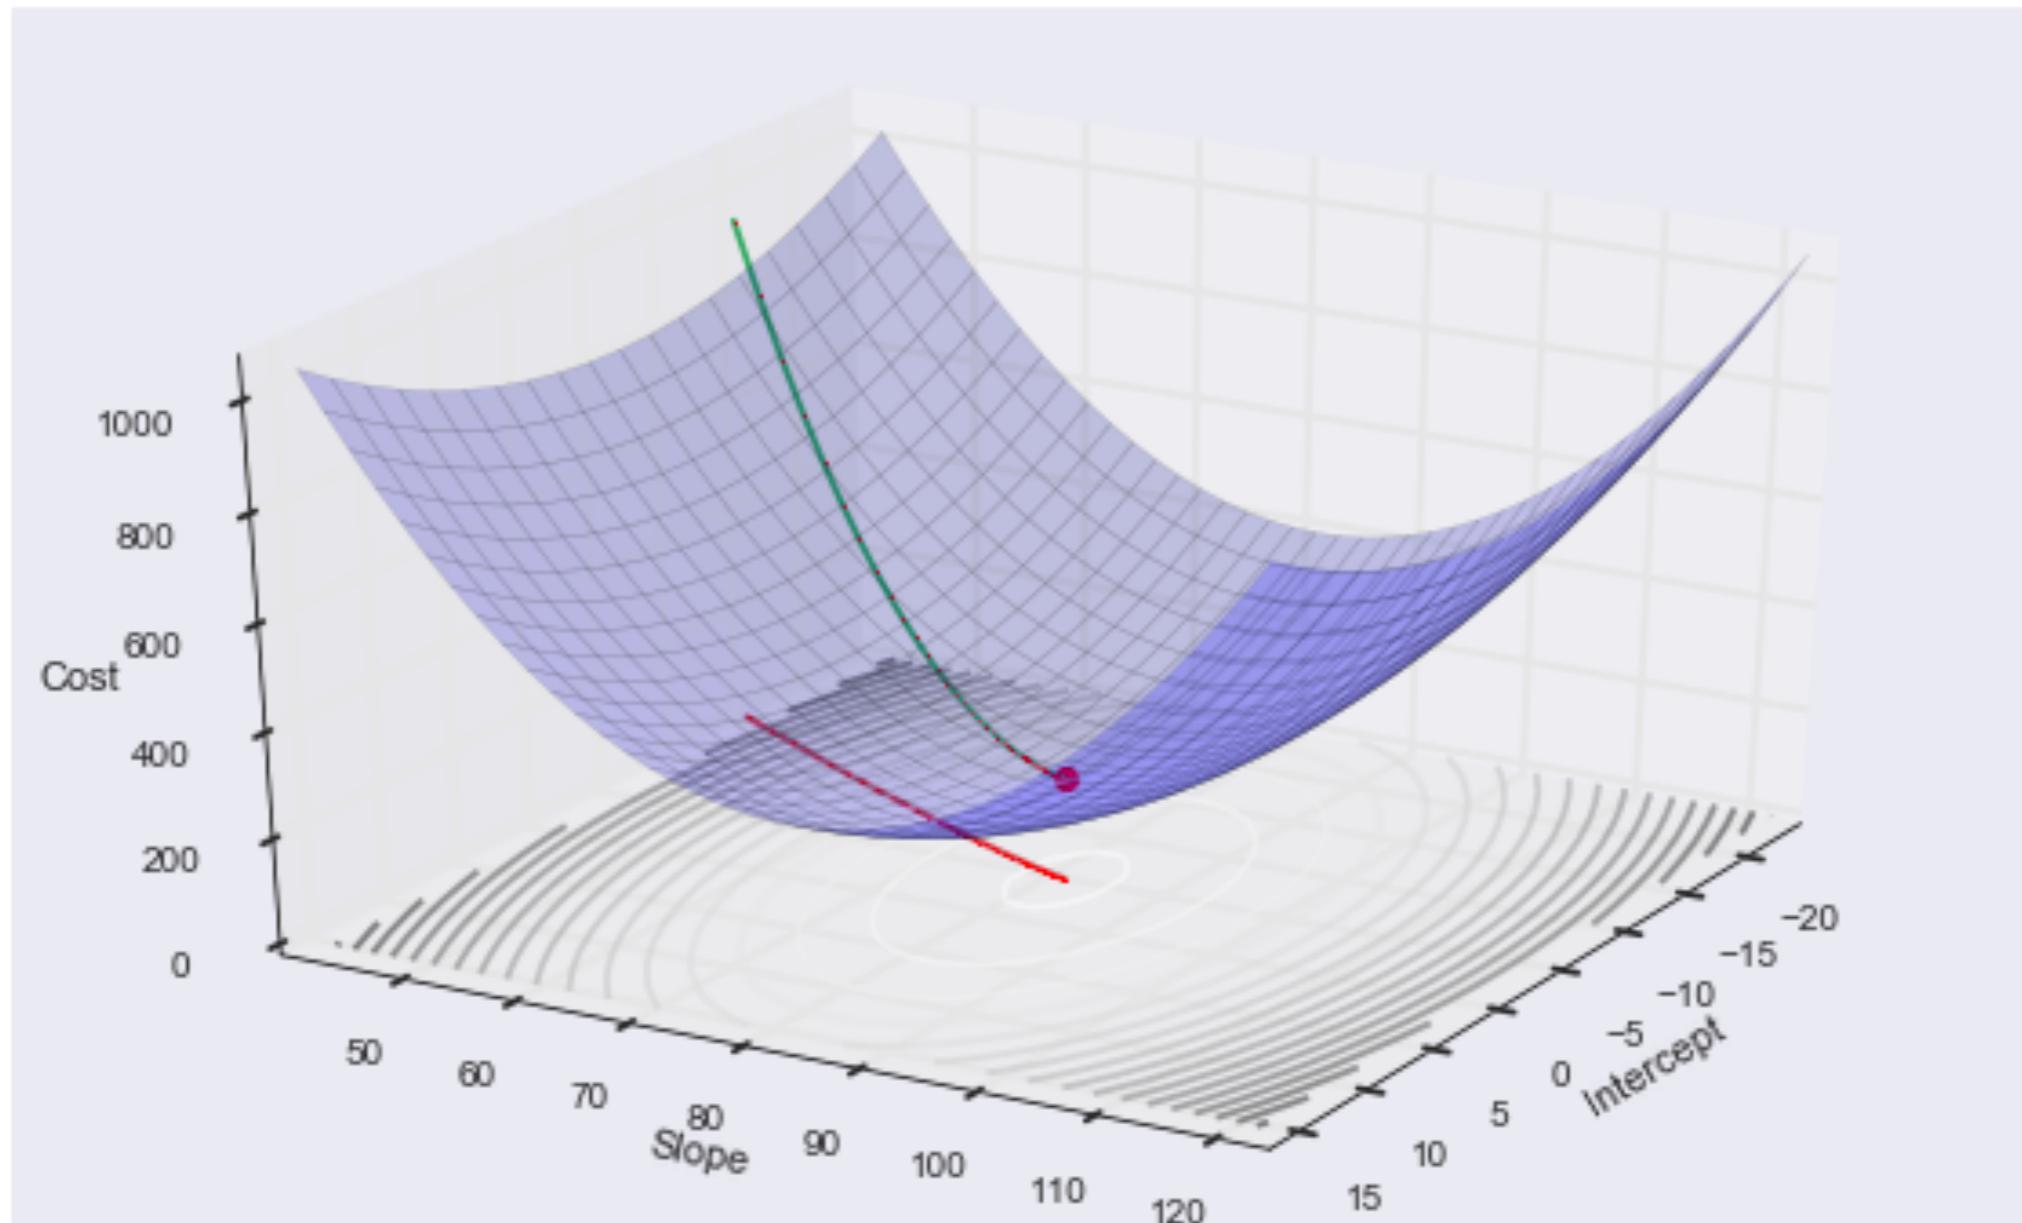
02-Classification1-assignment.ipynb

03-Classification2-assignment.ipynb

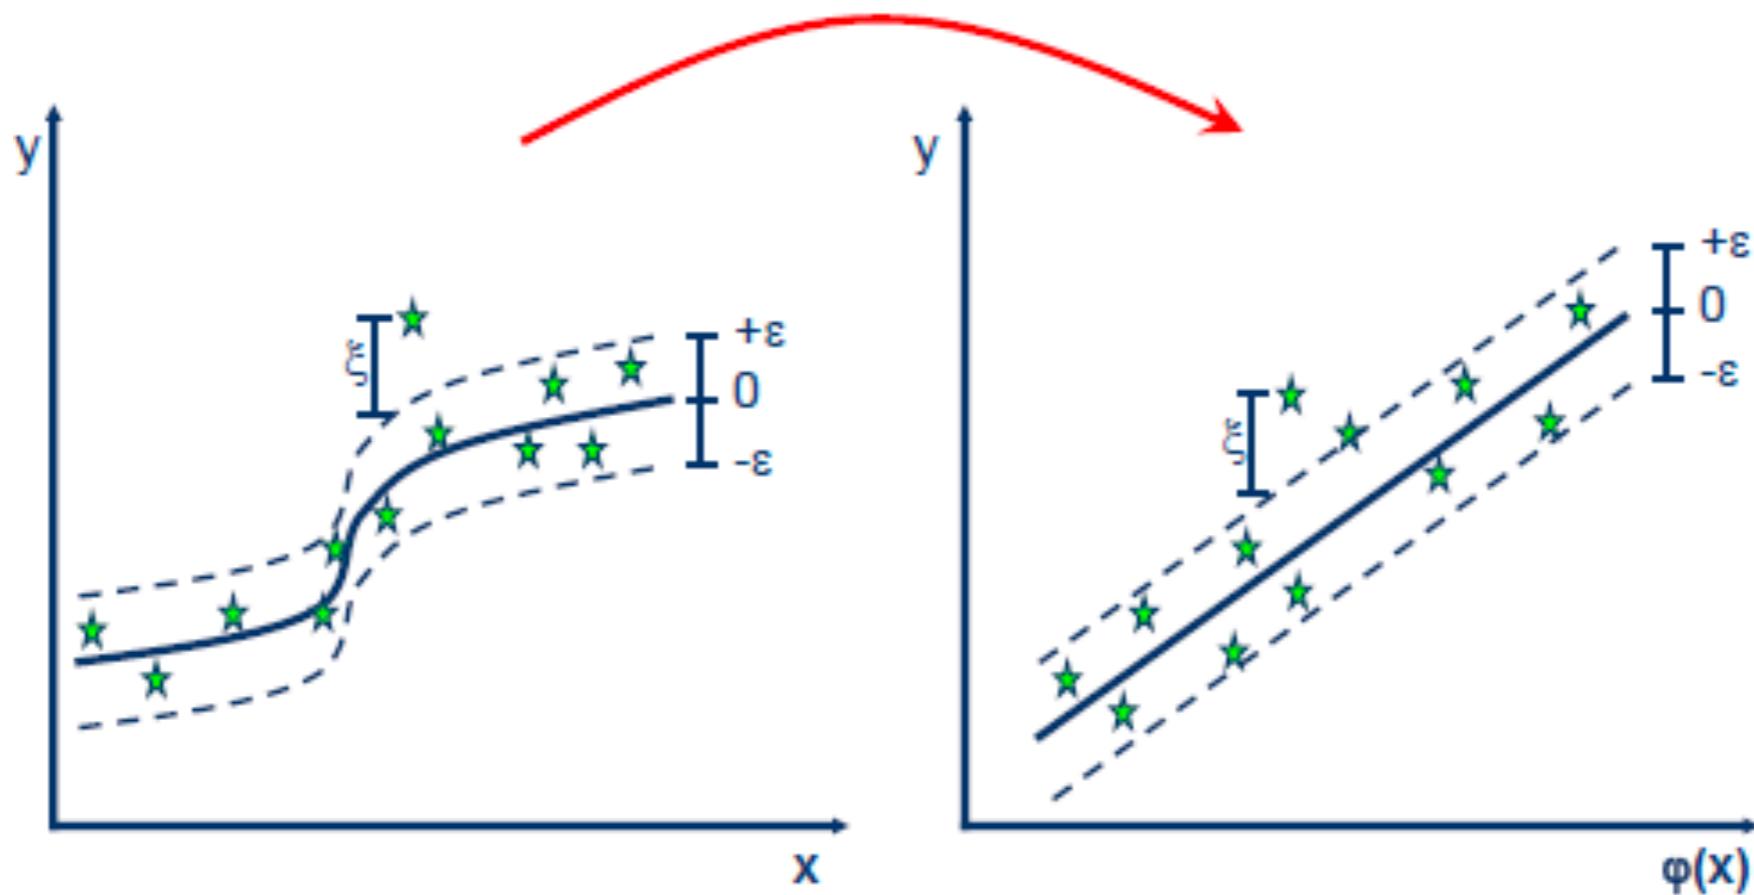
Regression



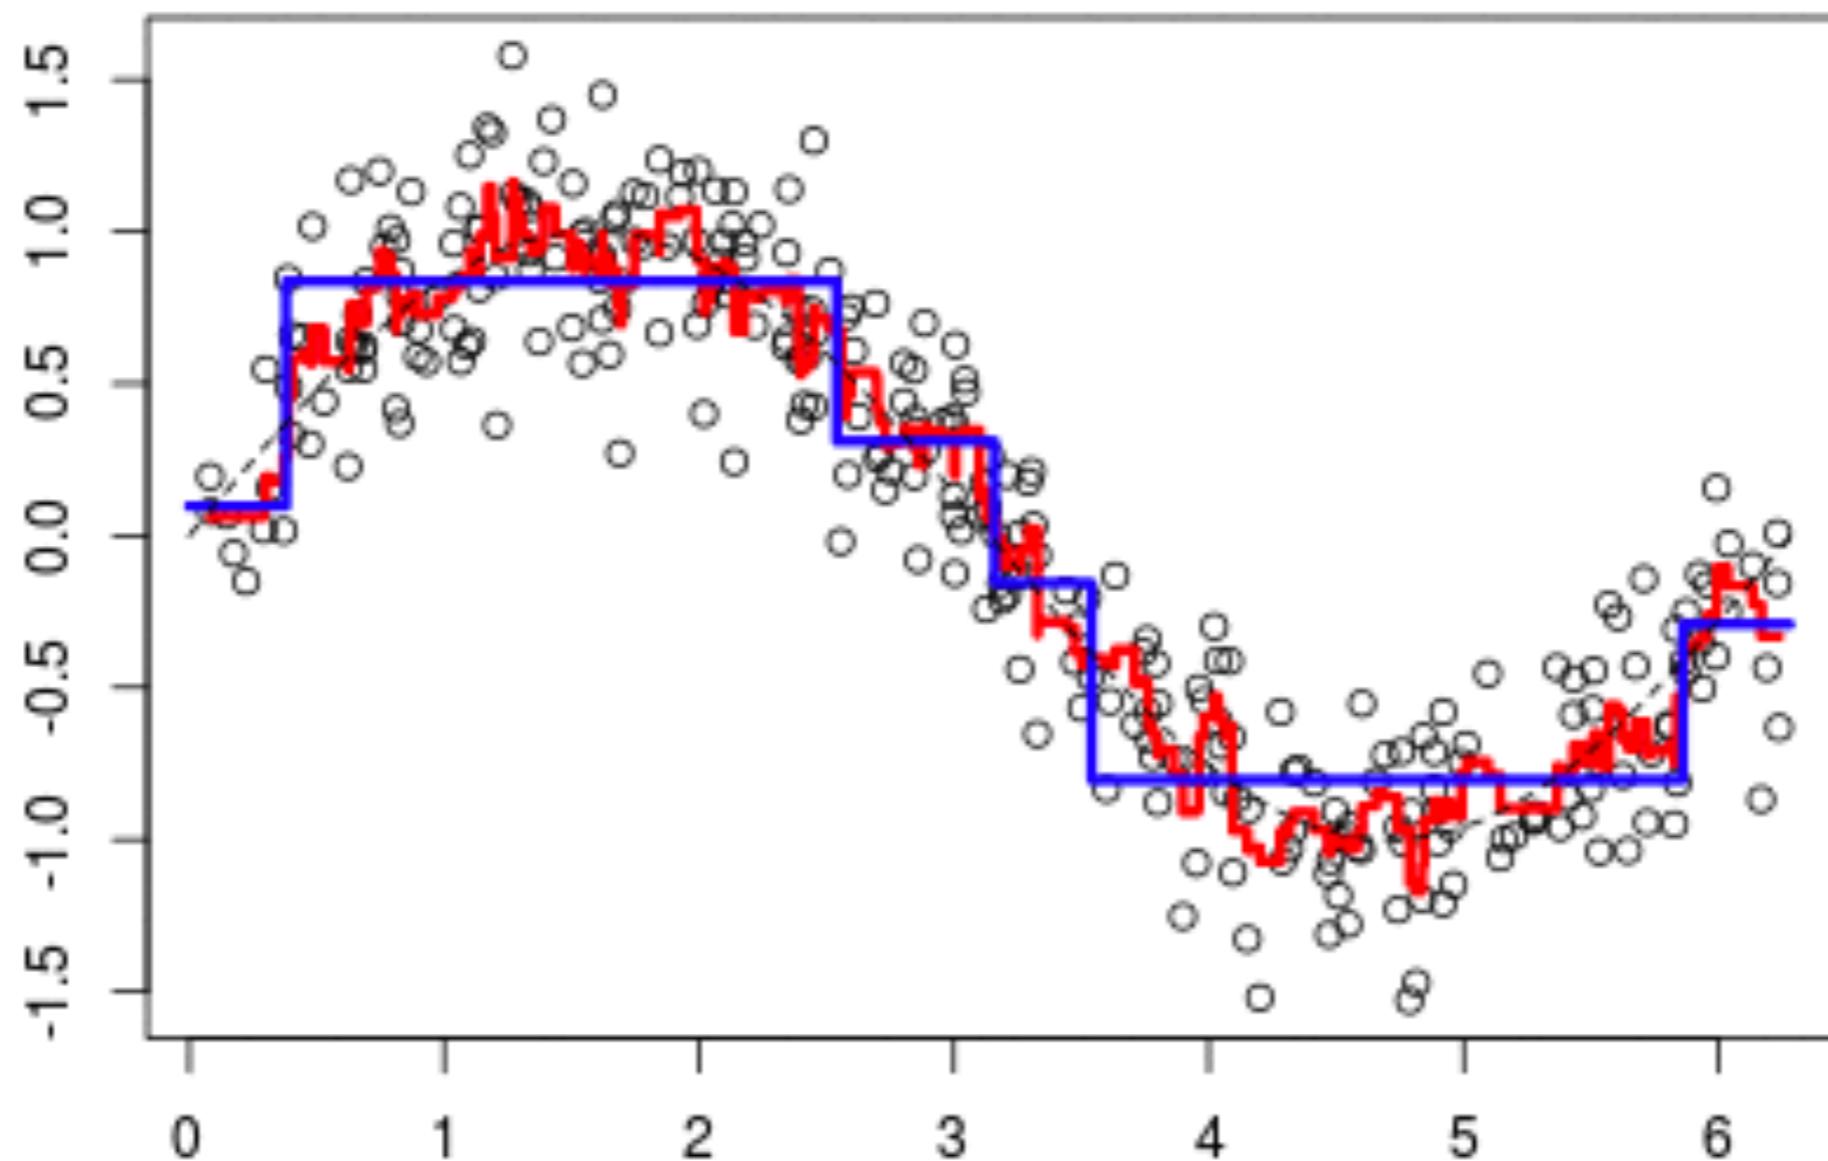
Linear regression with SGD



Support vector machines for Regression



Boosted regression trees



Evaluation of regression models

Root mean squared error

$$\text{RMSE} = \sqrt{\frac{\sum_i (y_i - \hat{y}_i)^2}{n}}$$

Mean absolute error

$$\text{MAE} = \frac{\sum_i |y_i - \hat{y}_i|}{n}$$

R Squared

$$\hat{R}^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Regression task

[04-Regression1-assignment.ipynb](#)

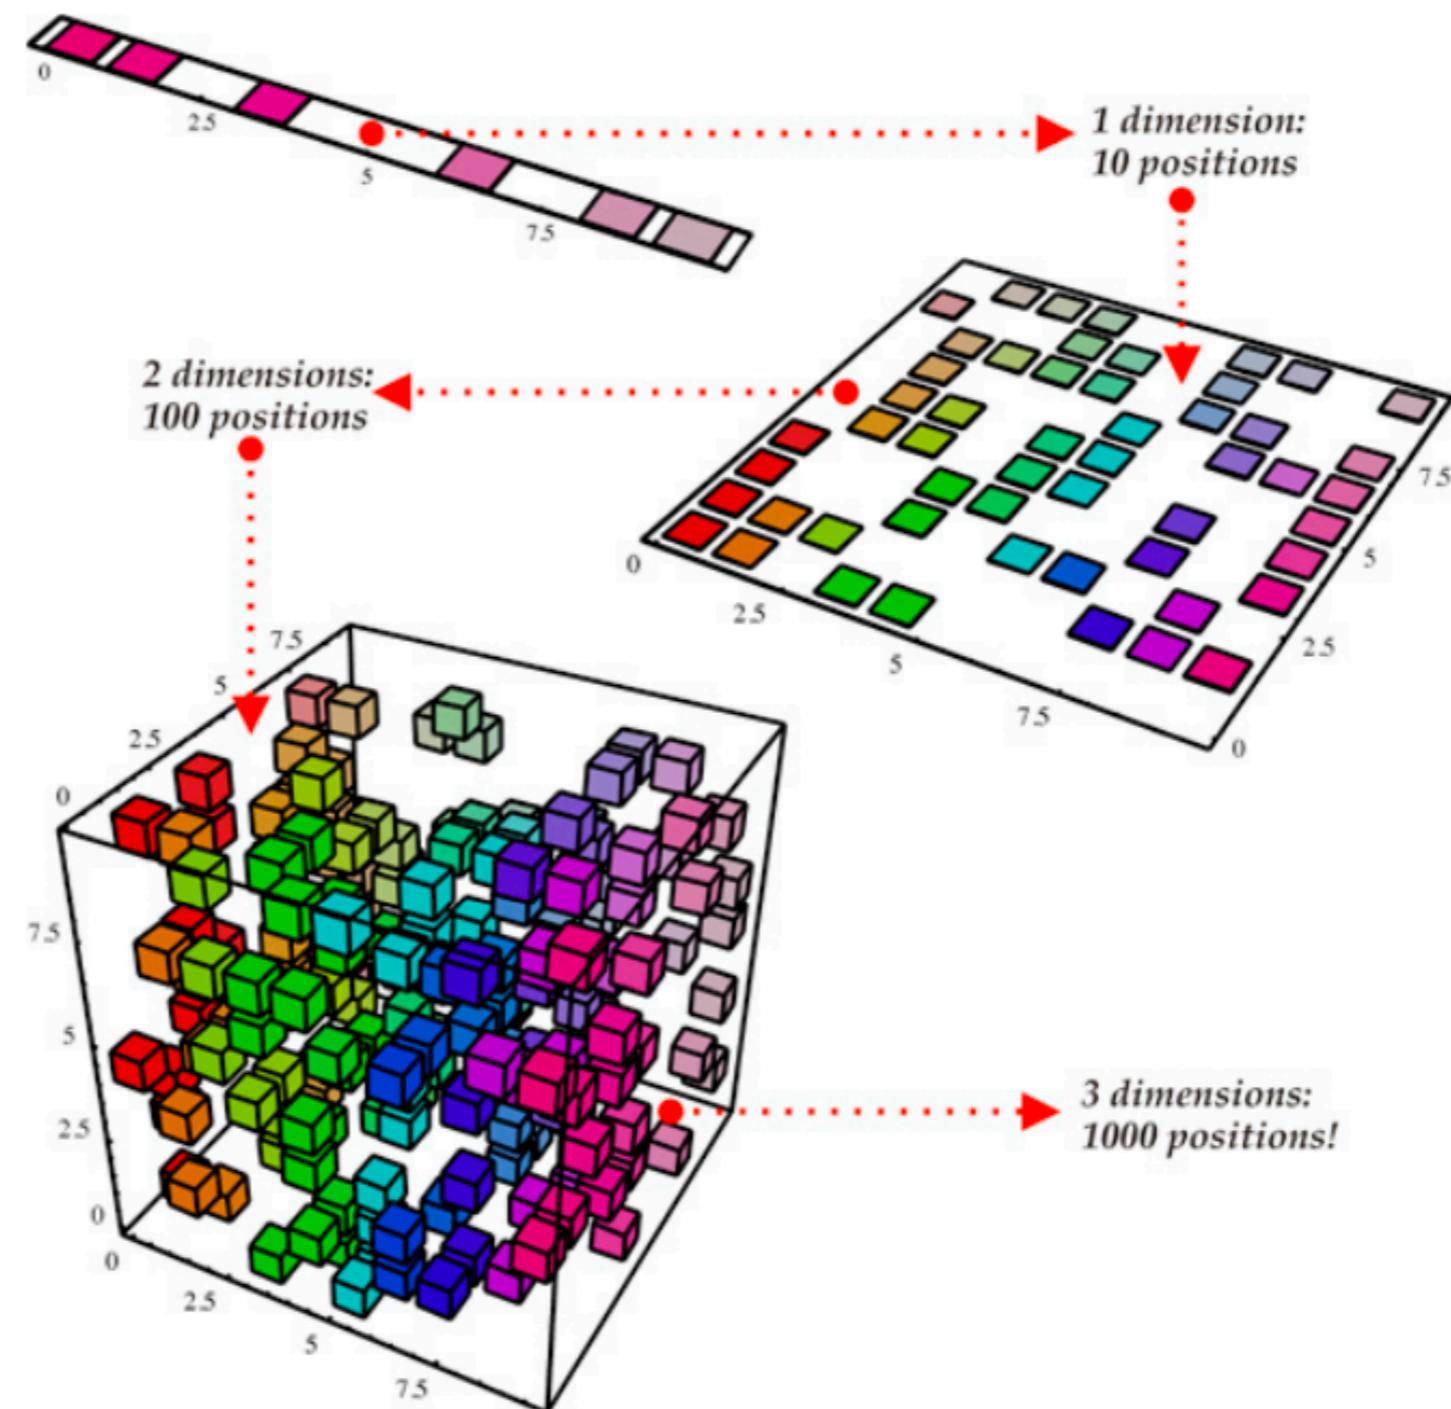
[05-Regression2-assignment.ipynb](#)

Unsupervised techniques



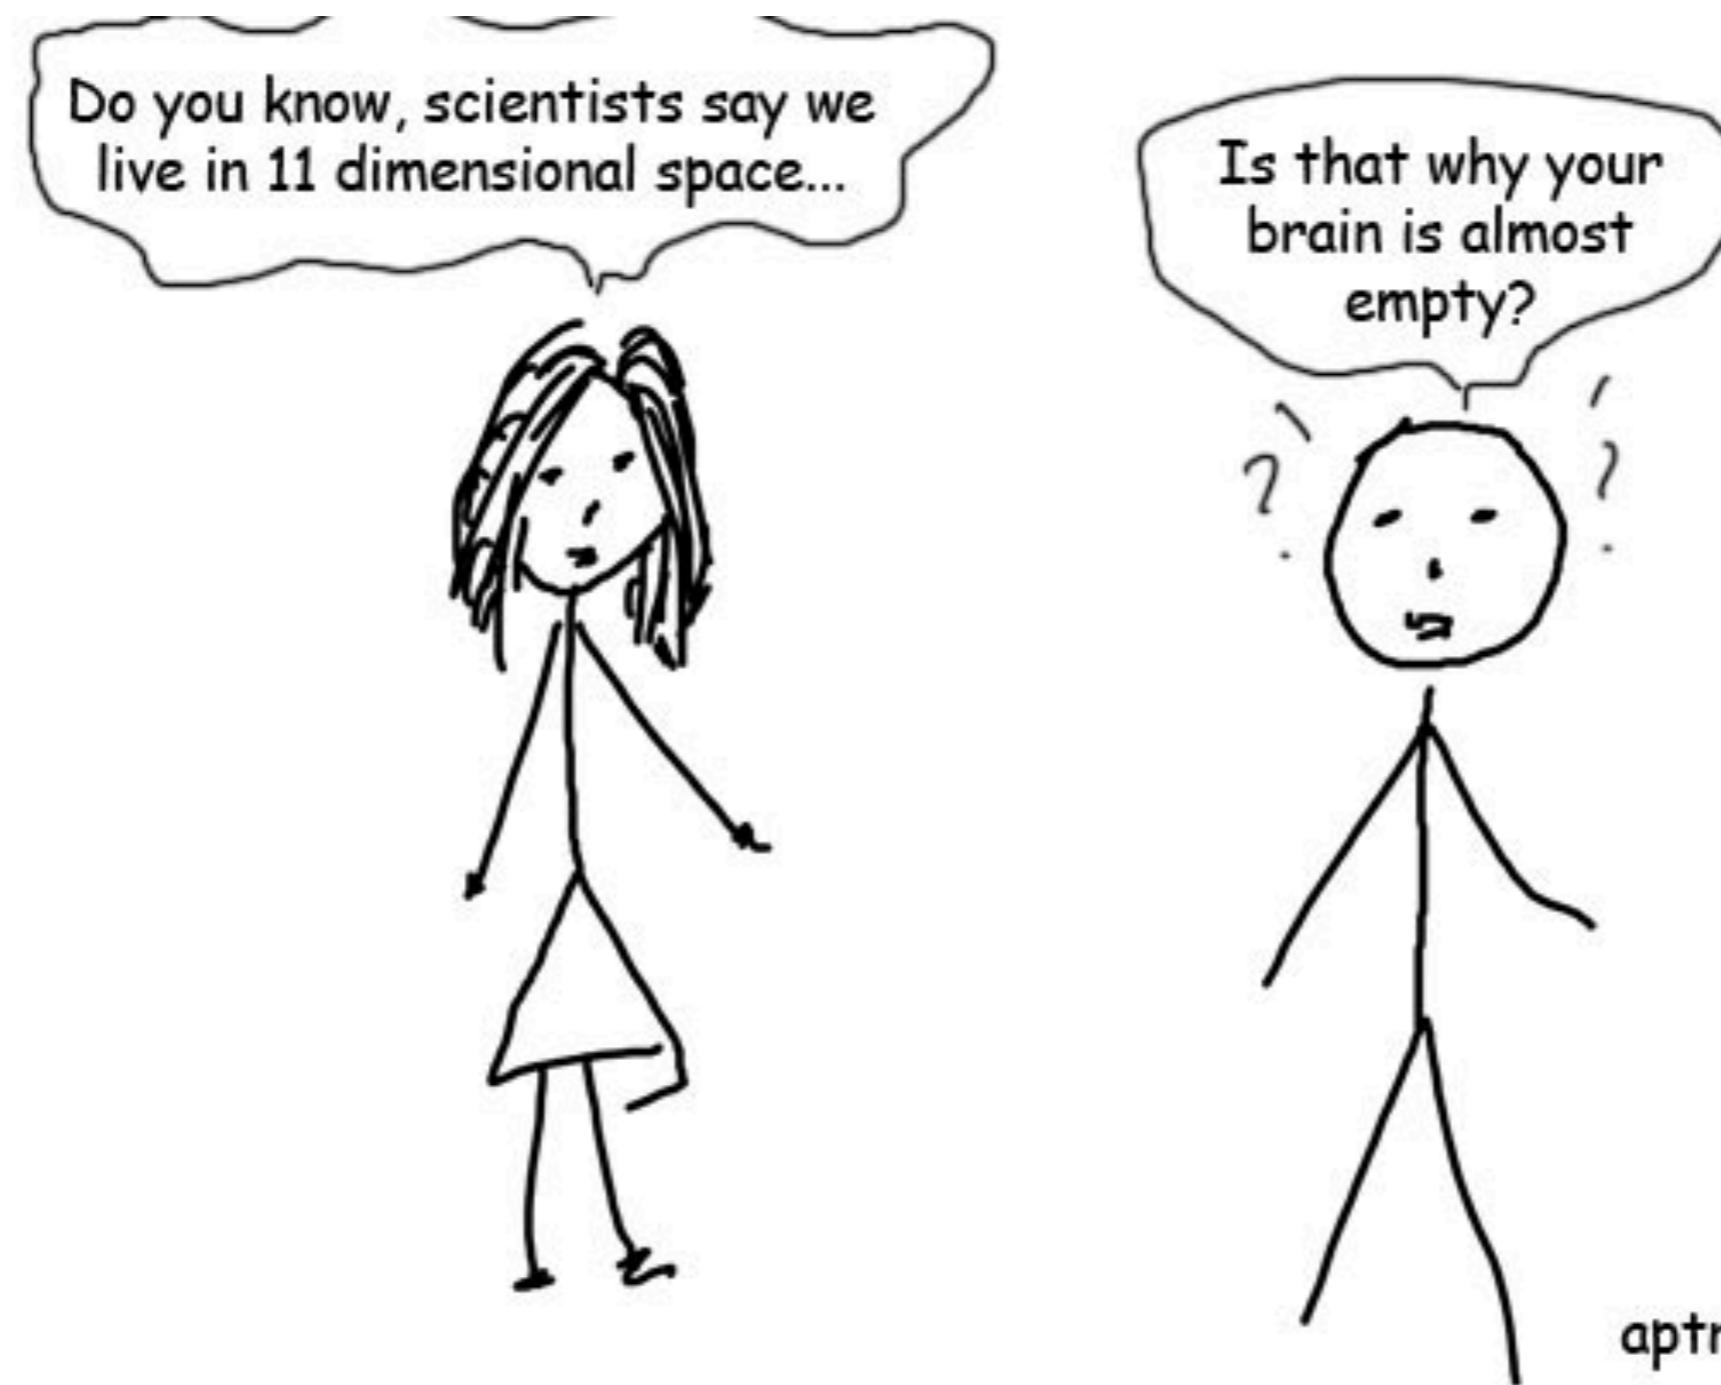
Source: tmdb.org

Dimensionality reduction

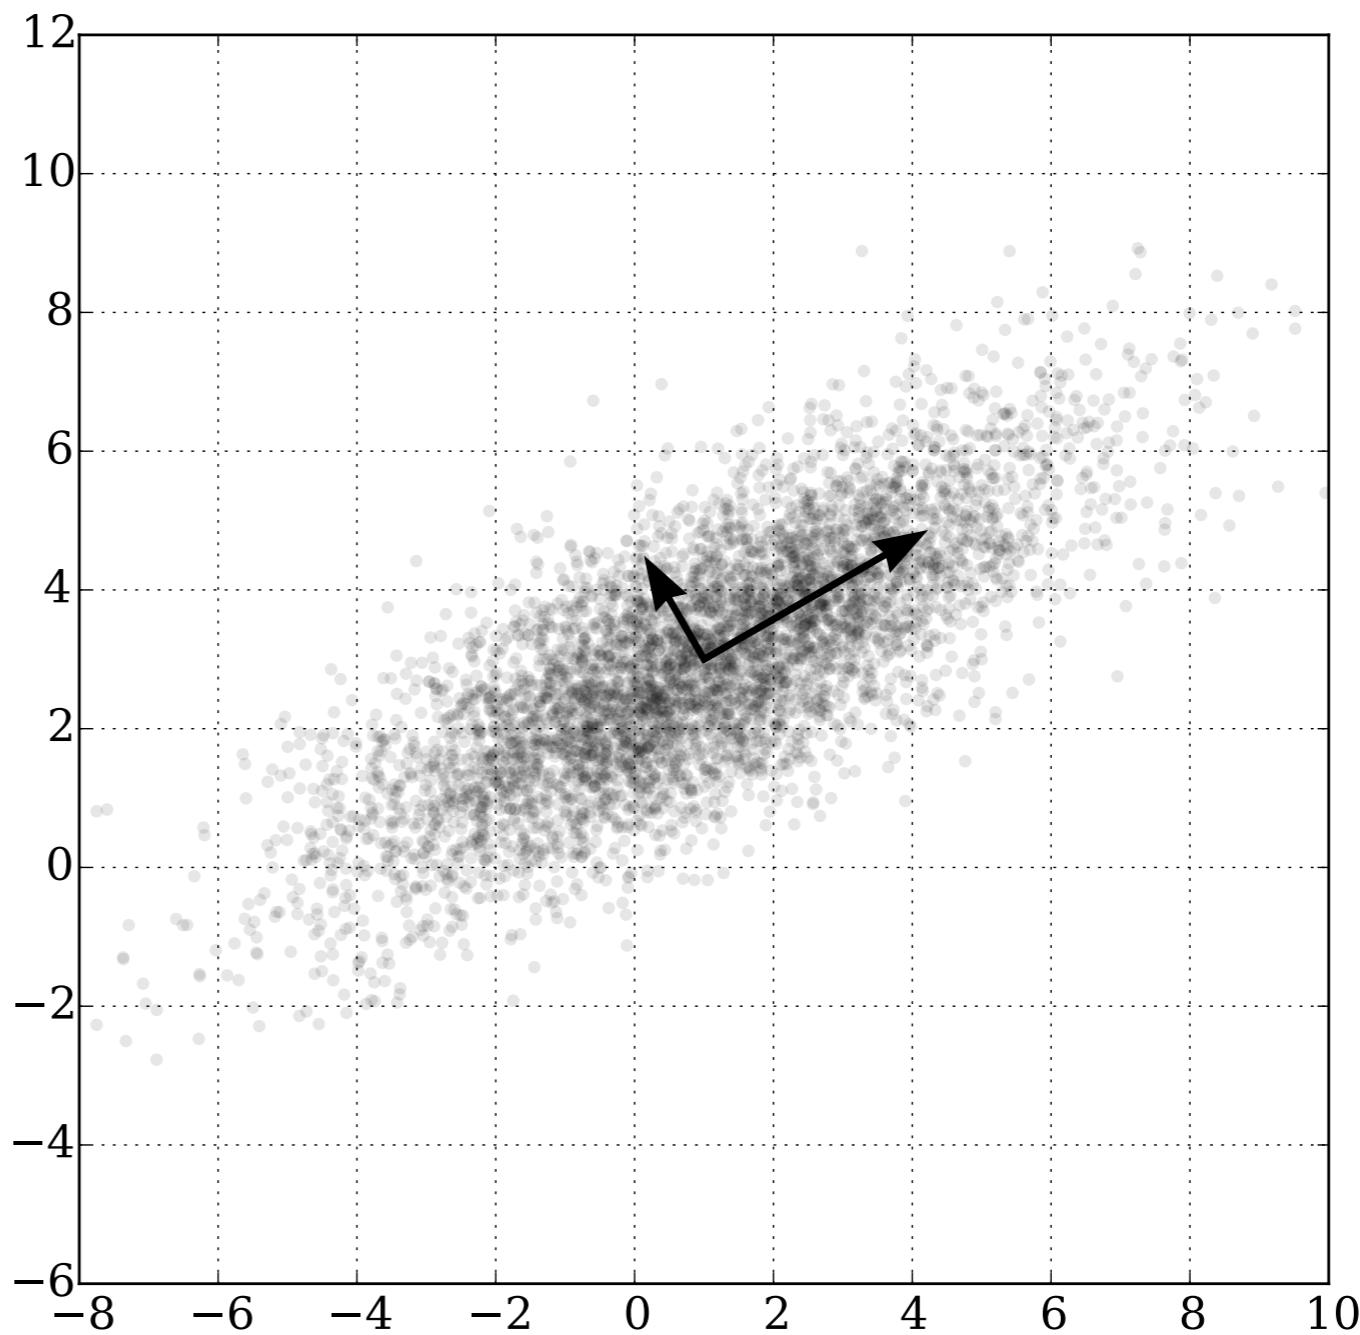


Source: <https://www.turingfinance.com/>

Curse of dimensionality



Principal Component Analysis



t-SNE and UMAP

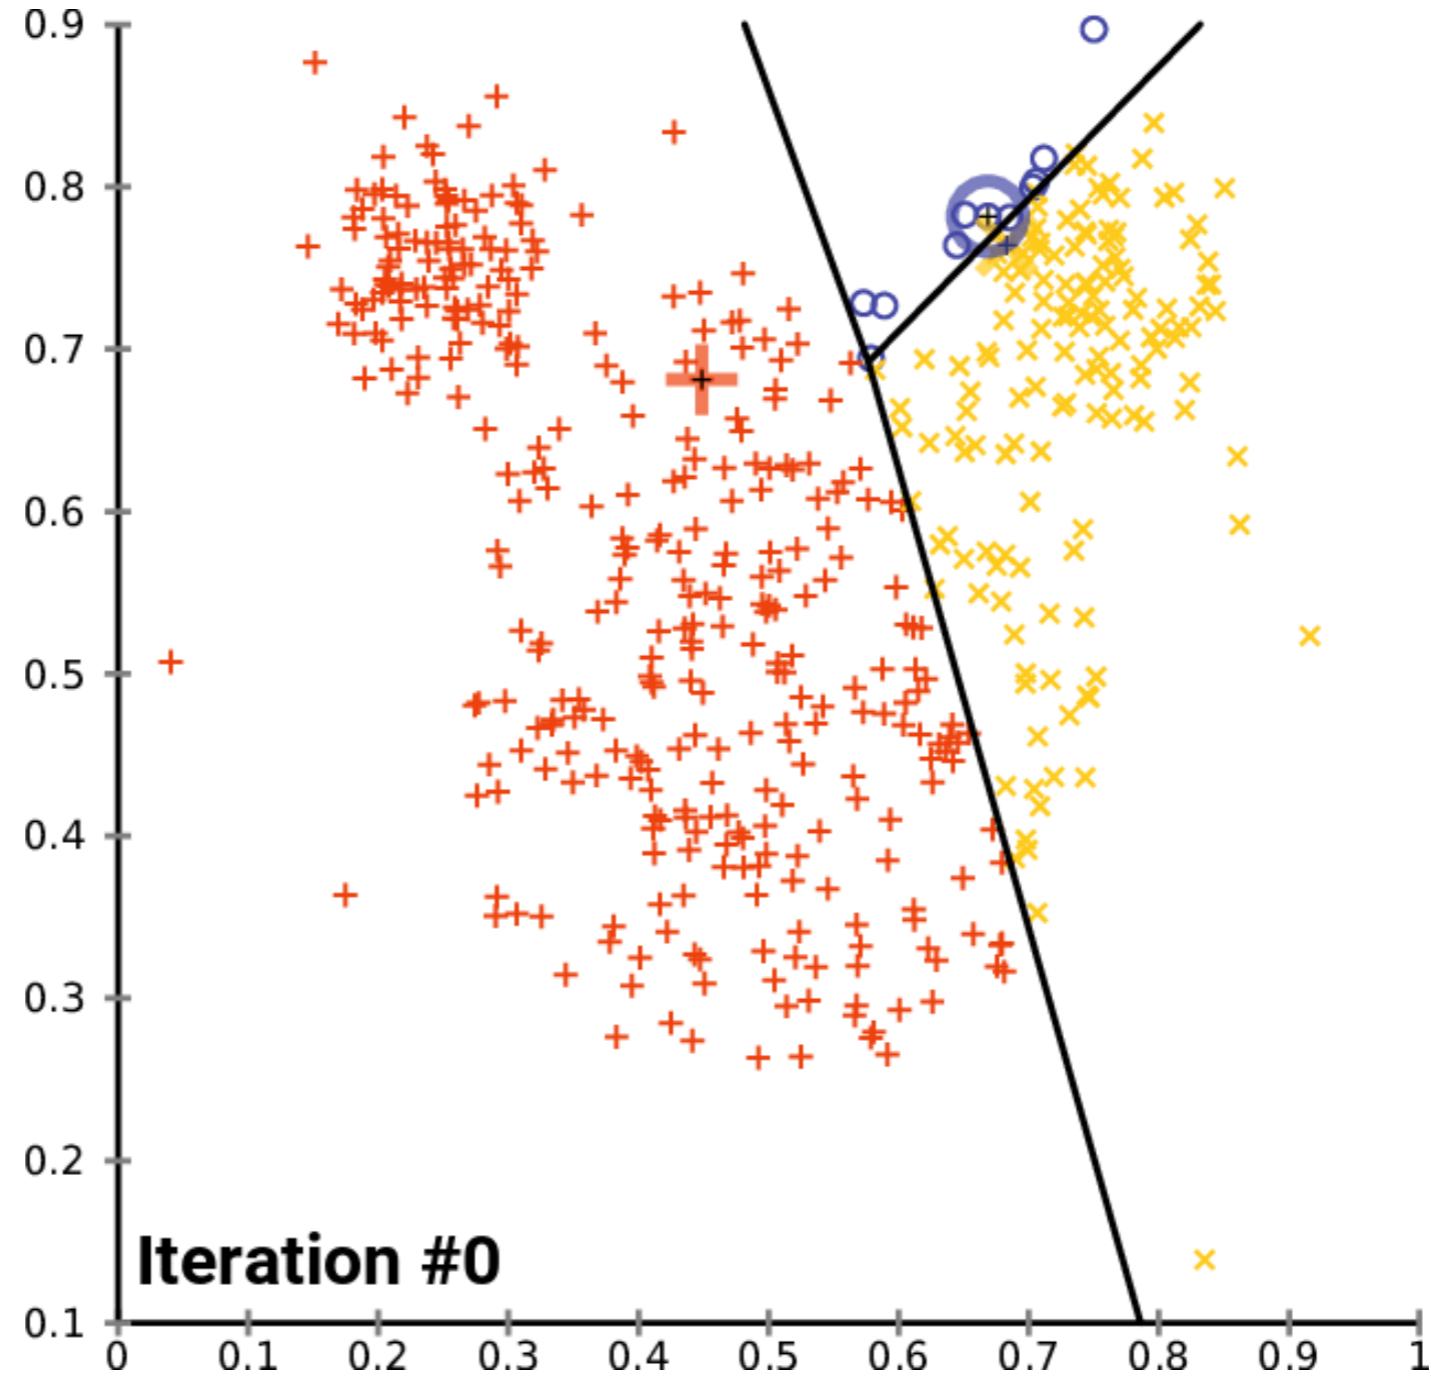
dimensionality reduction



Dimensionality reduction task

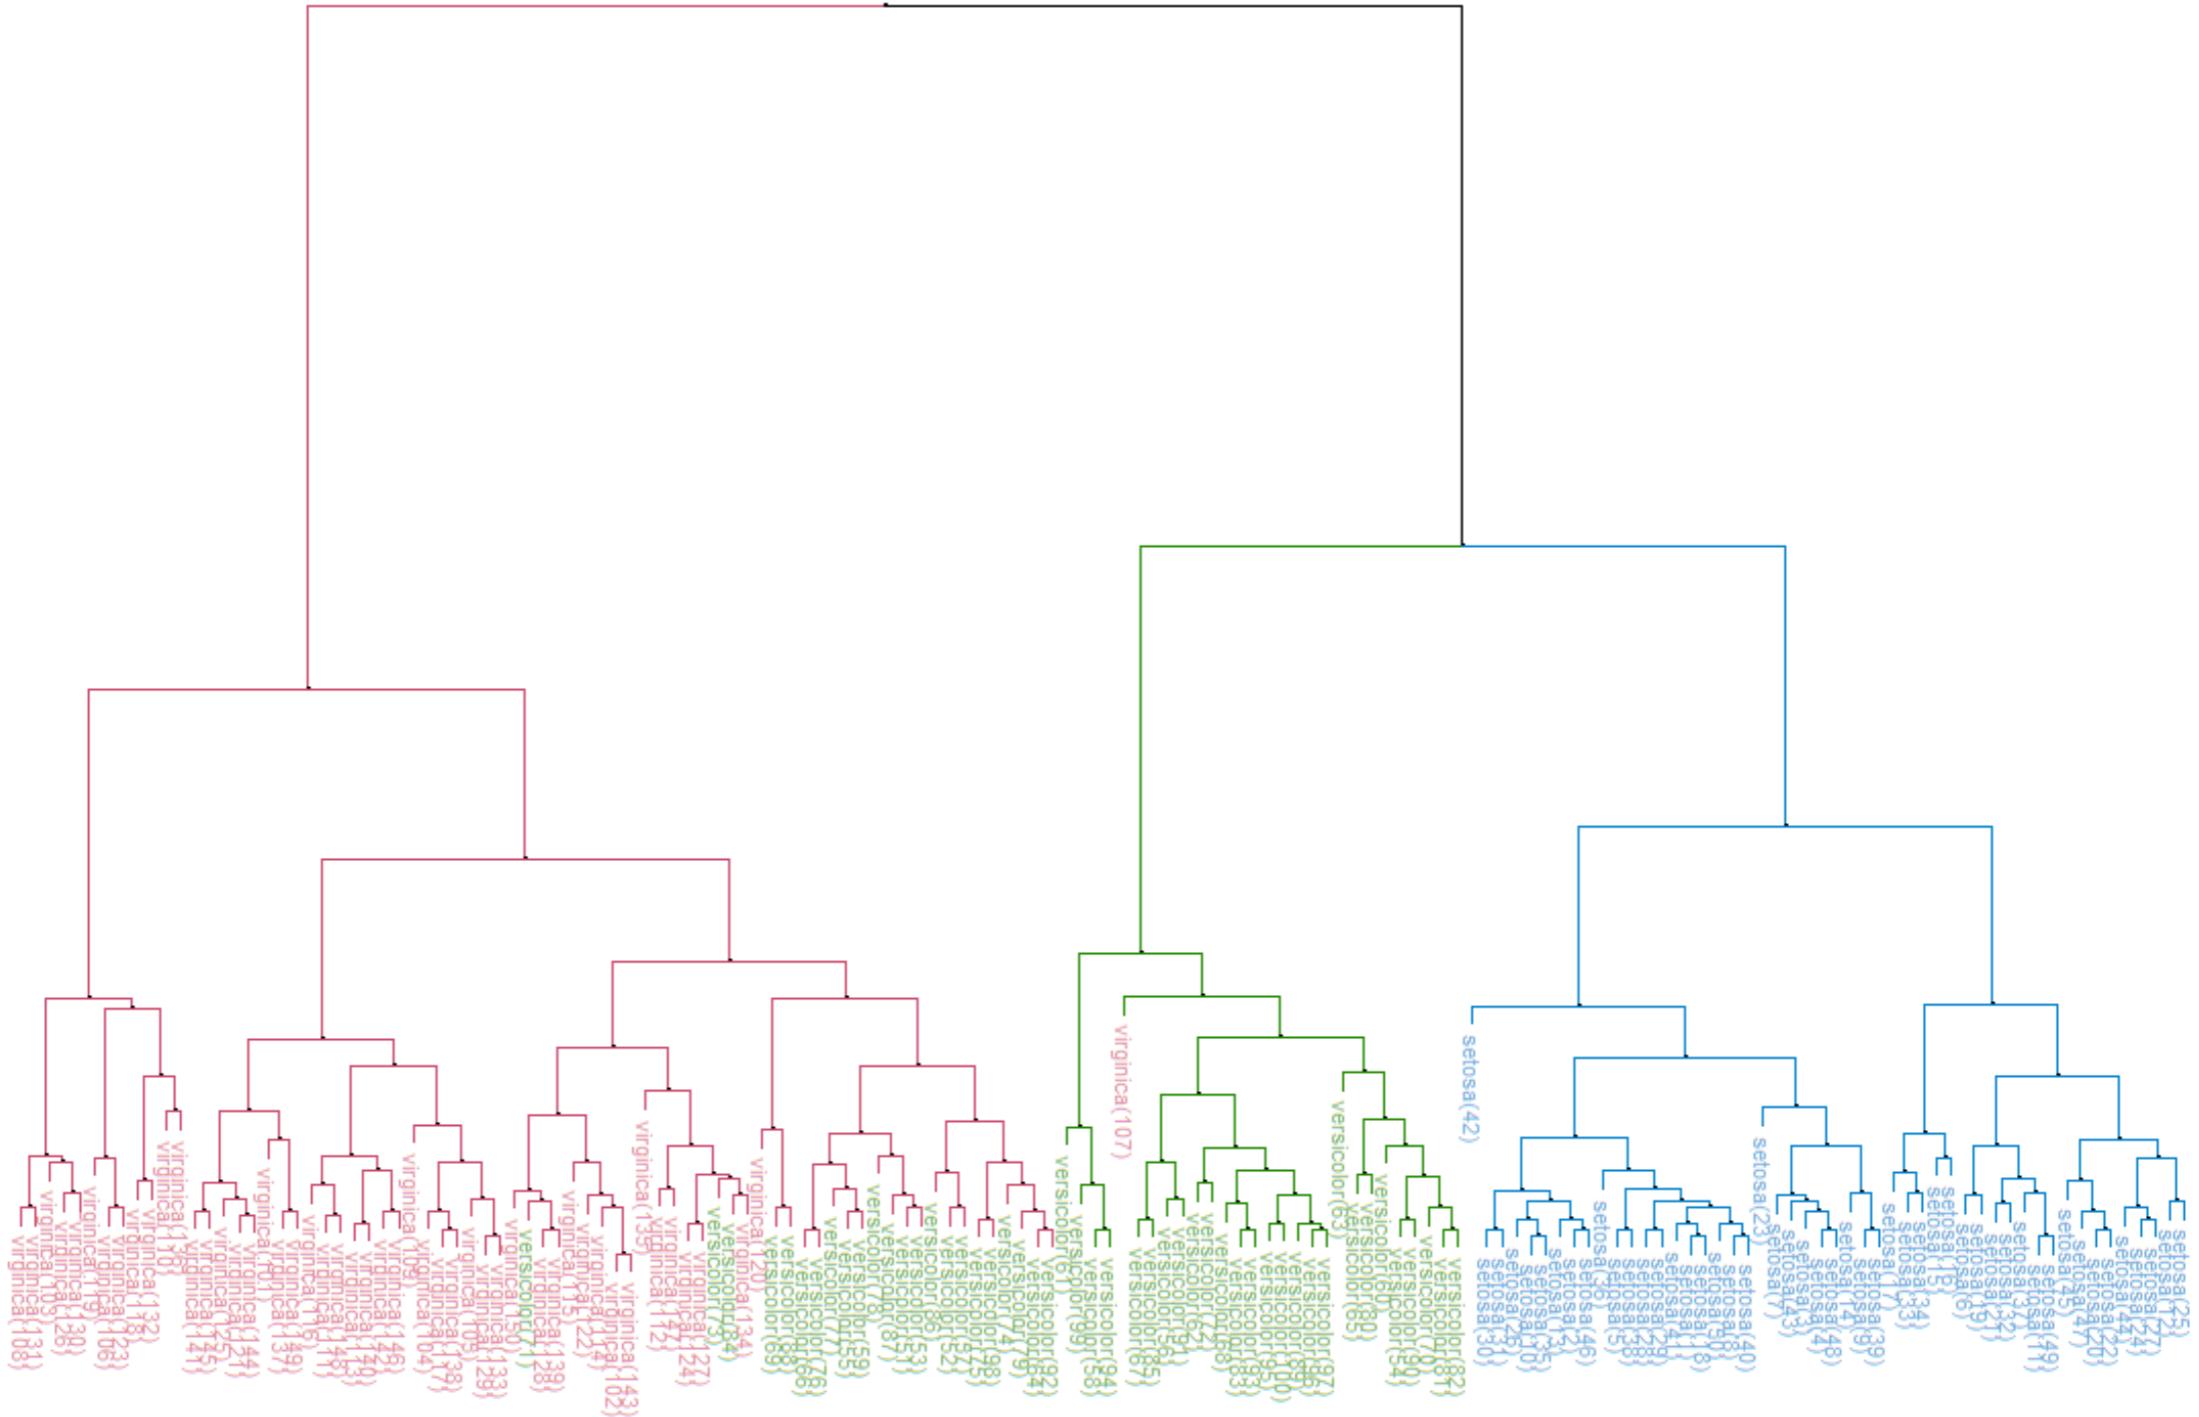
[06-Dim-reduction-assignment.ipynb](#)

k-means clustering



Source: <https://wikipedia.org>

Agglomerative clustering

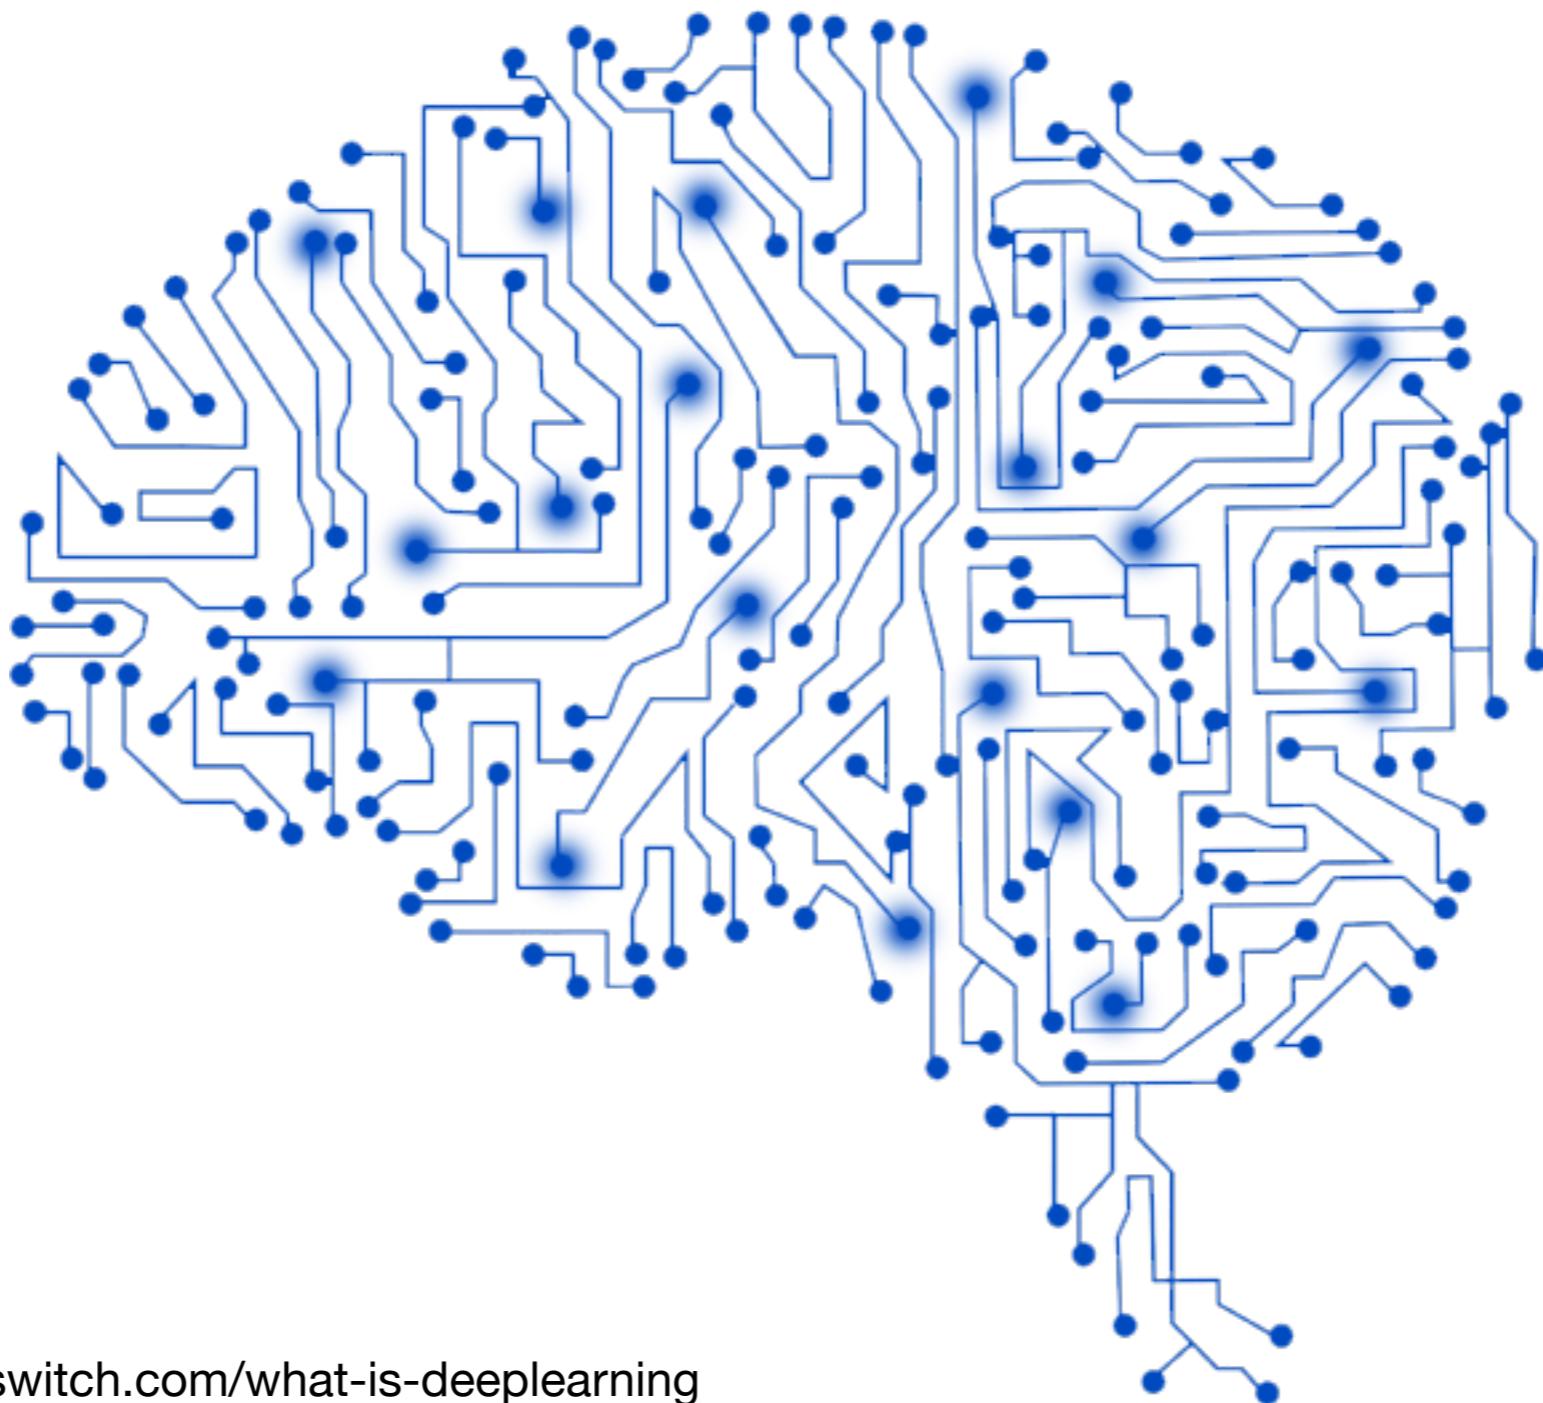


Source: <https://wikipedia.org>

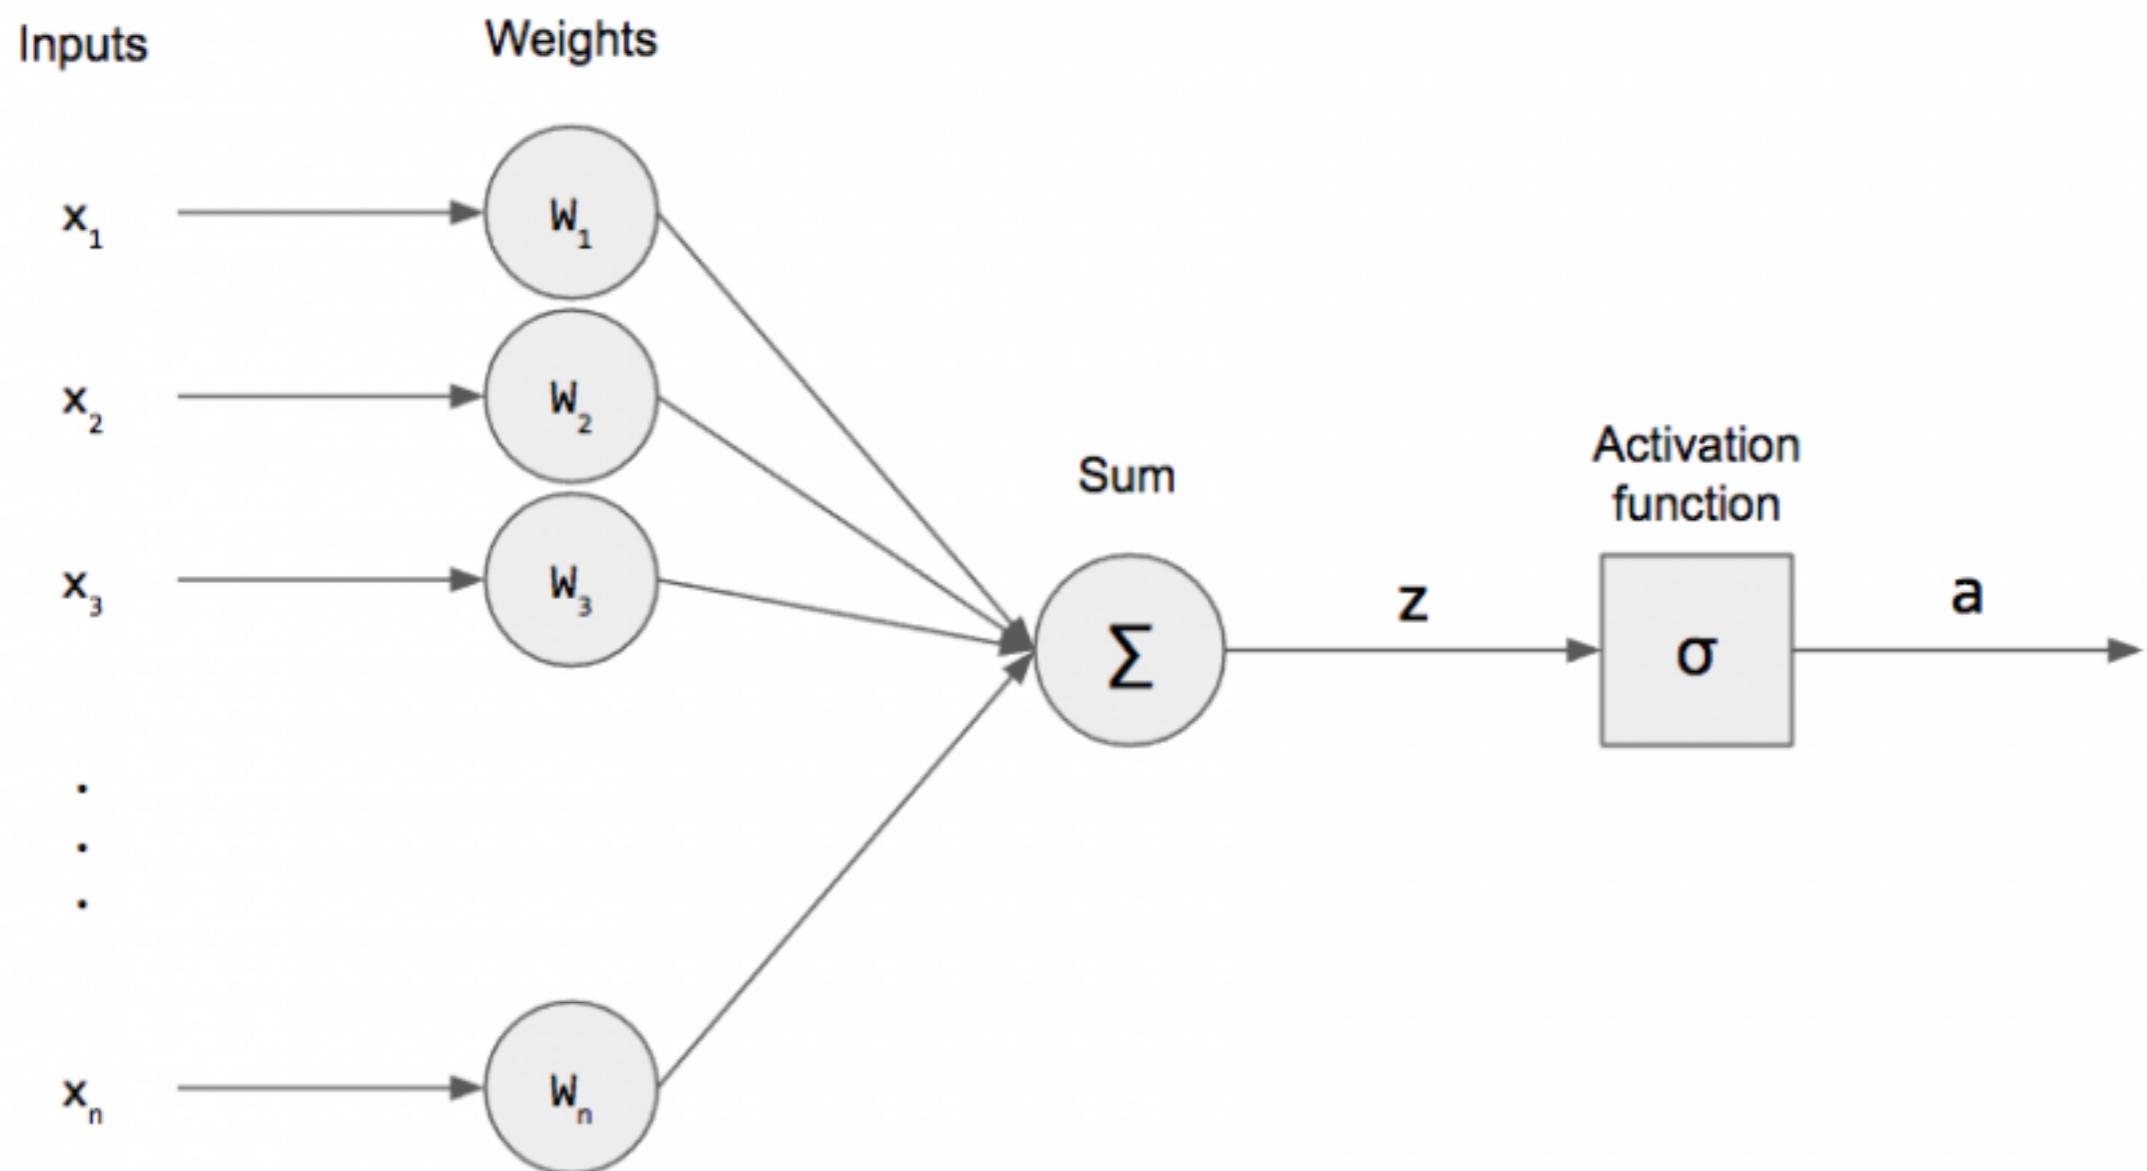
Clustering task

07-Clustering-assignment.ipynb

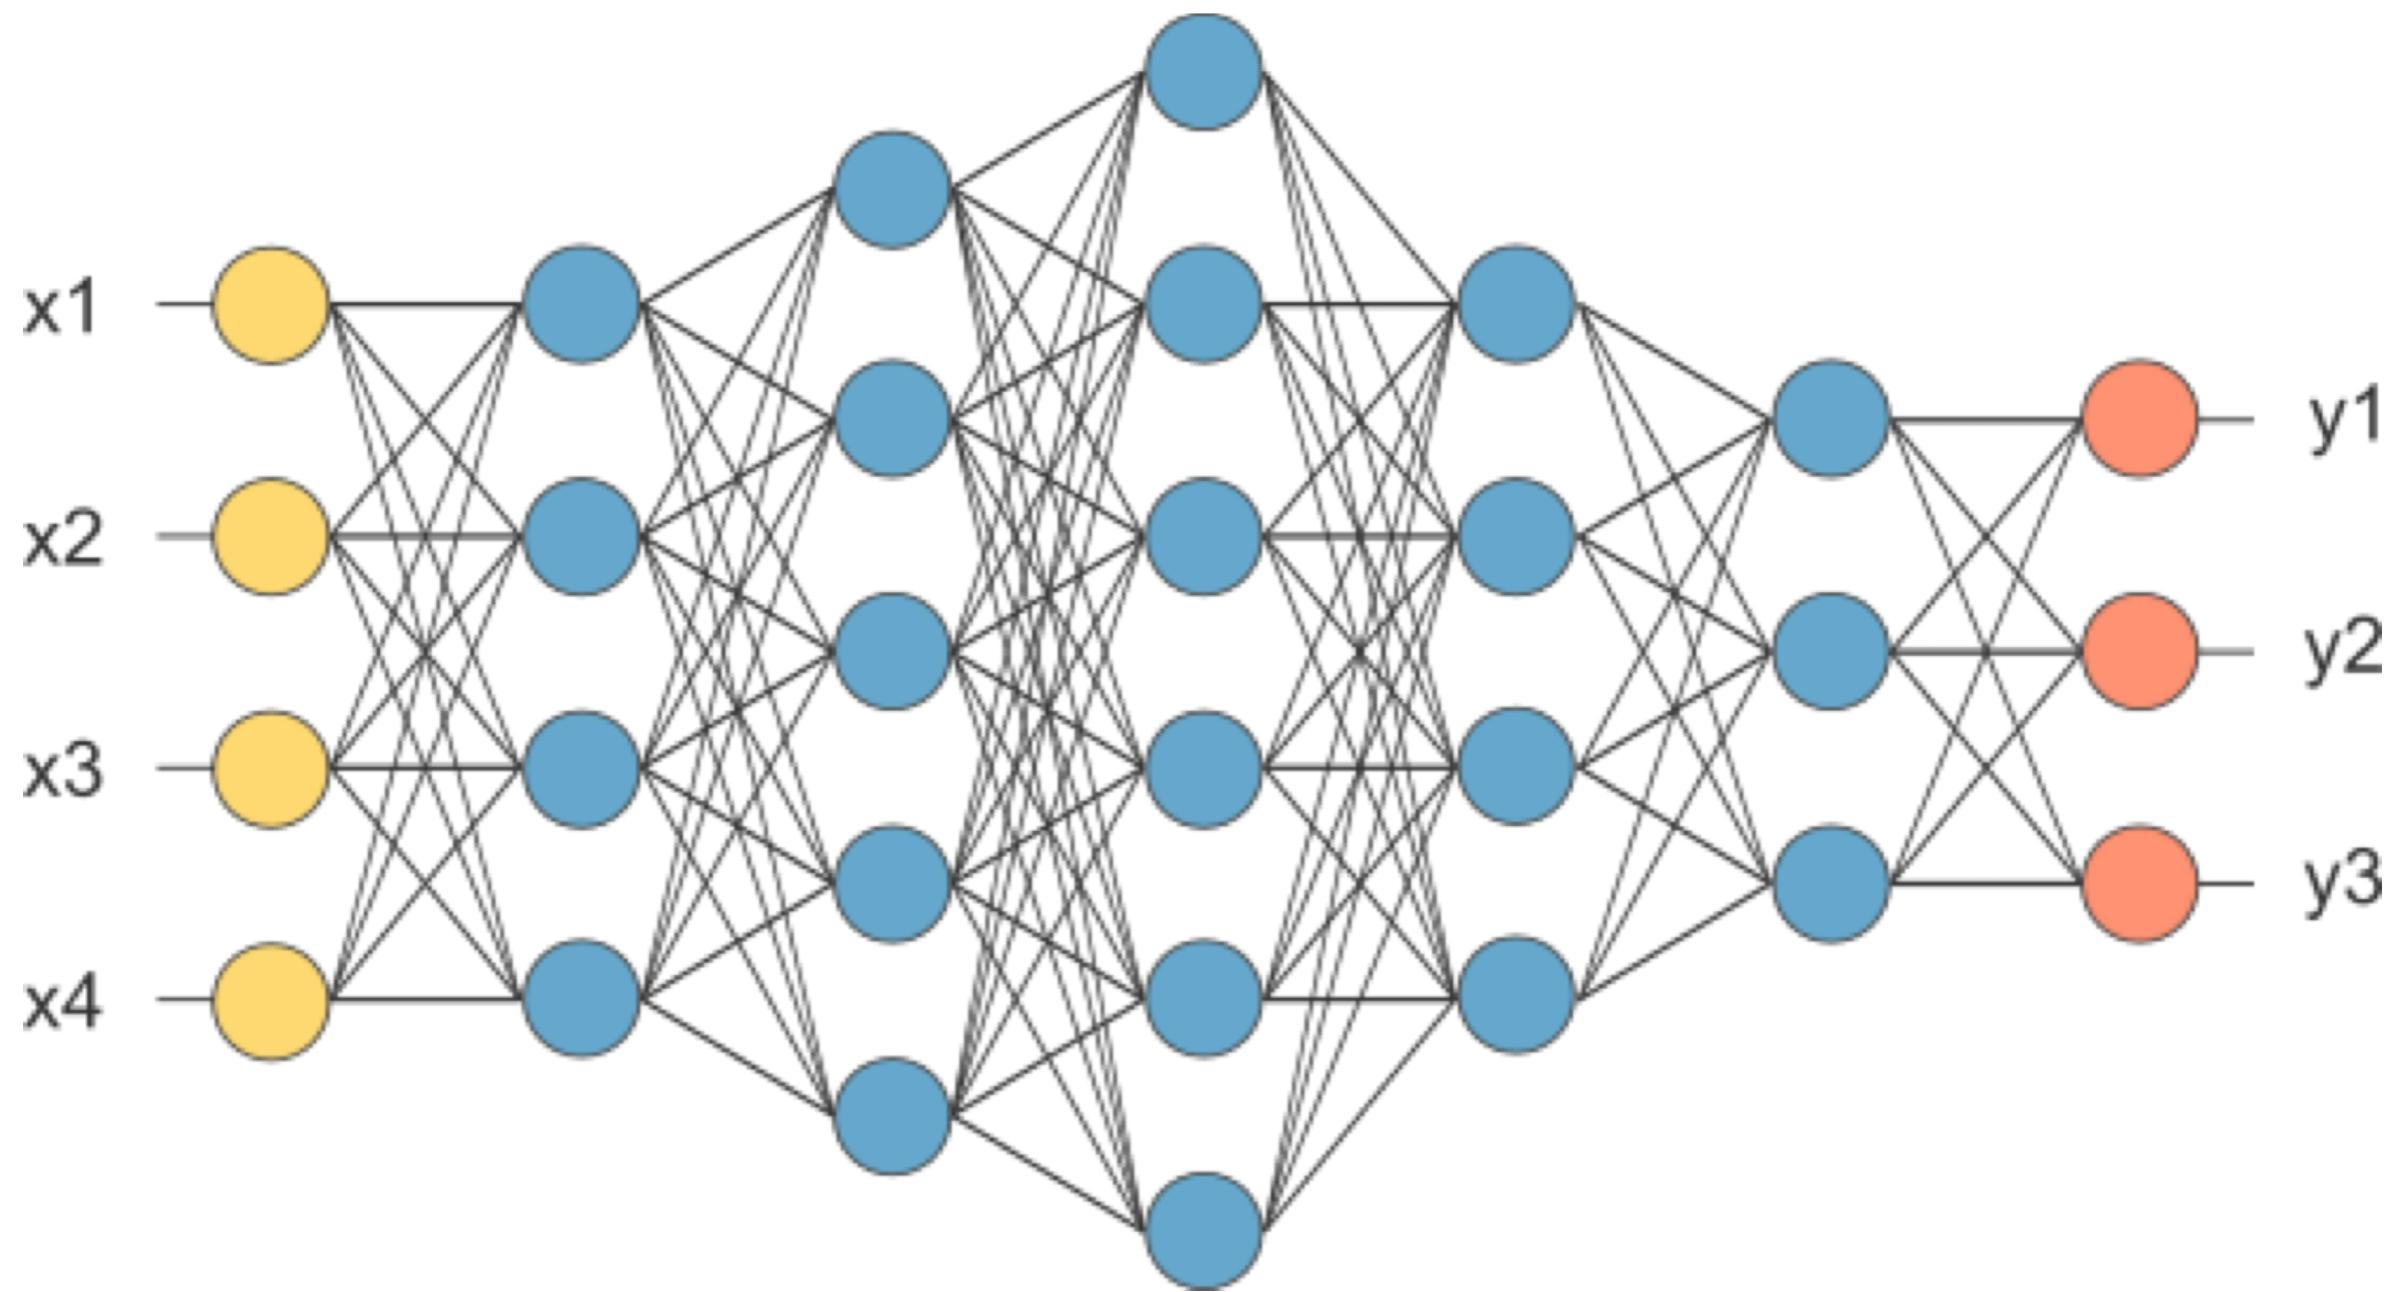
Neural networks and deep learning



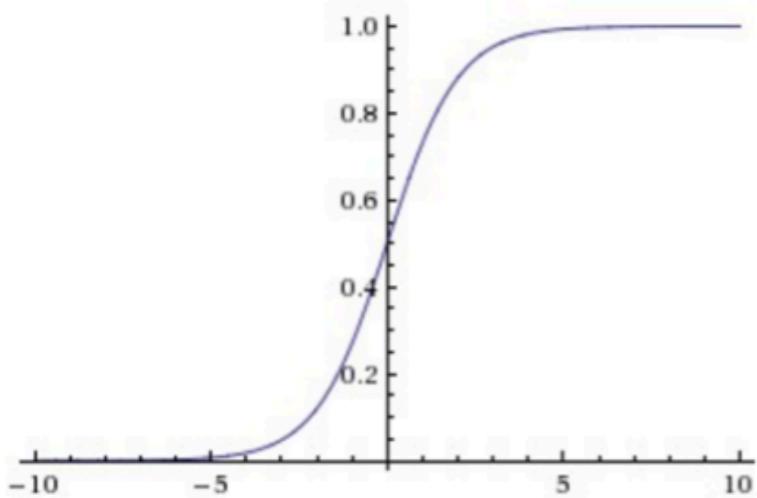
Perceptron



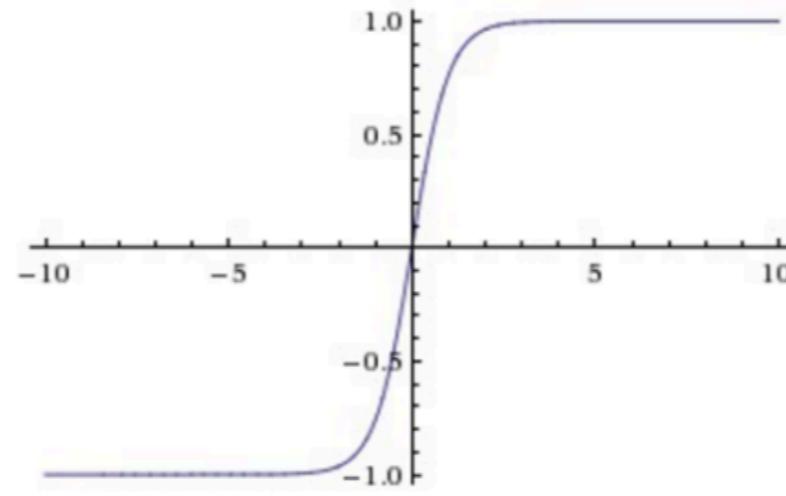
Multilayered Neural Networks



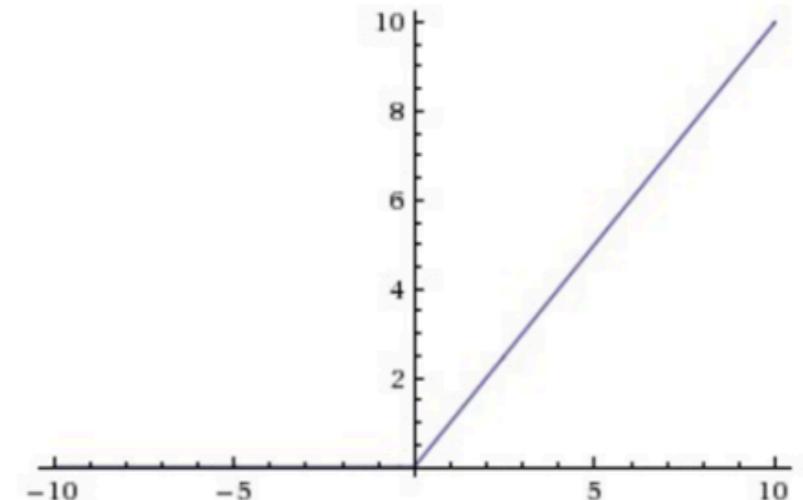
Activation functions



Sigmoid



tanh

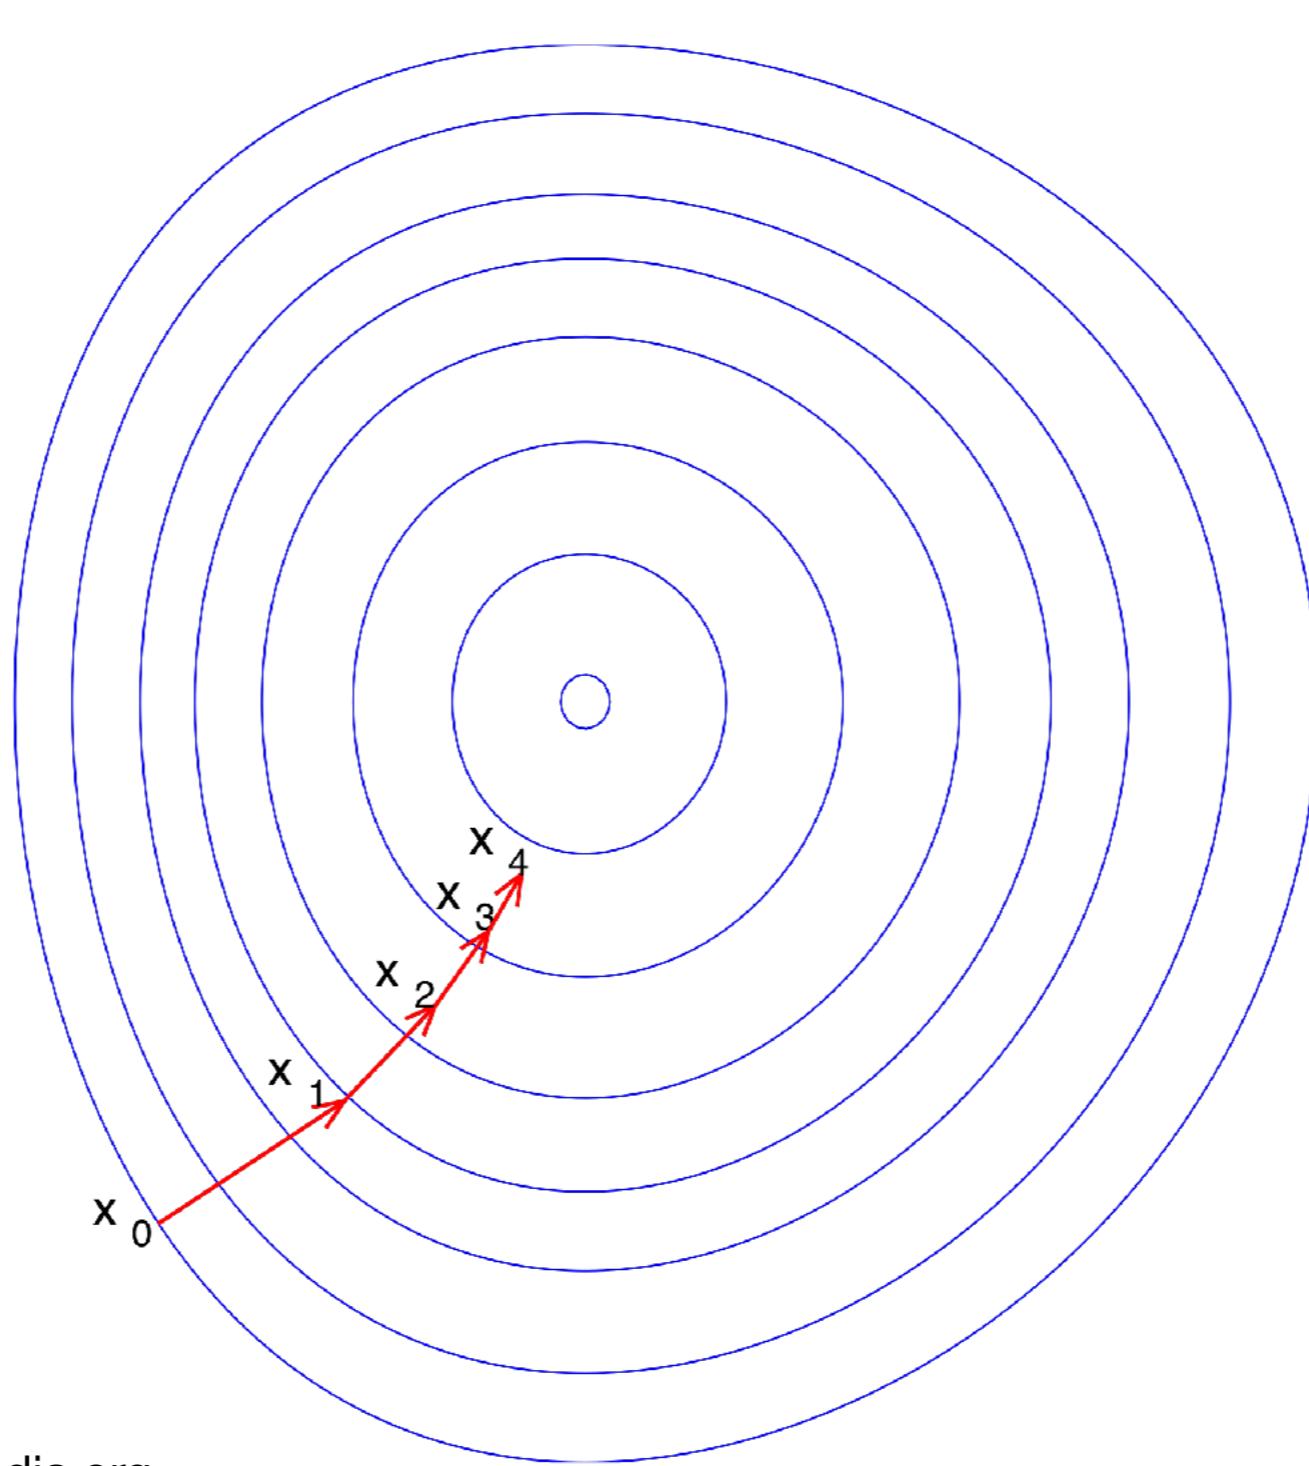


ReLU

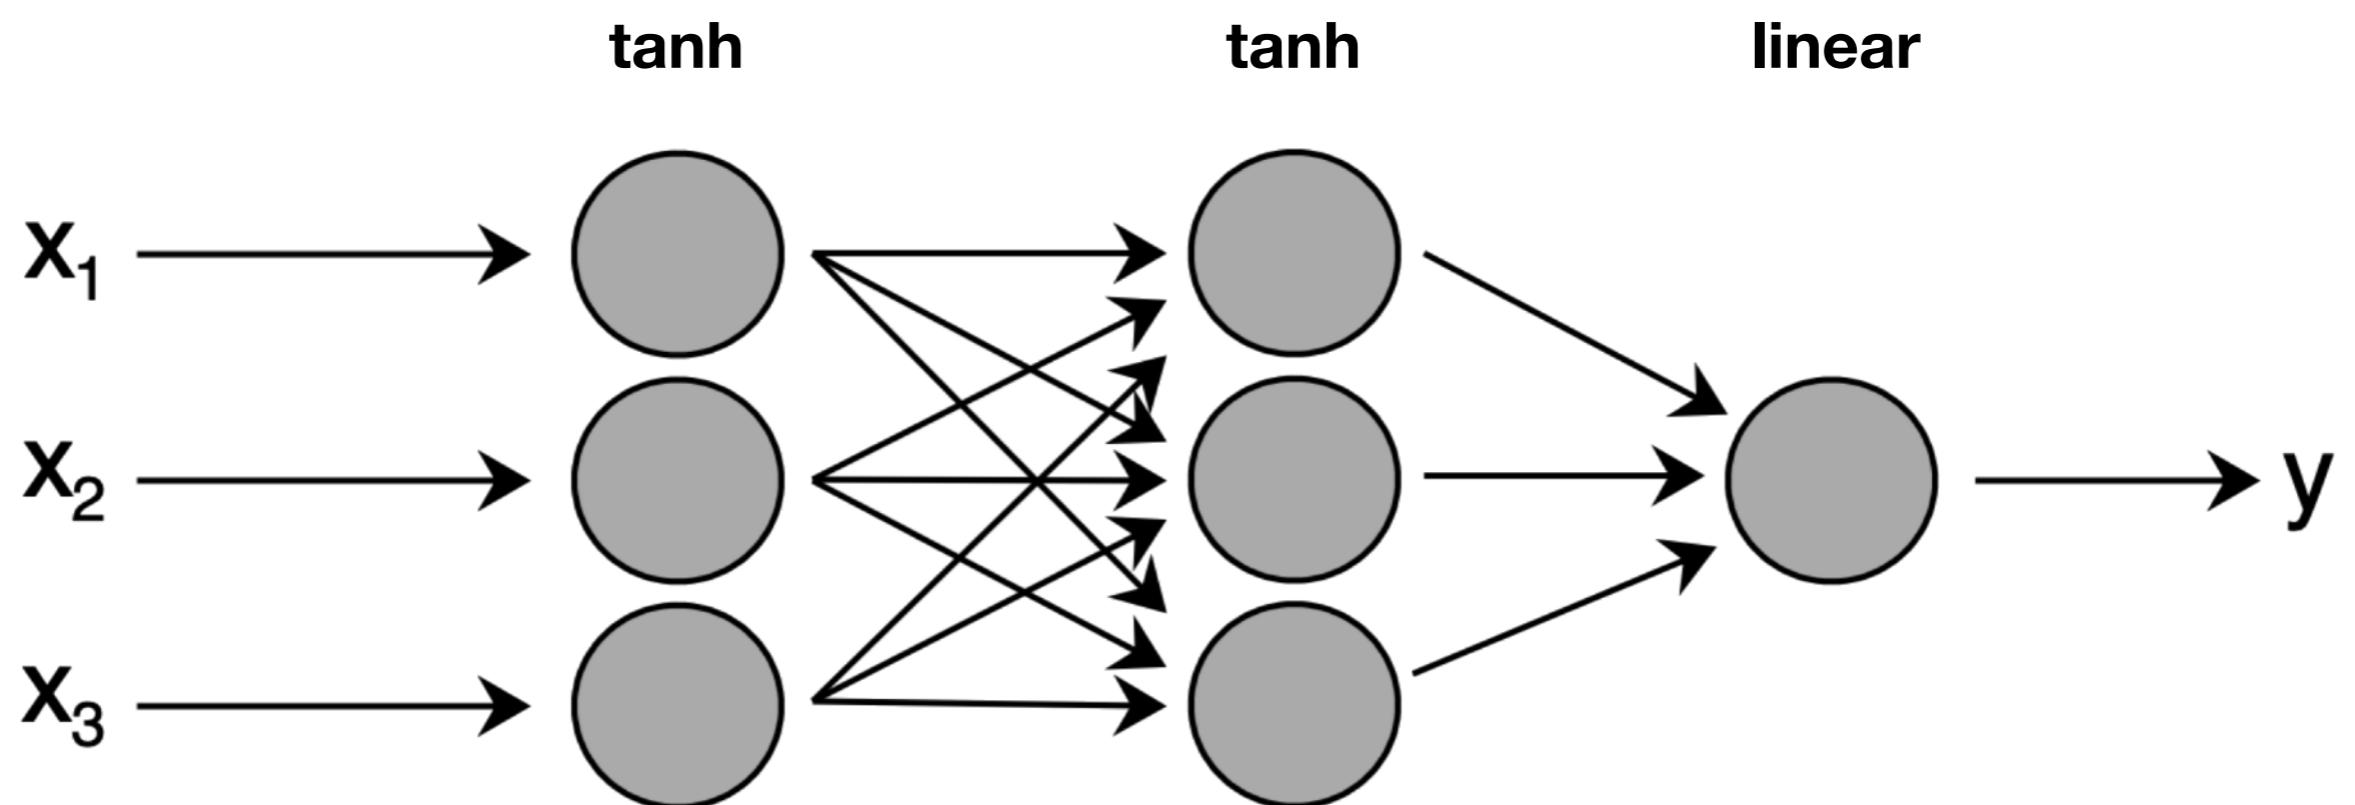
Softmax:

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

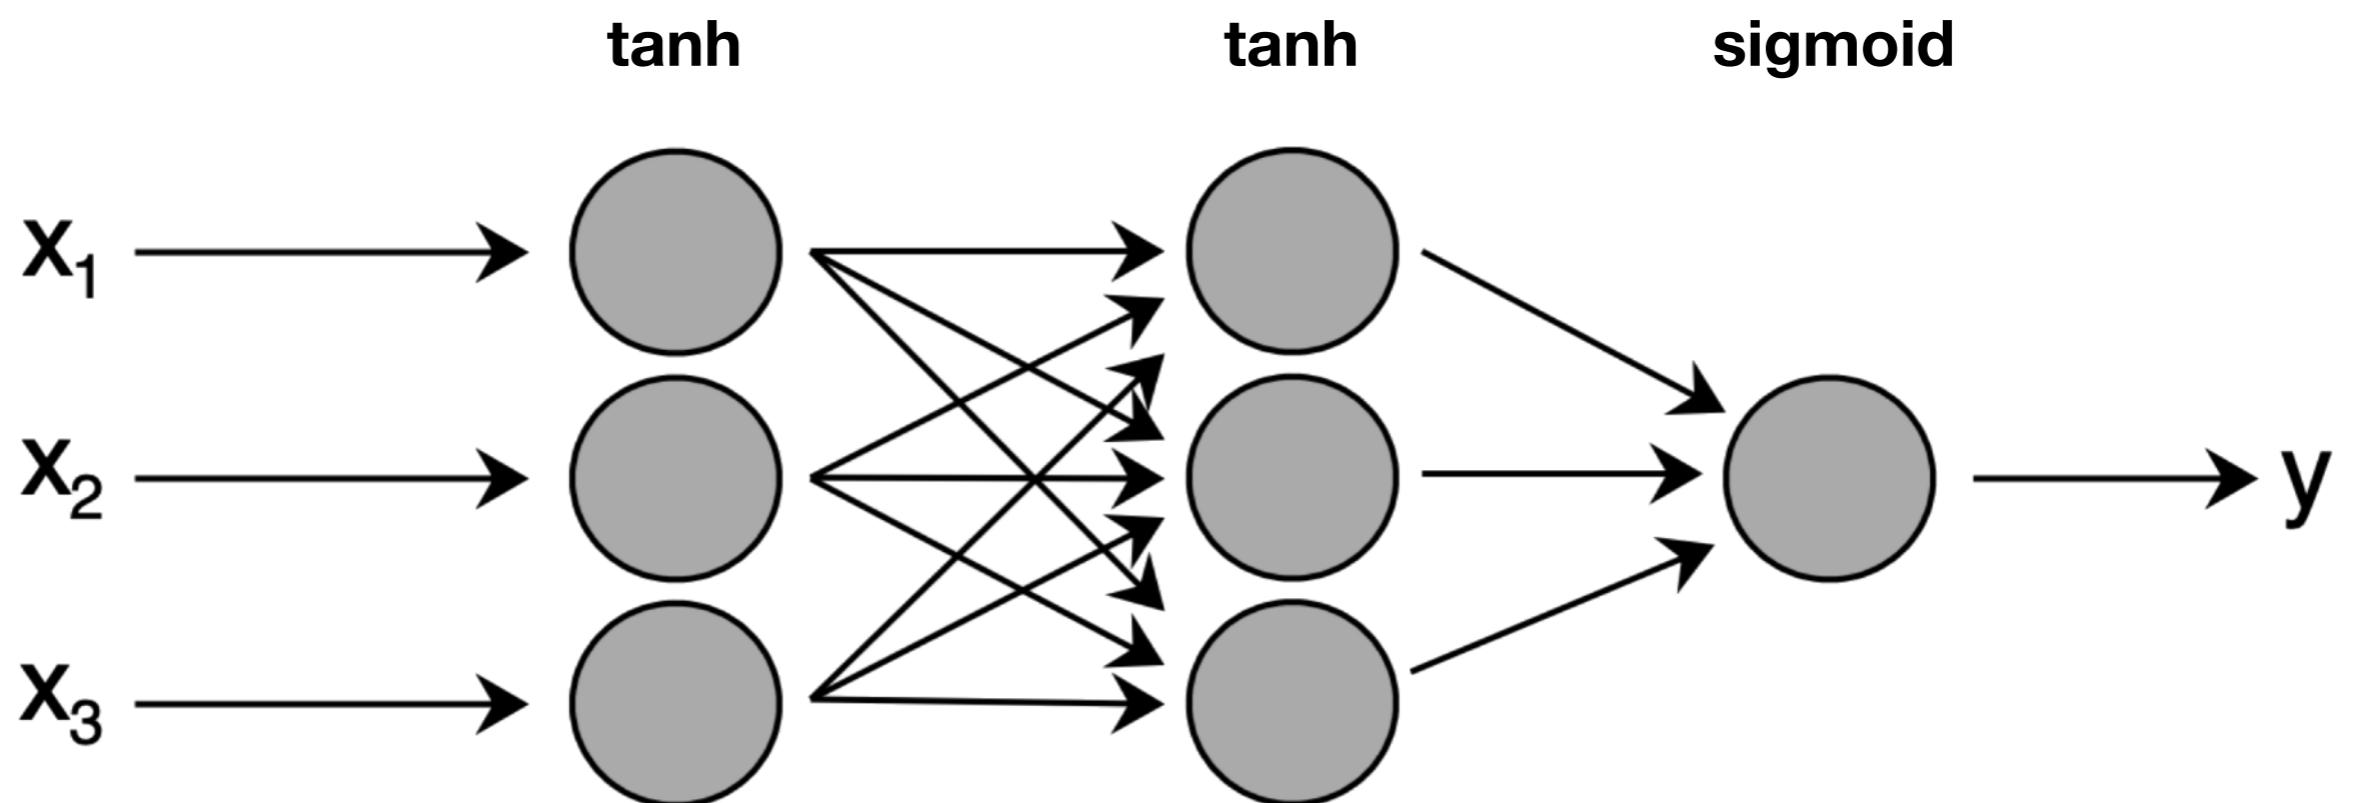
Steepest gradient descent



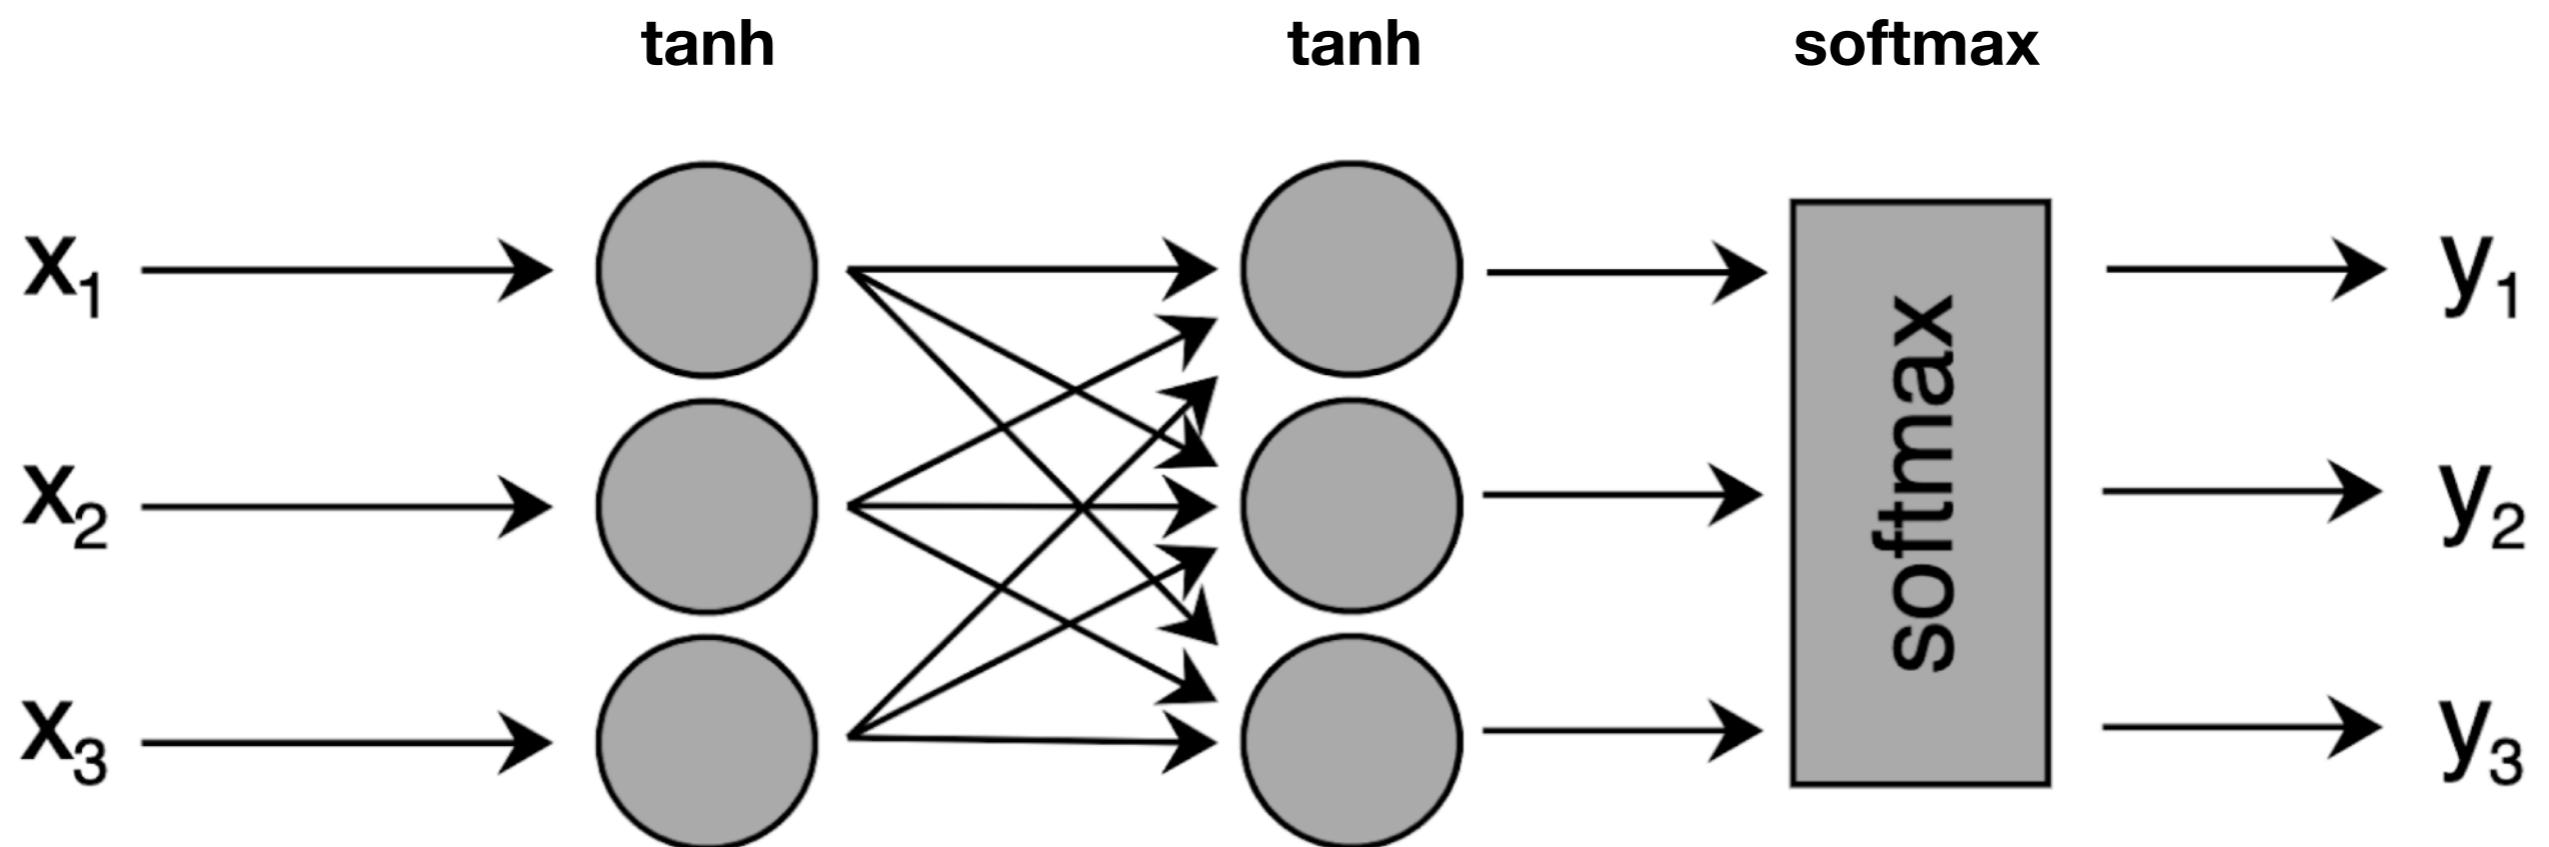
Example - NN for regression



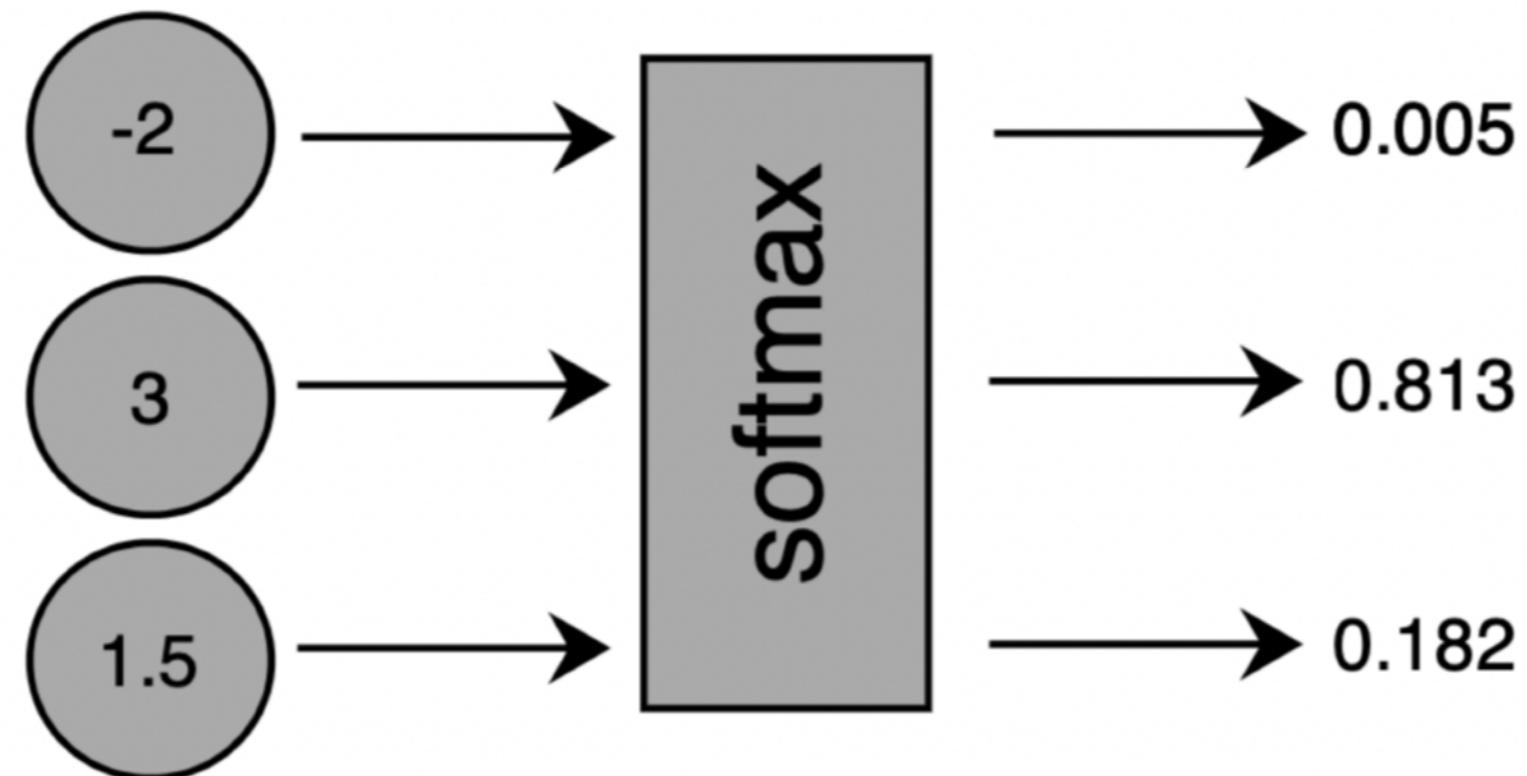
Example - NN for binary classification



Example - NN for multiclass classification



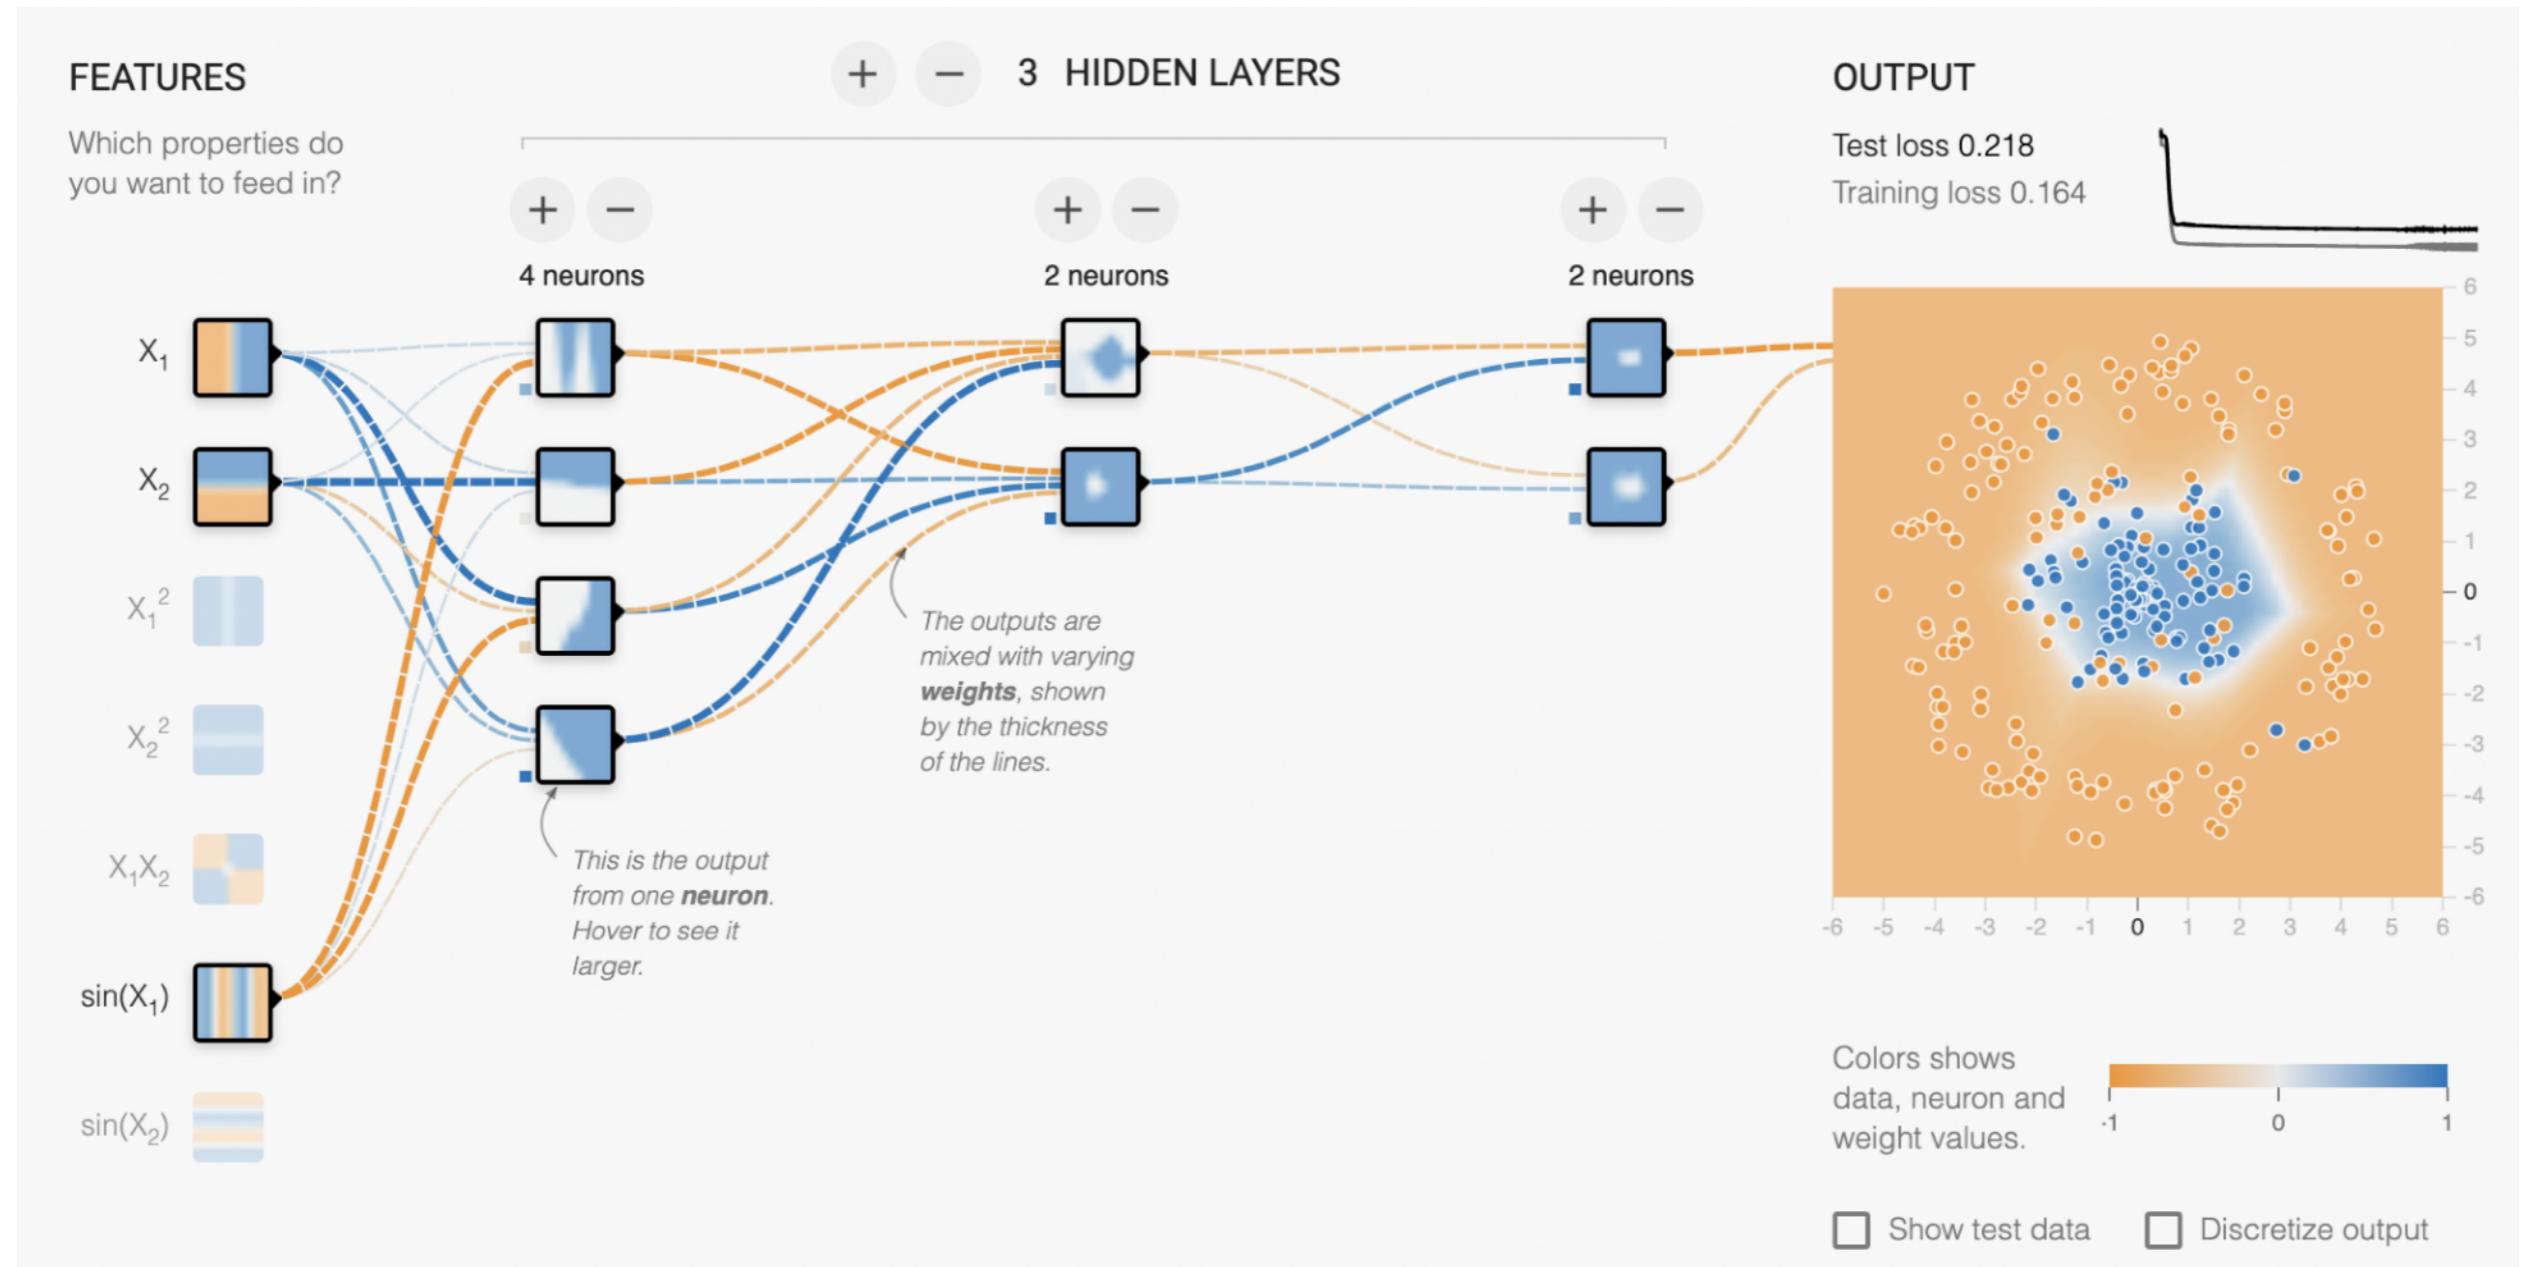
Softmax example



Important terms

- deep learning
- stochastic gradient descent
- batch and mini-batch learning
- epoch

Neural network playground



<https://playground.tensorflow.org/>

Keras tutorial

[08-Keras-introduction.ipynb](#)

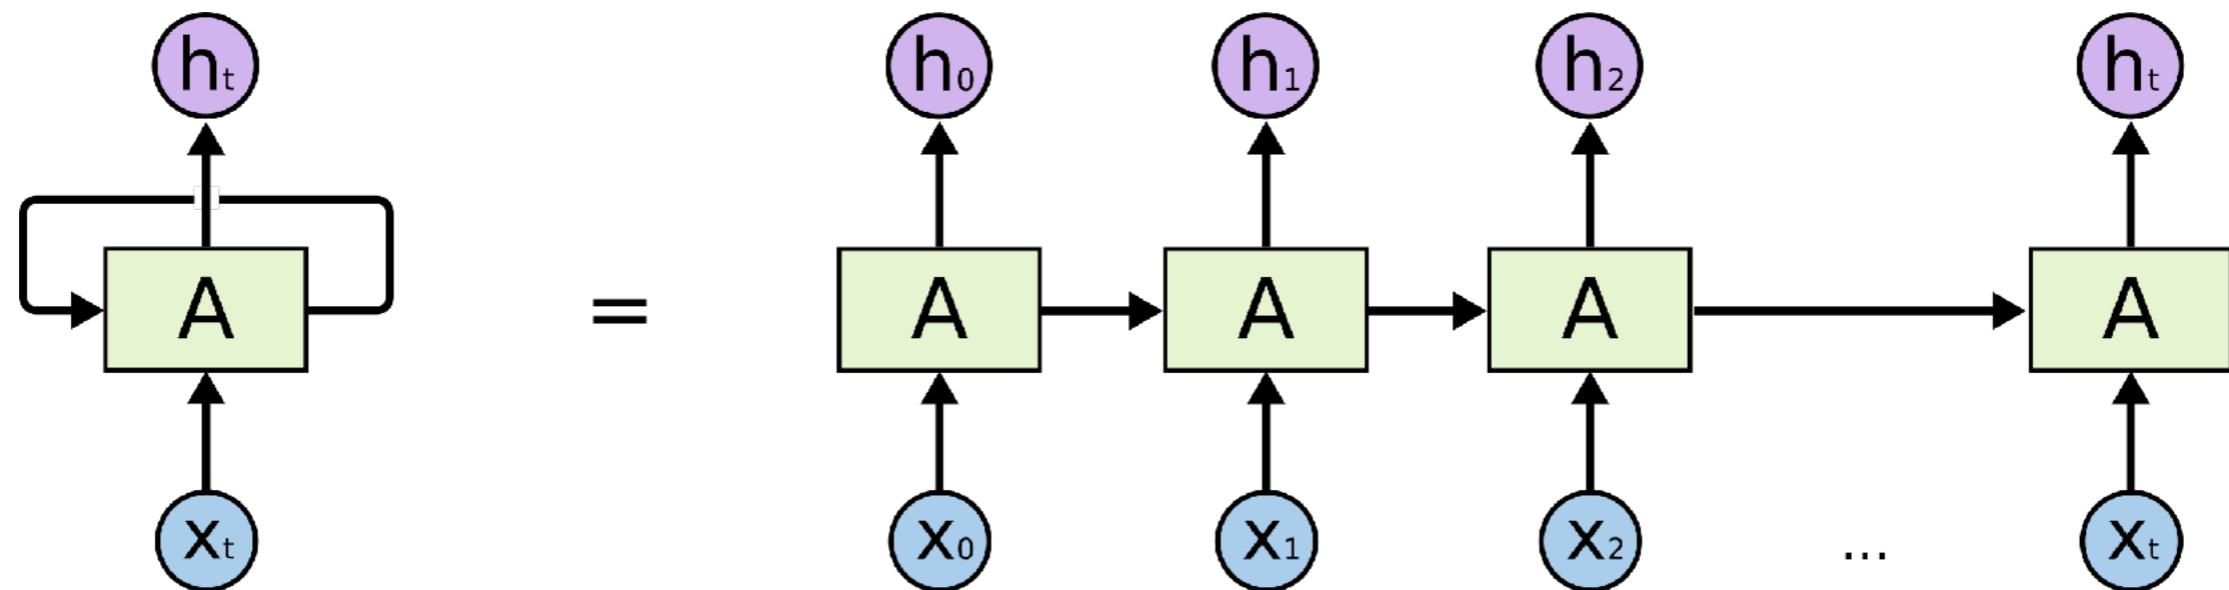
Reimplementation of the classification and regression task using NN

[**09-Classification-nn-assignment.ipynb**](#)

[**10-Regression-nn-assignment.ipynb**](#)

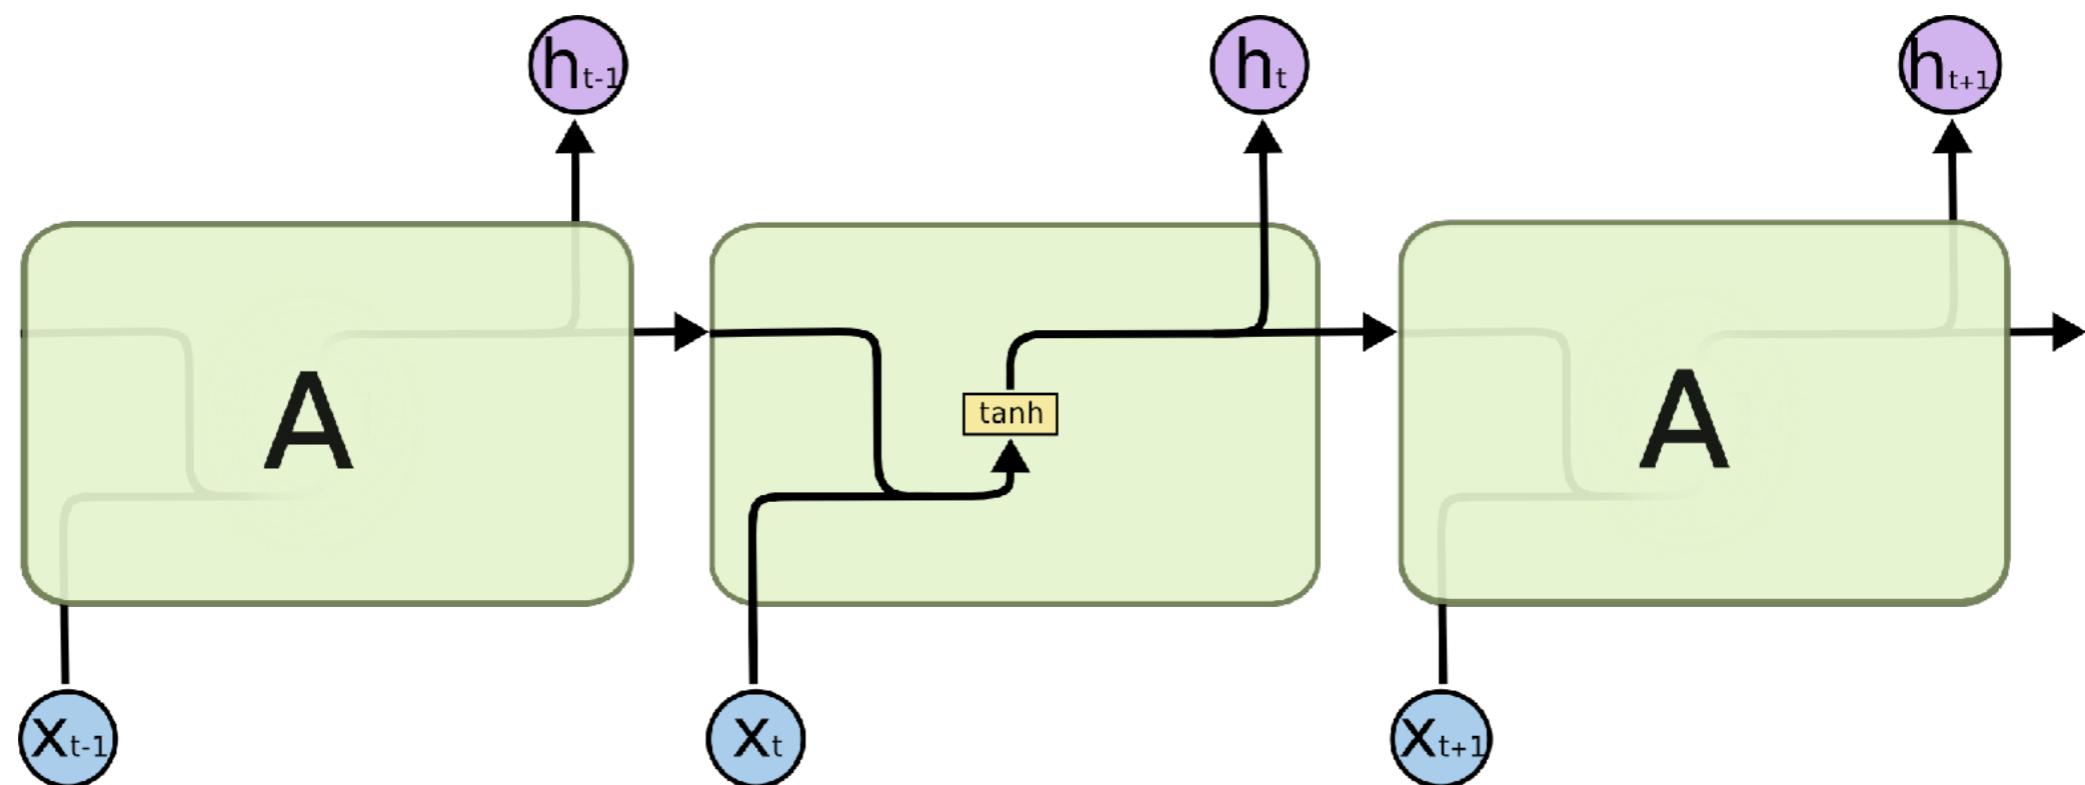
Recurrent Neural Networks

1/2 (bonus)

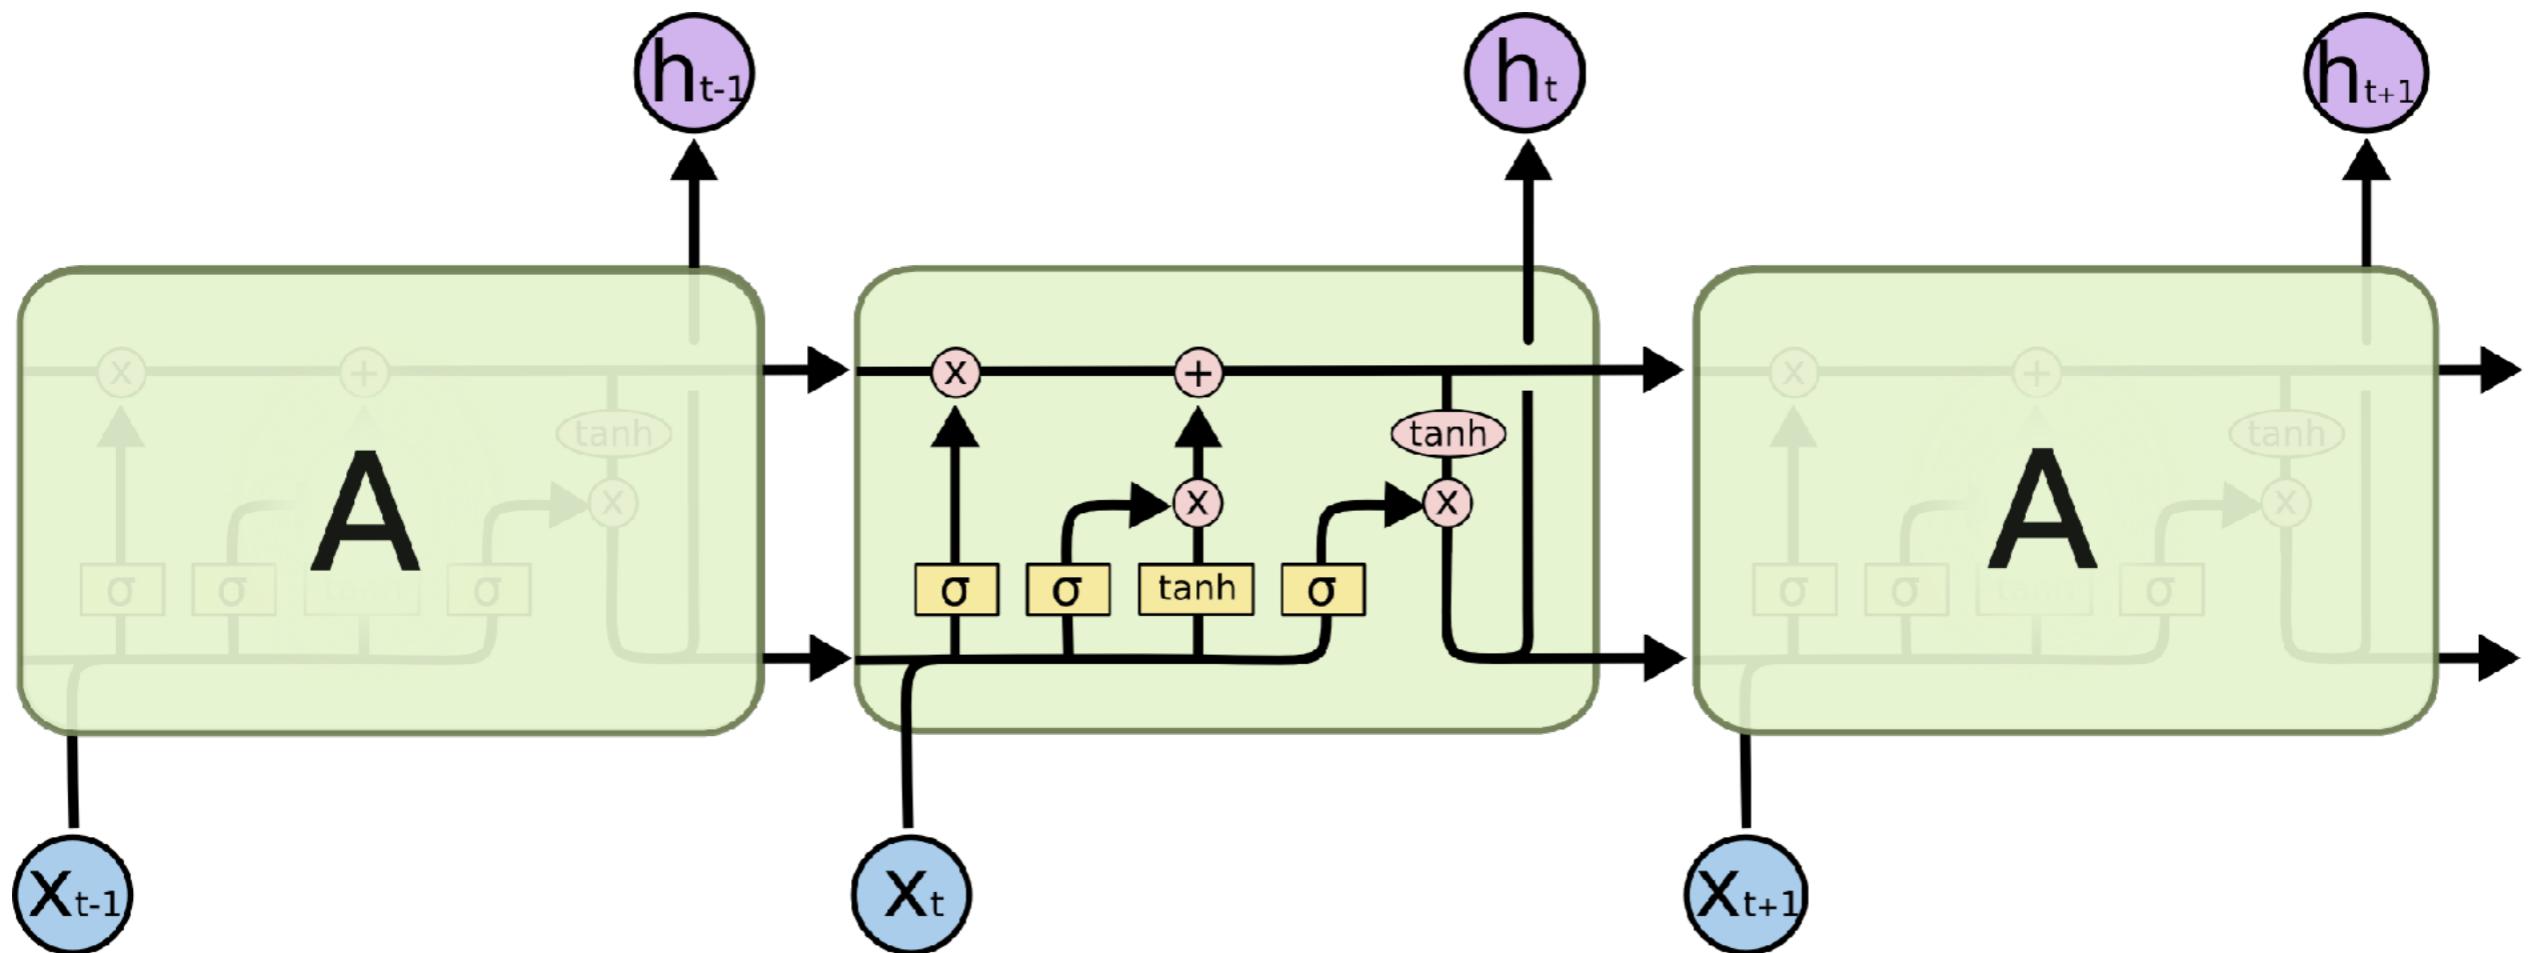


Recurrent Neural Networks

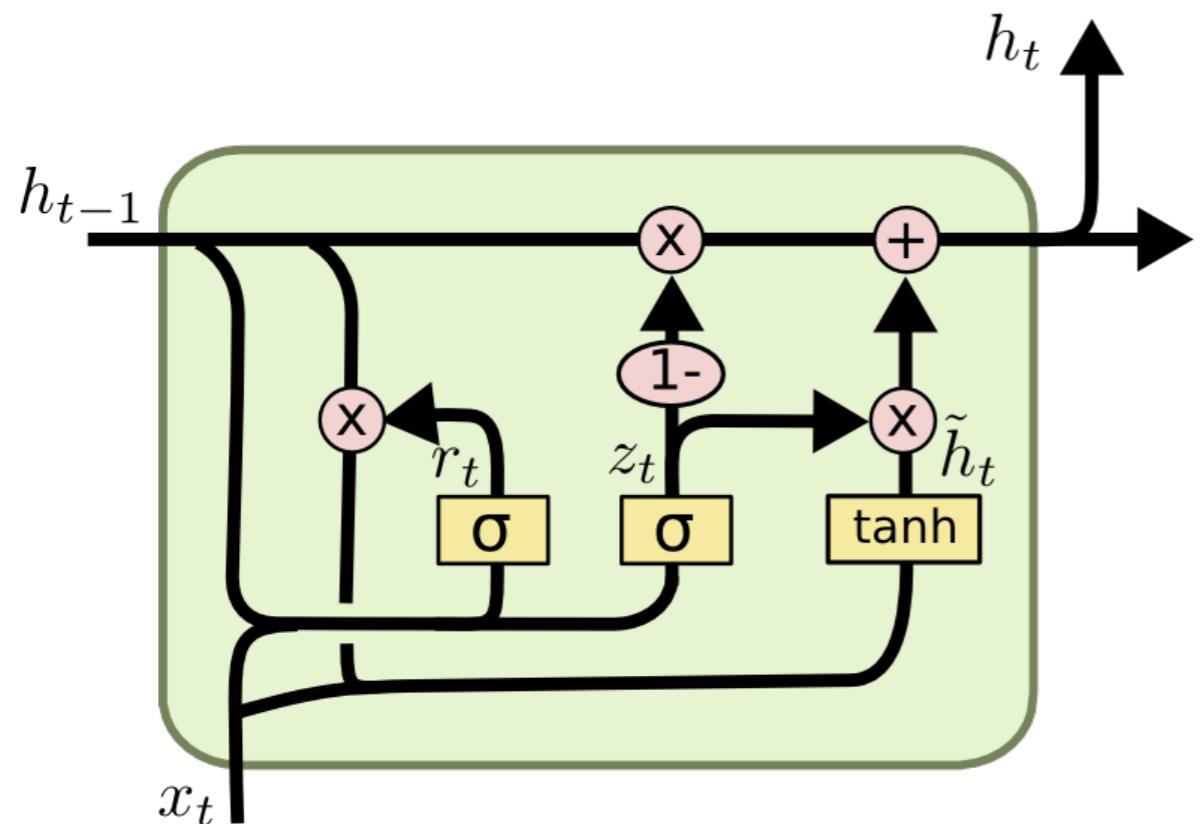
2/2 (bonus)



Long Short-Term Memory (bonus)



Gated Recurrent Unit (bonus)



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

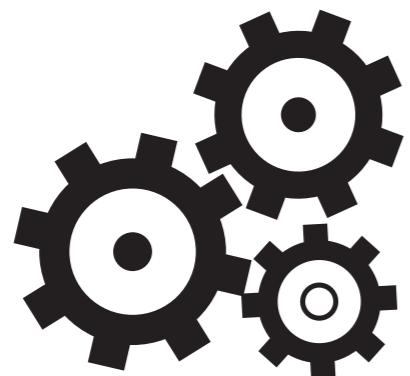
$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Reimplementation of the regression task using Recurrent neural network (bonus)

11-Regression-rnn.ipynb

What next?

<https://www.mlcollege.com/en/#courses>



Machine Learning Prague

ML MACHINE LEARNING
MU meetups

Thank you for your attention

e-mail: jiri@mlcollege.com

Web: www.mlcollege.com

Twitter: @JiriMaterna

Facebook: <https://www.facebook.com/maternajiri>

LinkedIn: <https://www.linkedin.com/in/jirimaterna/>